

AOS-CX 10.06 Quality of Service Guide

6200, 6300, 6400 Switch Series



Hewlett Packard
Enterprise

Part Number: 5200-7718
Published: November 2020
Edition: 1

Notices

The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Confidential computer software. Valid license from Hewlett Packard Enterprise required for possession, use, or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

Links to third-party websites take you outside the Hewlett Packard Enterprise website. Hewlett Packard Enterprise has no control over and is not responsible for information outside the Hewlett Packard Enterprise website.

Acknowledgments

Intel®, Itanium®, Optane™, Pentium®, Xeon®, Intel Inside®, and the Intel Inside logo are trademarks of Intel Corporation in the U.S. and other countries.

Microsoft® and Windows® are either registered trademarks or trademarks of Microsoft Corporation in the United States and/or other countries.

Adobe® and Acrobat® are trademarks of Adobe Systems Incorporated.

Java® and Oracle® are registered trademarks of Oracle and/or its affiliates.

UNIX® is a registered trademark of The Open Group.

All third-party marks are property of their respective owners.



Chapter 1 About this document	5
Applicable products	5
Latest version available online	5
Command syntax notation conventions	5
About the examples	6
Identifying switch ports and interfaces	7
Identifying modular switch components	7
 Chapter 2 QoS overview	 9
End-to-end QoS behavior	9
QoS on the switch	11
QoS trust	12
Port rate limiting	14
Queue profile	14
Schedule profiles	15
Egress queue shaping	16
Egress port shaping	16
Virtual output queues	16
Terms	17
 Chapter 3 QoS configuration	 19
Configuring QoS	19
Configuring expedited forwarding for VoIP traffic	20
Configuring rate limiting	21
Configuring egress queue shaping	23
Configuring egress port shaping	23
Supporting Ethernet 802.1D Class of Service	24
Monitoring queue operation	24
 Chapter 4 QoS commands	 26
apply qos	26
dwrr queue	28
map queue	28
name queue	30
qos cos	30
qos cos-map	31
qos dscp	32
qos dscp-map	34
qos queue-profile	35
qos schedule-profile	36
qos shape	37
qos trust	38
rate-limit	39
show interface queues	40
show interface qos	42
show qos cos-map	43

show qos dscp-map.....	44
show qos queue-profile.....	45
show qos schedule-profile.....	46
show qos trust.....	47
strict queue.....	48

Chapter 5 Support and other resources..... 50

Accessing Aruba Support.....	50
Accessing updates.....	50
Warranty information.....	51
Regulatory information.....	51
Documentation feedback.....	51

This document describes features of the AOS-CX network operating system. It is intended for administrators responsible for installing, configuring, and managing Aruba switches on a network.

Applicable products

This document applies to the following products:

- Aruba 6200 Switch Series (JL724A, JL725A, JL726A, JL727A, JL728A)
- Aruba 6300 Switch Series (JL658A, JL659A, JL660A, JL661A, JL662A, JL663A, JL664A, JL665A, JL666A, JL667A, JL668A, JL762A)
- Aruba 6400 Switch Series (JL741A, R0X26A, R0X27A, R0X29A, R0X30A)

Latest version available online

Updates to this document can occur after initial publication. For the latest versions of product documentation, see the links provided in [Support and other resources](#).

Command syntax notation conventions

Convention	Usage
<code>example-text</code>	Identifies commands and their options and operands, code examples, filenames, pathnames, and output displayed in a command window. Items that appear like the example text in the previous column are to be entered exactly as shown and are required unless enclosed in brackets ([]).
example-text	In code and screen examples, indicates text entered by a user.
Any of the following: • <i><example-text></i> • <code><example-text></code> • <i>example-text</i> • <code>example-text</code>	Identifies a placeholder—such as a parameter or a variable—that you must substitute with an actual value in a command or in code: • For output formats where italic text cannot be displayed, variables are enclosed in angle brackets (< >). Substitute the text—including the enclosing angle brackets—with an actual value. • For output formats where italic text can be displayed, variables might or might not be enclosed in angle brackets. Substitute the text including the enclosing angle brackets, if any, with an actual value.
	Vertical bar. A logical OR that separates multiple items from which you can choose only one.
	Any spaces that are on either side of the vertical bar are included for readability and are not a required part of the command syntax.

Table Continued

Convention	Usage
{ }	Braces. Indicates that at least one of the enclosed items is required.
[]	Brackets. Indicates that the enclosed item or items are optional.
... or	Ellipsis:
...	- In code and screen examples, a vertical or horizontal ellipsis indicates an omission of information. - In syntax using brackets and braces, an ellipsis indicates items that can be repeated. When an item followed by ellipses is enclosed in brackets, zero or more items can be specified.

About the examples

Examples in this document are representative and might not match your particular switch or environment.

The slot and port numbers in this document are for illustration only and might be unavailable on your switch.

Understanding the CLI prompts

When illustrating the prompts in the command line interface (CLI), this document uses the generic term `switch`, instead of the host name of the switch. For example:

```
switch>
```

The CLI prompt indicates the current command context. For example:

```
switch>
```

Indicates the operator command context.

```
switch#
```

Indicates the manager command context.

```
switch(CONTEXT-NAME)#
```

Indicates the configuration context for a feature. For example:

```
switch(config-if) #
```

Identifies the `interface` context.

Variable information in CLI prompts

In certain configuration contexts, the prompt may include variable information. For example, when in the VLAN configuration context, a VLAN number appears in the prompt:

```
switch(config-vlan-100) #
```

When referring to this context, this document uses the syntax:

```
switch(config-vlan-<VLAN-ID>) #
```

Where `<VLAN-ID>` is a variable representing the VLAN number.

Identifying switch ports and interfaces

Physical ports on the switch and their corresponding logical software interfaces are identified using the format:

member/slot/port

On the 6200 Switch Series

- *member*: Member number of the switch in a Virtual Switching Framework (VSF) stack. Range: 1 to 8. The primary switch is always member 1. If the switch is not a member of a VSF stack, then member is 1.
- *slot*: Line module number. Always 1.
- *port*: Physical number of a port on a line module.

For example, the logical interface 1/1/4 in software is associated with physical port 4 in slot 1 on member 1.

On the 6300 Switch Series

- *member*: Member number of the switch in a Virtual Switching Framework (VSF) stack. Range: 1 to 10. The primary switch is always member 1. If the switch is not a member of a VSF stack, then member is 1.
- *slot*: Line module number. Always 1.
- *port*: Physical number of a port on a line module.

For example, the logical interface 1/1/4 in software is associated with physical port 4 in slot 1 on member 1.

On the 6400 Switch Series

- *member*: Always 1. VSF is not supported on this switch.
- *slot*: Specifies physical location of a module in the switch chassis.
 - Management modules are on the front of the switch in slots 1/1 and 1/2.
 - Line modules are on the front of the switch starting in slot 1/3.
- *port*: Physical number of a port on a line module.

For example, the logical interface 1/3/4 in software is associated with physical port 4 in slot 3 on member 1.

Identifying modular switch components

- Power supplies are on the front of the switch behind the bezel above the management modules. Power supplies are labeled in software in the format: *member/power supply*.
 - *member*: 1.
 - *power supply*: 1 to 4.
- Fans are on the rear of the switch and are labeled in software as: *member/tray/fan*.
 - *member*: 1.
 - *tray*: 1 to 4.
 - *fan*: 1 to 4.
- Fabric modules are not labeled on the switch but are labeled in software in the format: *member/module*:

- *member*: 1.
- *member*: 1 or 2.
- The display module on the rear of the switch is not labeled with a member or slot number.

Quality of Service (QoS) enables network administrators to customize how different types of traffic are serviced on a network, taking into account the unique characteristics of each traffic type and its importance within an organization's infrastructure. QoS ensures uniform and efficient traffic handling, keeping the most important traffic moving at an acceptable speed, regardless of current bandwidth usage. It also provides methods for administrators to control the priority settings of inbound traffic arriving at each network device.

End-to-end QoS behavior

The QoS settings on each network device must be aligned to achieve the desired end-to-end QoS behavior for a network. Three service types can be used to categorize and prioritize network traffic:

- Best Effort Service
- Ethernet Class of Service (CoS)
- Internet Differentiated Services (DiffServ)

For a network as a whole, it is best to select one service type to use as the primary end-to-end behavior, and then use the other two service types as needed.

Best effort service

This is the simplest service type. All traffic is treated equally in a first-come, first-served manner. If the traffic load is low in relation to the capacity of the network links, then there is no need for the administrative complexity and costs of maintaining a more complex end-to-end policy. This is sometimes called over-provisioning, as all link speeds are much higher than peak loads on the network.

Class of Service

Class of Service (CoS) is a method for classifying network traffic at layer 2 by marking 802.1Q VLAN Ethernet frames with one of eight service classes.

CoS	Traffic type	Example protocols
7	Network Control	STP, PVST
6	Internetwork Control	BGP, OSPF, PIM
5	Voice (<10ms latency)	VoIP(UDP)
4	Video (<100ms latency)	RTP
3	Critical Applications	SQL RPC, SNMP
2	Excellent Effort	NFS, SMB
0	Best Effort	HTTP, TELNET
1	Background	SMTP, IMAP

CoS 1 is deliberately set as the lowest CoS. This enables a traffic service level below the default (best effort) traffic level to be specified.

The 3-bit Priority Code Point (PCP) field within the 16-bit Ethernet VLAN tag is used to mark the CoS.




Differentiated services

Differentiated services (DiffServ) is a method for classifying network traffic at layer 3 by marking packets with one of 64 different service classes. Services classes are identified by the Differentiated services Code Point (DSCP) value. Some common DSCP values are:

DSCP	Name	Service class	RFC
56	CS6	Network Control	2474
46	EF	Telephony	3246
40	CS5	Signaling	2474
34, 36, 38	AF41, AF42, AF43	Multimedia Conferencing	2597
32	CS4	Real-Time Interactive	2474
26, 28, 30	AF31, AF32, AF33	Multimedia Streaming	2597
24	CS3	Broadcast Video	2474
18, 20, 22	AF21, AF22, AF23	Low-Latency Data	2597
16	CS2	OAM	2474
00	CS0,BE,DF	Best Effort	2474
10, 12, 14	AF11, AF12, AF13	Bulk Data	2597
08	CS1	Low-Priority Data	3662

DSCP CS1 (08) CoS 1 is deliberately set as the lowest priority. This enables a traffic service level below the standard (best effort or default forwarding) level to be specified.

The DSCP value is carried within the IPv4 DSCP field or the upper 6-bits of the 8-bit IPv6 Traffic Class (TC) field.




QoS on the switch

There are five key stages a packet passes through when traversing a switch: ingress, prioritization, destination determination, egress queuing, and transmission. The following table provides an overview of each stage, and lists the commands that can be used to configure QoS settings.

Ingress	Prioritization	Destination Determination	Queueing	Transmission
Packets arrive at a switch interface (port).	A local priority value is assigned to the packet. This is used internally by the switch to classify traffic.	The egress interface for each packet is determined to allow queuing and packets are sent to the appropriate interface. QoS does not affect this stage.	Packets are queued based on the selected destination interface and local priority to queue mapping.	A scheduler defines the order in which packets are selected from the queues and are transmitted on the egress interface.
rate-limit Rate limiting can control ingress traffic flow by packet type.	qos trust Determines which packet values are used to create local priority: CoS, DSCP, or none (in which case the default value of CoS 0 is used). qos cos-map Defines how CoS values are mapped to local priority values. qos dscp-map Defines how DSCP values are mapped to local priority values. qos cos Remarks CoS egress values. qos dscp Remarks DSCP egress value.		qos queue-profile Creates a queue profile which defines eight transmit queues. map queue Assigns a local priority value to a queue number. Packets with a matching local priority value are placed in the queue. name queue Assigns a name to a queue. apply qos Assigns a queue profile and schedule profile globally to all interfaces, or a schedule profile to a specific interface.	qos schedule-profile Sets the order that queues are selected to transmit packets. qos shape Limits egress bandwidth on an interface. strict queue Assigns the strict priority algorithm to a queue number in a schedule profile dwrr queue Assigns the DWRR algorithm to a queue number in schedule profile.

The following diagram shows how different packets might traverse a switch. It also shows how QoS configuration settings apply at each stage.






On the 6400 Switch Series, interface identification differs.

QoS trust

Traffic priorities for networks can be carried in VLAN tags, using the CoS Priority Code Point (PCP), or in IP packet headers, using the Differentiated Services Code Point (DSCP). Whether these priorities affect how traffic is serviced, depends on how QoS trust mode is configured on the switch. QoS trust mode specifies how the switch assigns local priority values to ingress packets. Trust mode can be set globally for all interfaces, or individually for each interface. By default, trust mode is set to `none`, meaning that any QoS information in the packet (CoS or DSCP) is ignored, and local priority values are assigned from the CoS map value for code point 0. An exception to this can be configured, allowing a QoS remark to be applied to CoS or DSCP values when trust mode is `none`.

When trust mode is set to CoS or DSCP, the switch translates the QoS settings in VLAN tags (for CoS), or the DS field in an IP header (for DSCP), to local priority values on the switch. Translation is controlled by the CoS map or DSCP map tables.

For example:



Dynamic QoS trust mode

The device profile feature can dynamically set the QoS trust mode on an interface based on the LLDP information exchanged with a link partner. The device profile's trust mode temporarily overrides the static trust mode configured for an interface. The override remains in place as long as that link partner is connected and its link state is **up**. Use command `show interface IFNAME qos` to view the current QoS trust mode for an interface.

Port rate limiting

Port rate limiting helps control undesirable traffic. Its purpose is to allow enough broadcast, multicast, and ICMP rate-limit traffic for the network to function properly, while preventing flooding and traffic storms.

A certain amount of each type of traffic is required for normal network operation. Broadcast packets may include ARP and DHCP traffic, for instance. Video streams, and certain types of network protocol packets, are multicasts. Configuring rate limits can help provide the balance between necessary and flooded traffic.

Queue profile

A queue profile defines the queues that are associated with an interface to control the transmission of packets. Each profile supports up to eight queues, numbered 0 to 7. The larger the queue number, the higher its priority during transmission scheduling. Packets are assigned to a queue based on their local priority value (0 to 7). A queue profile must map all eight local priority values to whatever queues are being used on the switch, and a schedule profile must specify the configuration for those same queues. A queue without a local priority value assigned to it is not used to store packets.

The switch is automatically provisioned with an initial queue profile named `factory-default` which assigns each local priority to the queue of the same number. To see the default queue profile, use the command `show qos queue-profile factory-default`:

```
switch# show qos queue-profile factory-default
queue_num local_priorities name
-----
0          0                  Scavenger_and_backup_data
1          1
2          2
3          3
4          4
5          5
6          6
7          7
```

More than one local priority value can be assigned to the same queue. For example,

Local Priority	Queue
0	0
1	1
2	2
3	3
4	4
5	5
6	5
7	5

Queues must be consecutively defined starting at queue number zero. For example, a four-queue profile with priority values defined for queues 0, 1, 2, 3 is valid, but a four-queue profile which defines priority values for queues 1, 3, 5, and 7 is not.

Valid

Local Priority	Queue
0	0
1	1
2	1
3	1
4	2
5	2
6	3
7	3

Invalid

Local Priority	Queue
0	0
1	1
2	1
3	3
4	3
5	5
6	5
7	7

Commonly used commands for working with QoS queues are as follows:

- `qos queue-profile`: Creates an empty queue-profile and enters the profile configuration context.
- `name queue`: Assigns a descriptive name to a queue.
- `map queue`: Assigns a local-priority to a queue.
- `apply qos queue-profile`: Applies a queue-profile globally to all interfaces.

Schedule profiles

A schedule profile determines the order in which queues are selected for transmission, and the amount of service available for each queue. A schedule profile must be configured on every interface at all times. A schedule profile can be applied globally to all interfaces, or only to specific interfaces.

A weighted schedule profile assigns relative servicing for each queue. The amount of service per weight is relative to the underlying hardware implementation, and to the weights assigned to the other non-empty queues. Strict scheduling can be used to service queues purely on the basis of highest priority first (at the risk of starving lower-priority queues during high stress periods). A combination of strict and weighted

scheduling offers more service to the highest priority queue when needed, while preserving scheduling between the remaining queues, thus decreasing the risk of starvation.

The switch is automatically provisioned with a schedule profile named **factory-default**, which assigns DWRR to all queues with a weight of 1. Use the command `show schedule-profile factory-default` to view the default schedule profile. (Do not use `show running-configuration`, as it only displays changes from the initial settings.) The 6400 Switch Series only has four queues.

```
switch# show qos schedule-profile factory-default
queue_num algorithm weight
-----
0          dwrr        1
1          dwrr        1
2          dwrr        1
3          dwrr        1
4          dwrr        1
5          dwrr        1
6          dwrr        1
7          dwrr        1
```

Egress queue shaping

Egress queue shaping limits the amount of traffic transmitted per strict output queue. The buffer associated with each egress queue stores excess traffic to absorb bursts and smooths the output rate. For example, an administrator might limit strict-priority queue traffic to prevent low-priority queue starvation in the event that a device inappropriately sends too many higher-priority packets.

Egress queue shaping can be configured on an Ethernet port or on a link aggregation group (LAG). To configure egress queue shaping, define a schedule profile with the strict priority algorithm assigned to each queue.

Egress port shaping

Egress port shaping limits the amount of aggregate traffic transmitted through a port.

To be effective, the egress port-shaping rate must be less than the port's line rate. By default, the egress port-shaping rate is the same as the line-rate of the port. Buffers associated with each port store excess traffic.

When both egress port-shaping and egress queue-shaping are configured on the same interface, the switch respects the minimum of both configurations.

Virtual output queues

6200 and 6300 Switch Series

The 6200 and 6300 Switch Series contain 8 MB of packet buffer memory. Approximately 5 MB are used for buffering traffic arriving on network ports. Each port can buffer at most 1 MB of packets divided equally across all the active queues in a queue profile. For example, using the default eight-queue profile, each virtual output queue (VOQ) can buffer a maximum of 128 KB of packets.

Replicated traffic (L2 broadcasts, L2 multicasts, IP multicasts, mirrored packets, and unknown-unicasts) requires an extra step. There are eight additional replication queues initially used for replicated traffic. The switch services the replication queues strictly at the rate of 100 G. For ports on the same switch, pointers to the same packet are placed in the VOQ for each destination port. For ports on remote VSF member switches, only one copy of the packet data is moved across the VSF link to the destination replication VOQ to optimize link bandwidth. From there, the destination switch puts pointers to the same packet on the VOQ for each destination port.

6400 Switch Series

The 6400 Switch Series use a VOQ architecture where most packet buffering occurs on the ingress line module. Only shallow buffers are used on the egress line module. This optimizes usage of fabric bandwidth. Traffic destined for one port (unicast) uses different buffering and scheduling than traffic destined for multiple ports (replication).

Each line module contains eight VOQs for every destination port in the chassis, and eight egress physical queues (EPQs) for the ports local to the line module, regardless of the number of queues in the queue profile. The 6400 Switch Series R0X38-R0X45 line modules contain 8 MB of packet buffer memory. Approximately 5 MB is used for buffering traffic arriving on network ports. The remainder is used for buffering traffic arriving across the fabric. On each line module, at most 1 MB of buffering is available for every destination port in the chassis. This 1 MB buffer is equally divided across all the active queues in the queue profile.

For example, using the default eight-queue profile, each VOQ can buffer at most 128 KB of packets. Packets wait in VOQs on the arrival line module until the destination port scheduler selects them to be moved to one of its EPQs. If the destination port is on another line module, the packet is moved across the fabric. If the destination port EPQ is on the same line module, only a pointer to the packet is moved from the VOQ to the EPQ.

The schedule profile determines the order the destination port scheduler uses for VOQ servicing. The scheduler moves only enough packets to EPQs as are needed to keep the destination port transmitting at the configured rate.

Replicated traffic (L2 broadcasts, L2 multicasts, IP multicasts, mirrored packets, and unknown-unicasts) requires an extra step. There are eight additional internal replication VOQs per line module initially used for replicated traffic. The line module services these replication VOQs strictly at a rate of 100 G. For ports on the same line module, pointers to the same packet are placed in the VOQ for each destination port. For ports on remote line modules, only one copy of the packet data is moved across the fabric to the destination replication VOQ, to optimize fabric bandwidth. From there, the destination line module will put pointers to the same packet on the VOQ for each destination port.

Terms

Class

For networking, a set of packets sharing a common characteristic. For example, all IPv4 packets.

Code point

The name of a packet header field, or the value carried within a packet header field:

Example 1: Priority code point (PCP) is the name of a field in the IEEE 802.1Q VLAN tag.

Example 2: Differentiated services code point (DSCP) is the name of a field carried within the DS field of an IP packet header.

Color

A metadata label associated with each packet within the switch. It has two values: green (0) or yellow (1). When packets encounter congestion for a resource (queue), the switch uses packet color to distinguish which packets must be dropped, and is mostly used for packets marked with Assured Forwarding (AF) DSCP values.

Not supported in this release.

Class of service (CoS)

A 3-bit value used to mark packets with one of eight classes (levels of priority). It is carried within the priority code point (PCP) field of the IEEE 802.1Q VLAN tag.

Differentiated services code point (DSCP)

A 6-bit value used to mark packets for different per-hop behavior as originally defined by IETF RFC 2474. It is carried within the differentiated services (DS) field of the IPv4 or IPv6 header.

Local priority

A meta-data label associated with a packet within the switch which is used to classify packets for different treatment (such as queue assignment). Eight local priorities are defined on the switch, numbered from 0 to 7. A queue profile must map all eight local priorities to whatever queues are in use on the switch, and a schedule profile must specify the configuration for these same queues.

Metadata

Information labels associated with each packet in the switch, separate from the packet headers and data. These labels are used by the switch in its handling of the packet. For example: arrival port, egress port, VLAN membership, and local priority.

Priority code point (PCP)

The name of a 3-bit field in the IEEE 802.1Q VLAN tag. It carries the CoS value to mark a packet with one of eight classes (priority levels).

Quality of service (QoS)

General term used when describing or measuring performance. For networking, it means how different classes of packets are treated when traversing a network or device.

Traffic class (TC)

General term for a set of packets sharing a common characteristic. It used to be the name of an 8-bit field in the IPv6 header originally defined by IETF RFC 2460. This field name was changed to differentiated services by IETF RFC 2474.

Type of service (ToS)

General term when there are different levels of treatment (fare class). It used to be the name of an 8-bit field in the IPv4 header originally defined by IETF RFC 791. This field name was changed to differentiated services by IETF RFC 2474.

Configuring QoS

Procedure

1. Configure how local priority values are assigned to ingress packets with the commands `qos cos-map`, `qos dscp-map`, and `qos trust`.
2. Optionally, add a rate limit for ingress traffic on one or more interfaces with the command `rate-limit`.
3. If you do not want to use the default QoS queue profile to map local priority to queue, create one or more custom queue profiles with the command `qos queue-profile`. For each queue in a custom queue profile:
 - a. Assign a local priority value with the command `map queue`.
 - b. Optionally, define a descriptive name with the command `name queue`. All local priorities (0 to 7) must be mapped to a queue, and the queues selected for use must be in contiguous order starting at 0.
4. If you do not want to use the default QoS schedule profile to determine the order in which queues are selected to transmit a packet, create one or more custom schedule profiles with the command `qos schedule-profile`. For each queue in a custom schedule queue profile, define scheduling priority with the commands `strict queue` and `dwrr queue`.
5. Optionally for strict queues, configure egress queue shaping to limit egress bandwidth on an interface to a value that is less than its line rate. Use the `max-bandwidth` parameter of the `strict queue` command.
6. Activate QoS settings with the command `apply qos`. This command lets you apply a queue profile and schedule profile globally to all interfaces, or a schedule profile override to individual interfaces.
7. View QoS configuration settings with the provided `show` commands.

Examples

This example creates the following configuration:

- Configures CoS to be used to assign local priority to ingress packets.
- Modifies the default CoS map to assign CoS 1 to local priority 1.
- Creates a queue profile named Q1 and assigns local priorities as follows:

Queue	Local Priority
0	0
1	1
1	2

Table Continued

Queue	Local Priority
2	3
3	4
4	5
5	6
5	7

- Creates a schedule profile named **s1** and assigns DWRR to all queues in the schedule profile with the following weights:

Queue	Weight
0	5
1	10
2	15
3	20
4	25
5	50

- Applies **q1** and **s1** to all interfaces that do not have a QoS override applied.

```
switch(config)# qos trust cos
switch(config)# qos cos-map 1 local-priority 1
switch(config)# qos queue-profile Q1
switch(config)# map queue 0 local-priority 0
switch(config)# map queue 1 local-priority 1
switch(config)# map queue 1 local-priority 2
switch(config)# map queue 2 local-priority 3
switch(config)# map queue 3 local-priority 4
switch(config)# map queue 4 local-priority 5
switch(config)# map queue 5 local-priority 6
switch(config)# map queue 5 local-priority 7
switch(config)# qos schedule-profile S1
switch(config)# dwrr queue 0 weight 5
switch(config)# dwrr queue 1 weight 10
switch(config)# dwrr queue 2 weight 15
switch(config)# dwrr queue 3 weight 20
switch(config)# dwrr queue 4 weight 25
switch(config)# dwrr queue 5 weight 50
switch(config)# apply qos queue-profile Q1 schedule-profile S1
```

Configuring expedited forwarding for VoIP traffic

Voice over IP (VoIP) traffic is delay and jitter sensitive. Therefore, for optimum transmission of VoIP traffic, dwell time in network devices must be kept to a minimum, and all network devices in the data path must have identical per-hop behaviors. To configure a dedicated queue on the switch to handle VoIP traffic with priority service before all other queues, follow these steps.

Prerequisites

This scenario assumes that VoIP packets are uniquely identified using DiffServ code point 46, Expedited Forwarding (EF).

Procedure

1. Map DSCP EF packets exclusively to local priority 5. The default DSCP map has eight code points (40 through 47), that are mapped to local priority 5. To reserve local priority 5 for VoIP traffic, the other code points must be reassigned. In this scenario, local priority 6 is used for all reassignments, including for code point 40, Call Signaling protocol (CS5).

```
switch# config
switch(config)# qos dscp-map 40 local-priority 6 name CS5
switch(config)# qos dscp-map 41 local-priority 6
switch(config)# qos dscp-map 42 local-priority 6
switch(config)# qos dscp-map 43 local-priority 6
switch(config)# qos dscp-map 44 local-priority 6
switch(config)# qos dscp-map 45 local-priority 6
switch(config)# qos dscp-map 47 local-priority 6
```

2. Queue 7 is the highest priority queue, so for best throughput, create a queue profile that maps local priority 5 to queue 7.

```
switch(config)# qos queue-profile ef_priority
switch(config-queue)# name queue 7 Voice_Priority_Queue
switch(config-queue)# map queue 7 local-priority 5
switch(config-queue)# map queue 6 local-priority 7
switch(config-queue)# map queue 5 local-priority 6
switch(config-queue)# map queue 4 local-priority 4
switch(config-queue)# map queue 3 local-priority 3
switch(config-queue)# map queue 2 local-priority 2
switch(config-queue)# map queue 1 local-priority 1
switch(config-queue)# map queue 0 local-priority 0
switch(config-queue)# exit
switch(config)#
```

3. Create a schedule profile that services queue 7 using strict priority (SP), and the remaining queues with DWRR. This scenario gives all DWRR queues equal weight.

```
switch(config)# qos schedule-profile voip
switch(config-schedule)# strict queue 7
switch(config-schedule)# dwrr queue 6 weight 1
switch(config-schedule)# dwrr queue 5 weight 1
switch(config-schedule)# dwrr queue 4 weight 1
switch(config-schedule)# dwrr queue 3 weight 1
switch(config-schedule)# dwrr queue 2 weight 1
switch(config-schedule)# dwrr queue 1 weight 1
switch(config-schedule)# dwrr queue 0 weight 1
switch(config-schedule)# exit
switch(config)#
```

4. Apply the profiles to all interfaces.


```
switch(config)# apply qos queue-profile ef_priority schedule-profile voip
```

5. Configure DSCP trust mode on all ports

```
switch(config)# qos trust dscp
```

Configuring rate limiting

This scenario illustrates how to use rate limiting to manage the traffic from various devices connected to a switch. The physical topology of the network looks like this:



A certain amount of broadcast traffic is necessary to maintain healthy network operation, particularly from routers and across service boundaries. In this scenario, both the service cloud and the router connections limit this traffic to 1 Gbps. The server has a smaller limit, as it does not require as much network protocol traffic as the service cloud and router.

A multicast server needs to be able to stream multicast traffic to clients, so a multicast rate limit may not be helpful. A computer, however, should not be generating large amounts of multicast traffic (it may be receiving streams, but typically not sending them). In this example, the computer is configured with a multicast rate limit to prevent malicious traffic from taking up network bandwidth.

Finally, while the service cloud and router may need to send traffic for unknown unicast addresses to resolve address forwarding, the server and computer should send very little of this type of traffic. Rate limiting unknown unicast traffic on those two devices enforces that.

On the 6400 Switch Series, interface identification differs.

Procedure

1. Configure broadcast and multicast rate limiting for the service cloud connection.

```
switch# config
switch(config)# interface 1/1/1
switch(config-if)# rate-limit broadcast 500 pps
switch(config-if)# rate-limit multicast 500 pps
switch(config-if)# exit
```

2. Configure broadcast rate limiting for the router connection.

```
switch(config-if)# interface 1/1/2
switch(config-if)# rate-limit broadcast 500 pps
switch(config-if)# exit
```

3. Configure broadcast rate limiting for the server connection.

```
switch(config-if)# interface 1/1/5
switch(config-if)# rate-limit broadcast 100 pps
switch(config-if)# exit
```

4. Configure broadcast, and multicast rate limiting for the computer connection.

```
switch(config-if)# interface 1/1/10
switch(config-if)# rate-limit broadcast 50 pps
switch(config-if)# rate-limit multicast 50 pps
```

Configuring egress queue shaping

This example shows how to apply egress queue shaping to an interface. First, a schedule profile is created that has per-queue bandwidth limits set on all queues with `strict` as the scheduling algorithm. Next, this profile is applied to an interface or LAG.

On the 6400 Switch Series, interface identification differs.

The following example creates a schedule profile named `EQSEExample`, which services all seven queues using `strict` priority. This profile configures queues 1, 4, and 7 with a bandwidth limit of 10 Gbps, 20 Gbps, and 30 Gbps respectively. The profile is then applied to interface 1/1/1.

```
switch(config)# qos schedule-profile EQSEExample
switch(config-schedule)# strict queue 0
switch(config-schedule)# strict queue 1 max-bandwidth 10000000
switch(config-schedule)# strict queue 2
switch(config-schedule)# strict queue 3
switch(config-schedule)# strict queue 4 max-bandwidth 20000000
switch(config-schedule)# strict queue 5
switch(config-schedule)# strict queue 6
switch(config-schedule)# strict queue 7 max-bandwidth 30000000
switch(config-schedule)# exit
switch(config)# interface 1/1/1
switch(config-if)# apply qos schedule-profile EQSEExample
```

Configuring egress port shaping

This example shows how to apply egress port shaping to an interface to limit the rate of egress traffic. Egress port shaping is configured by specifying the desired bandwidth rate in kilobits per second (kbps). To be effective, the egress rate must be less than the line rate of the egress interface. If the configured egress rate exceeds the interface's line rate, then egress shaping has no effect.

The configured egress rate on a specific interface can be found by using the `show interface` and `show interface qos` commands.

On the 6400 Switch Series, interface identification differs.

The following example configures an egress rate of 100 Mbps .

```
switch(config)# interface 1/1/1
switch(config-if)# qos shape 100000
```

In the next example, both egress port shaping and egress queue shaping are configured on the same interface.

The example creates a schedule profile named `EQSEExample` with `strict` priority for all seven queues. Queue 7 is configured with a bandwidth limit of 300 Mbps. The profile is then applied to interface 1/1/1 with egress port shaping of 400 Mbps . As egress queue shaping and egress port shaping are both configured on port 1/1/1, egress queue shaping is subject to the lower port or queue shape rate. The effective bandwidth for the traffic egressing on queue 7 will be 300 Mbps .

```
switch(config)# qos schedule-profile EQSEExample
switch(config-schedule)# strict queue 0
switch(config-schedule)# strict queue 1
switch(config-schedule)# strict queue 2
switch(config-schedule)# strict queue 3
switch(config-schedule)# strict queue 4
switch(config-schedule)# strict queue 5
switch(config-schedule)# strict queue 6
switch(config-schedule)# strict queue 7 max-bandwidth 300000
switch(config-schedule)# exit
```

```
switch(config)# interface 1/1/1
switch(config-if)# apply qos schedule-profile EQSExample
switch(config-if)# qos shape 400000
```

Supporting Ethernet 802.1D Class of Service

IEEE 802.1Q is the most current Ethernet standard for Class of Service (CoS). It superseded an earlier standard, 802.1D, in 2005. IEEE 802.1Q slightly changed the ordering of the classes of service from its predecessor IEEE 802.1D for CoS 2 and CoS 0:

CoS 802.1Q	CoS 802.1D
7 Network Control	7 Network Control
6 Internetwork Control	6 Voice (<10ms latency)
5 Voice (<10ms latency)	5 Video (<100ms latency)
4 Video (<100ms latency)	4 Controlled Load
3 Critical Applications	3 Excellent Effort
2 Excellent Effort	0 Best Effort
0 Best Effort	2 Spare
1 Background	1 Background

In 802.1D, both CoS 2 and CoS 1 are below CoS 0 (Best Effort).

When a switch is installed in a network of devices following 802.1D Class of Service, the QoS Cos map must be re-configured to follow the 802.1D standard by swapping the assignments of CoS 0 and 2:

```
switch# config
switch(config)# qos cos-map 0 local-priority 2 color green name Best_Effort
switch(config)# qos cos-map 2 local-priority 1 color green name Spare
switch(config)# exit
switch# show qos cos-map
code_point local_priority color name
-----
0          2          green Best_Effort
1          0          green Background
2          1          green Spare
3          3          green Critical_Applications
4          4          green Video
5          5          green Voice
6          6          green Internetwork_Control
7          7          green Network_Control
```

Monitoring queue operation

Use the show interface queues command to display the traffic transmitted per queue, and the number of packets dropped due to the queue being full. (Tx Bytes is available on the 6200 and 6300 Switch Series.) For example:

```
switch# show interface 1/1/1 queues
Interface 1/1/1 is (Administratively down)
Admin state is down
State information: admin_down
      Tx Packets          Tx Bytes          Tx Drops
Q0          100             8000             0
Q1       1234567       12345678908             5
```


Q2	0	0	0
Q3	0	0	0
Q4	0	0	0
Q5	0	0	0
Q6	0	0	0
Q7	0	0	0

- **Tx Bytes:** Total bytes transmitted (available on the 6200 and 6300 Switch Series). The byte count may include packet headers and internal metadata that are removed before the packet is transmitted. Packet headers added when the packet is transmitted may not be included.
- **Tx Packets:** Total packets transmitted.
- **Tx Drops:** The number of packets dropped by a queue before it was sent. When traffic cannot be forwarded out an egress interface, it backs up at ingress. The more servicing assigned to a queue by a schedule profile, the less likely traffic destined for that queue will back up and be dropped. Tx Drops shows the sum of packets that were dropped across all line modules (due to insufficient capacity) by the ingress Virtual Output Queues (VOQs) destined for the egress port. Applicable to the 6400 Switch Series: Because the counts are read separately from each line module, the sum is not an instantaneous snapshot.

apply qos

Syntax

```
apply qos [queue-profile <QUEUE-NAME>] schedule-profile <SCHEDULE-NAME>
```

```
no apply qos schedule-profile <SCHEDULE-NAME>
```

Description

Applies a queue profile and schedule profile globally to all Ethernet and LAG interfaces on the switch, or applies a schedule profile to a specific interface. When applied globally, the specified schedule profile is configured only on Ethernet interfaces and LAGs that do not already have their own schedule profile.

The same profile can be applied both globally and locally to an interface. This guarantees that an interface always uses the specified profile, even if the global profile is changed.

The `no` form of this command removes the specified schedule profile from an interface and the interface uses the global schedule profile. This is the only way to remove a schedule profile override from the interface.



NOTE: Interfaces may shut down briefly during reconfiguration.

Command context

`config`

`config-if`

`config-lag-if`

Parameters

queue-profile <QUEUE-NAME>

Specifies the name of the queue profile to apply. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-). This parameter is not supported in the `config-if` context.

schedule-profile <SCHEDULE-NAME>

Specifies the name of the schedule profile to apply. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Usage

- For a queue profile to be complete and ready to be applied, all eight local priorities must be mapped to a queue.
- For a schedule profile to be complete and ready to be applied, it must define all queues specified in the queue profile. All queues must use the same algorithm, except for the highest numbered queue, which can be **strict**.
- Both the queue profile and the schedule profile must specify the same number of queues.
- Schedule profiles can be modified while applied, but only in ways where a single command will not result in the profile becoming invalid. For example, queue 7 can have the algorithm changed, and weighted queues can have their weights changed.
- Queues must be consecutively defined starting at queue number zero. For example, a four-queue profile with priority values defined for queues 0, 1, 2, 3 is valid, but a four-queue profile which defines priority values for queues 1, 3, 5, and 7 is not.
- The following commands illustrate a valid configuration, where every local priority value is assigned to a queue, and all assigned queues are defined.

```
qos cos-map 1 local-priority 1
qos queue-profile Q1
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 2 local-priority 2
map queue 3 local-priority 3
map queue 4 local-priority 4
map queue 5 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
qos schedule-profile S1
dwrr queue 0 weight 5
dwrr queue 1 weight 10
dwrr queue 2 weight 15
dwrr queue 3 weight 20
dwrr queue 4 weight 25
dwrr queue 5 weight 50
```

The following commands illustrate an invalid configuration, because local priority 2 is not assigned to a queue.

```
qos cos-map 1 local-priority 1
qos queue-profile Q1
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 3 local-priority 3
map queue 4 local-priority 4
map queue 5 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
qos schedule-profile S1
dwrr queue 0 weight 5
dwrr queue 1 weight 10
dwrr queue 3 weight 15
dwrr queue 4 weight 25
dwrr queue 5 weight 50
```

Example

Applying the QoS profile **myprofile** and the schedule profile **myschedule** to all interfaces that do not have an applied interface-specific schedule profile:

```
switch(config)# apply qos queue-profile myprofile schedule-profile myschedule
```

dwrr queue

Syntax

```
dwrr queue <QUEUE-NUMBER> weight <WEIGHT>
```

```
no dwrr queue <QUEUE-NUMBER>
```

Description

Assigns the deficit weighted round robin (DWRR) algorithm and its weight to a queue in a schedule profile. DWRR allocates available bandwidth among all non-empty queues in relation to the queue weights.

The `no` form of this command removes the DWRR algorithm from a queue in a schedule profile.

Command context

config-schedule

Parameters

<QUEUE-NUMBER>

Specifies the queue number. Range: 0 to 7.

weight <WEIGHT>

Specifies the scheduling weight. Range: 1 to 1023.

Authority

Administrators or local user group members with execution rights for this command.

Examples

Assigning DWRR with a weight of 17 to queue 2 in the schedule profile `myschedule`:

```
switch(config)# qos schedule-profile myschedule  
switch(config-schedule)# dwrr queue 2 weight 17
```

Deleting DWRR for queue 2 from the schedule profile `myschedule`:

```
switch(config)# qos schedule-profile myschedule  
switch(config-schedule)# no dwrr queue 2
```

map queue

Syntax

```
map queue <QUEUE-NUMBER> local-priority <PRIORITY-NUMBER>
```

```
no map queue <QUEUE-NUMBER> [local-priority <PRIORITY-NUMBER>]
```

Description

Assigns a local priority to a queue in a queue profile. By default, the larger the queue number the higher its priority. A queue without a local priority value assigned to it is not used to store packets. The same queue can be assigned multiple local priorities.

Queues must be consecutively defined starting at queue number zero. For example, a four-queue profile with priority values defined for queues 0, 1, 2, 3 is valid, but a four-queue profile which defines priority values for queues 1, 3, 5, and 7 is not.

The `no` form of this command removes the specified local priority from a specific queue. If no local priority number is specified, then all local priorities are removed from the queue.

Command context

`config-queue`

Parameters

<QUEUE-NUMBER>

Specifies the queue number. Range: 0 to 7.

<PRIORITY-NUMBER>

Specifies the queue priority. Range: 0 to 7, where 0 is the lowest priority and 7 is the highest.

Authority

Administrators or local user group members with execution rights for this command.

Usage

For a queue profile to be complete and ready to be applied, all eight local priorities must be mapped to a queue.

Queues must be consecutively defined starting at queue number zero. For example, a four-queue profile with priority values defined for queues 0, 1, 2, 3 is valid, but a four-queue profile which defines priority values for queues 1, 3, 5, and 7 is not.

The following commands illustrate an invalid configuration, where every local priority value is assigned to a queue.

```
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 1 local-priority 2
map queue 3 local-priority 3
map queue 4 local-priority 4
map queue 5 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
```

The following commands illustrate an invalid configuration, because local priority 2 is not assigned to a queue.

```
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 2 local-priority 3
map queue 3 local-priority 4
map queue 4 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
```

Examples

Assigning priority 7 to queue 7 in profile `myprofile`:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# map queue 7 local-priority 7
```

Removing priority 7 from queue 7 in profile **myprofile**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# no map queue 7 local-priority 7
```

name queue

Syntax

```
name queue <QUEUE-NUMBER> <DESCRIPTION>
```

```
no name queue <QUEUE-NUMBER>
```

Description

Assigns a description to a queue in a queue profile. This is for identification purposes and has no effect on configuration.

The **no** form of this command removes the description associated with a queue.

Command context

config-queue

Parameters

<QUEUE-NUMBER>

Specifies the queue number. Range: 0 to 7.

<DESCRIPTION>

Specifies a queue description for identification purposes. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Examples

Assigning the description **priority-traffic** to queue 7:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# name queue 7 priority-traffic
```

Removing the description from queue 7:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# no name queue 7
```

qos cos

Syntax

```
qos cos <CODE-POINT>
```

```
no qos cos
```

Description

Configures a CoS PCP remark for an Ethernet or LAG interface. Packets that ingress on the interface are remarked at egress using the configured CoS PCP value.

The remark only occurs when QoS trust mode on the interface is set to `none`.

If QoS trust mode is not set to `none`, then the remark is ignored, and the following commands will show the CoS remark status as `ignored` (incompatible Port Access Trust configuration) or `not applied` (incompatible QoS global/port Trust configuration):

- `show running-configuration`
- `show interface <PORT-NUM>`
- `show interface <PORT-NUM> qos`

The `no` form of this command removes a CoS remark on an interface.

Command context

`config-if`

Parameters

<CODE-POINT>

Specifies an 802.1 VLAN priority CoS value. Range: 0 to 7.

Authority

Administrators or local user group members with execution rights for this command.

Examples

On the 6400 Switch Series, interface identification differs.

Configuring a CoS remark of 3 on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# qos trust none
switch(config-if)# qos cos 3
```

Deleting a CoS remark of 3 on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# no qos cos
```

qos cos-map

Syntax

```
qos cos-map <CODE-POINT> local-priority <PRIORITY-NUMBER> [color <COLOR>] [name <DESCRIPTION>]
```

```
no qos cos-map <CODE-POINT>
```

Description

Defines the local priority assigned to incoming packets for a specific 802.1 VLAN priority code point (CoS) value. The CoS map values are used to mark incoming packets when QoS trust mode is set to `cos`. In **trust none** mode, CoS map entry 0 is used to set the port default local priority and color.

To see the default CoS map settings, use the following command:

```
switch# show qos cos-map default
code_point local_priority color name
-----
0          1              green Best_Effort
1          0              green Background
2          2              green Excellent_Effort
3          3              green Critical_Applications
4          4              green Video
5          5              green Voice
6          6              green Internetwork_Control
7          7              green Network_Control
```

The `no` form of this command restores the assignments for a CoS map value to the default setting.

Command context

config

Parameters

<CODE-POINT>

Specifies an 802.1 VLAN priority CoS value. Range: 0 to 7. Default 0.

local-priority <PRIORITY-NUMBER>

Specifies a local priority value to associate with the `CODE-POINT` value. Range: 0 to 7. Default: 0.

color <COLOR>

Reserved for future use.

name <DESCRIPTION>

Specifies a description for the CoS setting. The name is for identification only, and has no effect on queue configuration. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Examples

Mapping CoS value 1 to a local priority of 2:

```
switch(config)# qos cos-map 1 local-priority 2
```

Mapping CoS value 1 to the default local priority value:

```
switch(config)# no qos cos-map 1
```

qos dscp

Syntax

```
qos dscp <CODE-POINT>
```

```
no qos dscp
```

Description

Configures a differentiated services code point (DSCP) remark for an Ethernet or LAG interface. IPV4 and IPV6 packets that ingress on the interface are remarked at egress using the configured DSCP value.

The remark only occurs when QoS trust mode on the interface is set to `none`. If a DSCP remark is configured and then trust mode is subsequently set to `cos` or `dscp`, then the DSCP remark is ignored.

The following commands will show the remark status as `ignored` (incompatible Port Access Trust configuration) or `not applied` (incompatible QoS global/port Trust configuration):

- `show running-configuration`
- `show interface <PORT-NUM>`
- `show interface <PORT-NUM> qos`

The `no` form of this command removes a CoS remark on an interface.

Command context

`config-if`

Parameters

<CODE-POINT>

Specifies an IP differentiated services code point value. Range: 0 to 63.

Authority

Administrators or local user group members with execution rights for this command.

Usage

For all arriving IPv4 or IPv6 packets:

- Local priority and color are assigned from the CoS map entry 0.
- The packet's DSCP value in the IPv4 or IPv6 DS header field is remarked with the DSCP remark value.

For all arriving non-IP packets:

- Local priority and color are assigned from the CoS map entry 0.
- The CoS of all arriving tagged non-IP packets is unchanged.

Examples

On the 6400 Switch Series, interface identification differs.

Configuring a DSCP remark of **43** on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# qos trust none
switch(config-if)# qos dscp 43
```

Deleting a DSCP remark of **43** on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# no dscp 43
```

qos dscp-map

Syntax

```
qos dscp-map <CODE-POINT> local-priority <PRIORITY-NUMBER> [color <COLOR>] [cos <COS-CODE-POINT>] [name <DESCRIPTION>]
no qos dscp-map <CODE-POINT>
```

Description

Defines the local priority assigned to incoming packets for a specific IP differentiated services code point (DSCP) value. The DSCP map values are used to prioritize incoming packets when QoS trust mode is set to **dscp**.

To see the default DSCP map settings, use the following command:

```
switch# show qos dscp-map default
code_point local_priority cos color name
-----
0          1              green CS0
1          1              green
2          1              green
3          1              green
4          1              green
5          1              green
...
45         5              green
46         5              green EF
47         5              green
48         6              green CS6
...
61         7              green
62         7              green
63         7              green
```

The **no** form of this command restores the assignments for a code point to the default setting.

Command context

config

Parameters

<CODE-POINT>

Specifies an IP differentiated services code point. Range: 0 to 63. Default: 0.

local-priority <PRIORITY-NUMBER>

Specifies a local priority value to associate with the **CODE-POINT** value. Range: 0 to 7. Default: 0.

color <COLOR>

Reserved for future use.

cos <COS-CODE-POINT>

Specifies an 802.1 VLAN priority CoS remark value. Range: 0 to 7. Default 0

name <DESCRIPTION>

Specifies a description for the DSCP setting. The name is used for identification only, and has no effect on queue configuration. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Examples

Setting code point 1 to a local priority of 2 and a CoS of 0:

```
switch(config)# qos dscp-map 1 local-priority 2 cos 0
```

Setting code point 1 to the default value:

```
switch(config)# no qos dscp-map 1
```

qos queue-profile

Syntax

```
qos queue-profile <NAME>
```

```
no qos queue-profile <NAME>
```

Description

Creates a new QoS queue profile and switches to the `config-queue` context for the profile. Or, if the specified QoS queue profile exists, this command switches to the `config-queue` context for the profile.

A queue profile maps queues to local-priority values.

Each profile supports up to eight queues, starting at 0 and incrementing sequentially to the number of desired queues.

A queue profile named **factory-default** is defined by default and is automatically applied to all interfaces. It cannot be edited or deleted. Use the command `show qos queue-profile factory-default` to view this profile. Do not use `show running-configuration`, as it will only display changes from the initial values.

The `no` form of this command removes the specified QoS queue profile. Only profiles that are not currently applied can be removed.

Command context

```
config
```

Parameters

<NAME>

Specifies the name of the QoS queue profile to create or configure. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Examples

Creating the profile **myprofile**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)#
```

Deleting the profile **myprofile**:

```
switch(config)# no qos queue-profile myprofile
```

qos schedule-profile

Syntax

```
qos schedule-profile <NAME>
```

```
no qos schedule-profile <NAME>
```

Description

Creates a QoS schedule profile and switches to the `config-schedule` context for the profile. If the specified schedule profile exists, this command switches to the `config-schedule` context for the profile. The schedule profile determines the order in which queues are selected to transmit a packet, and the amount of service defined for each queue.

Queues in a schedule profile are numbered consecutively starting from zero. Queue zero is the lowest priority queue. The larger the queue number, the higher priority the queue has in scheduling algorithms.

A profile named **factory-default** is defined by default and applied to all interfaces. It cannot be edited or deleted. To see its settings, use the command:

```
switch# show qos schedule-profile factory-default
queue_num algorithm weight
-----
0          dwrr        1
1          dwrr        1
2          dwrr        1
3          dwrr        1
4          dwrr        1
5          dwrr        1
6          dwrr        1
7          dwrr        1
```

A profile named **strict** is predefined and cannot be edited or deleted. The strict profile services all queues of the queue profile to which it is applied, using the strict priority algorithm.

A schedule profile must be defined on all interfaces at all times.

There are two permitted configurations for a schedule profile:

1. All queues use the same scheduling algorithm (for example, DWRR).
2. The highest queue number uses strict priority, and all remaining (lower) queues use the same algorithm (for example, DWRR). This supports priority scheduling behavior necessary for the IETF RFC 3246 Expedited Forwarding specification (<https://tools.ietf.org/html/rfc3246>).

Only limited changes can be made to an applied schedule profile:

- The weight of a dwrr queue.
- The bandwidth of a strict queue.
- The algorithm of the highest numbered queue can be swapped between dwrr and strict, and vice versa.

Applicable to REST: Any other changes will result in an unusable schedule profile, and the switch will revert to the `factory-default` profile until the profile is corrected.

The `no` form of this command removes the specified QoS schedule profile when it is not applied. Only profiles that are not currently applied to an interface can be removed.

Command context

config

Parameters

<NAME>

Specifies the name of the QoS schedule profile to create or configure. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Authority

Administrators or local user group members with execution rights for this command.

Examples

Creating the schedule profile **myschedule**:

```
switch(config)# qos schedule-profile myschedule
switch(config-schedule)#
```

Deleting the schedule profile **myschedule**:

```
switch(config)# no qos schedule-profile myschedule
```

qos shape

Syntax

qos shape <RATE>

no qos shape

Description

Limits the egress bandwidth on an interface to a value that is lower than its line rate.

The **no** form of this command removes shaping from an interface.

Command context

config-if

Parameters

<RATE>

Specifies the maximum traffic rate in kbps. Range: 468 to 100000000.

Authority

Administrators or local user group members with execution rights for this command.

Examples

On the 6400 Switch Series, interface identification differs.

Configuring an egress port shaping rate of 400 Mbps on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# qos shape 400000
```

Deleting egress port shaping on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# no qos shape
```

qos trust

Syntax

```
qos trust {none | cos | dscp}
```

```
no qos trust
```

Description

Sets the trust mode. Trust mode determines whether CoS or DSCP values are used to assign local priority values to ingress packets. CoS values are taken from the CoS map, and DSCP values are taken from the DSCP map.

In the `config` context:

- This command sets the trust mode that is globally applied to all interfaces that do not have a trust mode configured.
- The `no` form of this command restores all interfaces that do not currently have a trust mode configured to the default setting (`none`).

In the `config-if` context:

- This command sets the trust mode override for a specific interface.
- The `no` form of this command clears a trust mode override. The interface then uses the global setting. This is the only way to remove a trust mode override.

Command context

```
config
```

```
config-if
```

Parameters

none

Ignore all packet headers. Ingress packets are assigned the priority value configured for CoS map value 0. Default.



NOTE: In the `config-if` context, the commands `qos cos` and `qos dscp` can be used to remark packets when `none` is set as the trust mode.

cos

For 802.1 VLAN-tagged packets, use the priority code point field value from the outermost VLAN header as the index into the CoS map. If the packet is untagged, use the priority configured for CoS map value 0.

dscp

For IP packets, used the code point field value as the index into the DSCP map.

For non-IP packets with 802.1 VLAN tags, use the priority code point field value of the outermost tag header as the index into the CoS map. For untagged, non-IP packets, use the priority configured for CoS map value 0.

Authority

Administrators or local user group members with execution rights for this command.

Example

Setting the global trust mode to **dscp**, which is applied to all interfaces that do not already have an individual trust mode configured. An override is then applied to interface **2/2/2**, and LAG 100, setting trust mode to **cos**:

```
switch(config)# qos trust dscp
switch(config)# interface 2/2/2
switch(config-if)# qos trust cos
switch(config-if)# interface lag 100
switch(config-if)# qos trust cos
```

rate-limit

Syntax

```
rate-limit {broadcast | multicast | icmp {ip-all | ipv4 | ipv6}} <RATE> {kbps | pps}

no rate-limit {broadcast | multicast | icmp}
```

Description

Sets the amount of traffic of a specific type that can ingress on an Ethernet interface, or on each port of a LAG interface. Rate limits are enforced separately on each individual member of a LAG, not on the LAG as a whole.

The **no** form of this command removes the traffic limit for the specified traffic type.

Command context

config-if

Parameters

{broadcast | multicast | icmp {ip-all | ipv4 | ipv6}}

Specifies the type of ingress traffic to which the rate limit applies: broadcast, multicast, or ICMP.

The multicast rate limit affects multicast and broadcast traffic. The broadcast rate limit only affects broadcast traffic. When both types are applied to the same interface, broadcast packets are limited to the lower of the two rate values. Layer 2 BPDU packets, like spanning tree, are also included in the multicast rate limit.

The ICMP rate limit can be configured to apply to IPv4, IPv6, or all IP traffic. Only one ICMP rate-limit can be configured at a time. Applying a new ICMP rate-limit replaces any previous ICMP rate-limit.

<RATE> {kbps | pps}

Specifies the rate limit. Range: 64 to 100000000 kbps (in steps of 64 kbps), or 64 to 209090910 kbps (in steps of 64 pps). The actual rate limit varies with steps approximately equal to the minimum value. Verify the actual rate limit using the command `show interface <INTERFACE-NAME>`.

Authority

Administrators or local user group members with execution rights for this command.

Examples

On the 6400 Switch Series, interface identification differs.

Limiting broadcast traffic to **2000 pps** on interface **1/1/3**:

```
switch(config)# interface 1/1/3  
switch(config-if)# rate-limit broadcast 500 kbps
```

Limiting all ICMP IPv4 traffic to **10000 kbps** on interface **1/1/3**:

```
switch(config)# interface 1/1/3  
switch(config-if)# rate-limit icmp ipv4 10000 kbps
```

Viewing the results of the previous configuration settings.

```
switch# show interface 1/1/3 qos  
Interface 1/1/3 is up  
Admin state is up  
Description:  
Hardware: Ethernet, MAC Address: 08:97:34:b1:20:00  
MTU 1500  
Type 1000BT  
qos trust none  
rate-limit broadcast 2000 pps  
(2000 actual)  
rate-limit icmp ip-all 10000 kbps  
(10000 actual)  
Speed 1000 Mb/s  
L3 Counters: Rx Disabled, Tx Disabled  
Auto-Negotiation is on  
Flow-control: off  
Rx  
          0 input packets          0 bytes  
          0 input error            0 dropped  
          0 CRC/FCS  
L3:  
          0 packets, 0 bytes  
Tx  
          127 output packets        16510 bytes  
          0 input error              0 dropped  
          0 collision  
L3:  
          0 packets, 0 bytes
```

show interface queues

Syntax

```
show interface <INTERFACE-NAME> queues [vsx-peer]
```

Description

Shows the traffic transmitted per queue and the number of packets dropped due to the queue being full.

Command context

Operator (>) or Manager (#)

Parameters

<INTERFACE-NAME>

Specifies the name of an interface on the switch. Format: member/slot/port.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Usage

Statistics include:

- **Tx Bytes:** Total bytes transmitted (available on the 6200 and 6300 Switch Series). The byte count may include packet headers and internal metadata that are removed before the packet is transmitted. Packet headers added when the packet is transmitted may not be included.
- **Tx Packets:** Total packets transmitted.
- **Tx Drops:** The number of packets dropped by a queue before it was sent. When traffic cannot be forwarded out an egress interface, it backs up at ingress. The more servicing assigned to a queue by a schedule profile, the less likely traffic destined for that queue will back up and be dropped. Tx Drops shows the sum of packets that were dropped across all line modules (due to insufficient capacity) by the ingress Virtual Output Queues (VOQs) destined for the egress port. Applicable to the 6400 Switch Series: Because the counts are read separately from each line module, the sum is not an instantaneous snapshot.

Examples

On the 6400 Switch Series, interface identification differs.

Showing settings for interface 1/1/5:

```
switch# show interface 1/1/5
Interface 1/1/5 is down (Administratively down)
Admin state is down
State information: admin_down
Hardware: Ethernet, MAC Address: aa:55:aa:55:00:29
MTU 1500
Full-duplex
qos trust cos
qos schedule-profile default
qos shape 400000
qos cos 4 (not applied)
qos dscp 46 (not applied)
rate-limit broadcast 500 kbps
rate-limit multicast 4000 pps
Speed 0 Mb/s
Auto-Negotiation is turned on
Flow-control: off
```

Showing queue statistics for interface 1/1/5:

(Tx Bytes is available on the 6200 and 6300 Switch Series.)

```
switch# show interface 1/1/5 queues
Interface 5 is down
Admin state is up
      Tx Bytes      Tx Packets      Tx Drops
Q0              0              0              3
Q1            15356              73              0
Q2              0              0              0
Q3              0              0              0
Q4              0              0              0
Q5              0              0              0
Q6              0              0              0
Q7              0              0              0
```

show interface qos

Syntax

```
show interface <INTERFACE-NAME> qos
```

Description

Shows various QoS settings for a specific interface.

Command context

Operator (>) or Manager (#)

Parameters

<INTERFACE-NAME>

Specifies the name of an interface on the switch. Format: member/slot/port or lag number.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Examples

On the 6400 Switch Series, interface identification differs.

Showing QoS settings for interface 1/1/5:

```
switch# show interface 1/3/5 qos
Interface 1/1/5 is up
Admin state is up
qos trust none (global)
qos queue-profile factory-default (global)
qos schedule-profile eqs (override)
qos cos 5
qos dscp 47
rate-limit broadcast 40000 kbps (40000 actual)
```

```
rate-limit icmp ip-all 10000 kbps (10000 actual)
```

	Max-Bandwidth Kbps
Q6	100000
Q7	150000

show qos cos-map

Syntax

```
show qos cos-map [default] [vsx-peer]
```

Description

Shows the global QoS CoS code point settings, or the factory default settings.

Command context

Operator (>) or Manager (#)

Parameters

default

Shows the factory default CoS code point settings.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Examples

Showing the current CoS map:

```
switch# show qos cos-map
code_point local_priority color name
-----
0          2              green Bulk
1          0              green Scavenger
2          1              green Best_Effort
3          3              green Critical_Applications
4          4              green Video
5          5              green Voice
6          6              green Internetwork_Control
7          7              green Network_Control
```

Showing the default CoS map:

```
switch# show qos cos-map default
code_point local_priority color name
-----
0          1              green Best_Effort
1          0              green Background
2          2              green Excellent_Effort
3          3              green Critical_Applications
4          4              green Video
```

5	5	green	Voice
6	6	green	Internetwork_Control
7	7	green	Network_Control

(Color is reserved for future use.)

show qos dscp-map

Syntax

```
show qos dscp-map [default] [vsx-peer]
```

Description

Shows the global QoS DSCP code point settings, or the factory default settings.

Command context

Operator (>) or Manager (#)

Parameters

default

Shows the factory default DSCP code point settings.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Examples

Showing the current DSCP map:

```
switch# show qos dscp-map
code_point  local_priority  cos  color  name
-----
0           1              3    green  CS0
1           1              3    green
2           1              3    green
3           1              3    green
4           1              3    green
5           1              3    green
6           1              3    green
7           1              3    green
8           0              3    green  CS1
...
45          5              7    green
46          7              7    green  EF
47          5              7    green
48          6              7    green  CS6
...
61          7              7    green
62          7              7    green
63          7              7    green
```

Showing the default DSCP map:

```
switch# show qos dscp-map default
code_point  local_priority  cos  color  name
-----
0           1              green CS0
1           1              green
2           1              green
3           1              green
4           1              green
5           1              green
...
45          5              green
46          5              green EF
47          5              green
48          6              green CS6
...
61          7              green
62          7              green
63          7              green
```

(Color is reserved for future use.)

show qos queue-profile

Syntax

```
show qos queue-profile [<NAME> | factory-default ] [vsx-peer]
```

Description

Shows the status of all queue profiles, or a specific queue profile.

Command context

Operator (>) or Manager (#)

Parameters

<NAME>

Specifies the name of a queue profile. Range 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

factory-default

Specifies the default profile.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Example

Showing the status of queue profile **myprofile**:

```
switch# show qos queue-profile myprofile
queue_num local_priorities name
-----
0          2                Normal_and_Bulk
1          0                Other
2          1                Best_Effort
3          3
4          4
5          5
6          6
7          7                High_Priority
```

Showing the status of the default queue profile **factory-default**:

```
switch# show qos queue-profile factory-default
queue_num local_priorities name
-----
0          0                Scavenger_and_backup_data
1          1
2          2
3          3
4          4
5          5
6          6
7          7
```

show qos schedule-profile

Syntax

```
show qos schedule-profile [<NAME> | factory-default] [vsx-peer]
```

Description

Shows the status of all schedule profiles, or a specific schedule profile.

Command context

Operator (>) or Manager (#)

Parameters

<NAME>

Specifies the name of a queue or schedule profile. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

factory-default

Specifies the default schedule profile.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Example

Showing the status of schedule profile **myschedule**:

```
switch# show qos schedule-profile myschedule
queue_num algorithm weight
-----
0          dwrr       1
1          dwrr       2
2          dwrr       4
3          dwrr       8
4          dwrr      16
5          dwrr      24
6          dwrr      32
7          dwrr      64
```

Showing the status of the default schedule profile:

```
switch# show qos schedule-profile factory-default
queue_num algorithm weight
-----
0          dwrr       1
1          dwrr       1
2          dwrr       1
3          dwrr       1
4          dwrr       1
5          dwrr       1
6          dwrr       1
7          dwrr       1
```

show qos trust

Syntax

```
show qos trust [default] [vsx-peer]
```

Description

Shows the global QoS trust settings, or the factory default settings.

Command context

Operator (>) or Manager (#)

Parameters

default

Shows the factory default QoS trust settings.

[vsx-peer]

Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Authority

Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

Examples

Showing the current QoS trust settings:

```
switch# show qos trust
qos trust cos
```

Showing the default QoS trust settings:

```
switch# show qos trust default
qos trust none
```

strict queue

Syntax

```
strict queue <QUEUE-NUMBER> [[max-bandwidth <BANDWIDTH>]]
```

```
no strict queue <QUEUE-NUMBER>
```

Description

Assigns the strict priority algorithm to a queue. Strict priority services all packets waiting in a queue, before servicing the packets in lower priority queues.

Egress queue shaping can be configured using the `max-bandwidth` option to limit the amount of traffic transmitted per output queue. The buffer associated with each egress queue stores the excess traffic to smooth the output rate. Sustained rates of traffic above the maximum bandwidth will eventually fill the output queue causing tail drops. Use the command `show interface` to determine if any tail drop errors have occurred.

The `no` form of this command removes the queue configuration from the schedule profile. To remove only egress queue shaping, re-enter the `strict queue` command without the `max-bandwidth` parameter.

Command context

```
config-schedule
```

Parameters

<QUEUE-NUMBER>

Specifies the number of the queue. Range: 0 to 7.

max-bandwidth <BANDWIDTH>

Specifies the maximum bandwidth allowed on the queue in Kbps. Range: 49 to 100000000.

Authority

Administrators or local user group members with execution rights for this command.

Examples

Assigning strict priority to queue 7 in the schedule profile `myschedule`:

```
switch(config)# qos schedule-profile myschedule
switch(config-schedule)# strict queue 7
```

Deleting strict priority from queue 7 in the schedule profile `myschedule`:

```
switch(config)# qos schedule-profile myschedule
switch(config-schedule)# no strict queue 7
```


Assigning strict priority to queue 7 in the schedule profile **myschedule** with a maximum bandwidth of **10000 Kbps**:

```
switch(config)# qos schedule-profile myschedule  
switch(config-schedule)# strict queue 7 max-bandwidth 10000
```

Accessing Aruba Support

Aruba Support Services	https://www.arubanetworks.com/support-services/
Aruba Support Portal	https://asp.arubanetworks.com/
North America telephone	1-800-943-4526 (US & Canada Toll-Free Number) +1-408-754-1200 (Primary - Toll Number) +1-650-385-6582 (Backup - Toll Number - Use only when all other numbers are not working)
International telephone	https://www.arubanetworks.com/support-services/contact-support/

Be sure to collect the following information before contacting Support:

- Technical support registration number (if applicable)
- Product name, model or version, and serial number
- Operating system name and version
- Firmware version
- Error messages
- Product-specific reports and logs
- Add-on products or components
- Third-party products or components

Other useful sites

Other websites that can be used to find information:

Airheads social forums and Knowledge Base	https://community.arubanetworks.com/
Software licensing	https://lms.arubanetworks.com/
End-of-Life information	https://www.arubanetworks.com/support-services/end-of-life/
Aruba software and documentation	https://asp.arubanetworks.com/downloads

Accessing updates

To download product updates:

Aruba Support Portal

<https://asp.arubanetworks.com/downloads>

If you are unable to find your product in the Aruba Support Portal, you may need to search My Networking, where older networking products can be found:

My Networking

<https://www.hpe.com/networking/support>

To view and update your entitlements, and to link your contracts and warranties with your profile, go to the Hewlett Packard Enterprise Support Center **More Information on Access to Support Materials** page:

<https://support.hpe.com/portal/site/hpsc/aae/home/>



IMPORTANT: Access to some updates might require product entitlement when accessed through the Hewlett Packard Enterprise Support Center. You must have an HP Passport set up with relevant entitlements.

Some software products provide a mechanism for accessing software updates through the product interface. Review your product documentation to identify the recommended software update method.

To subscribe to eNewsletters and alerts:

<https://asp.arubanetworks.com/notifications/subscriptions> (requires an active Aruba Support Portal (ASP) account to manage subscriptions). Security notices are viewable without an ASP account.

Warranty information

To view warranty information for your product, go to <https://www.arubanetworks.com/support-services/product-warranties/>.

Regulatory information

To view the regulatory information for your product, view the *Safety and Compliance Information for Server, Storage, Power, Networking, and Rack Products*, available at <https://www.hpe.com/support/Safety-Compliance-EnterpriseProducts>

Additional regulatory information

Aruba is committed to providing our customers with information about the chemical substances in our products as needed to comply with legal requirements, environmental data (company programs, product recycling, energy efficiency), and safety information and compliance data, (RoHS and WEEE). For more information, see <https://www.arubanetworks.com/company/about-us/environmental-citizenship/>.

Documentation feedback

Aruba is committed to providing documentation that meets your needs. To help us improve the documentation, send any errors, suggestions, or comments to Documentation Feedback (docsfeedback-switching@hpe.com). When submitting your feedback, include the document title, part number, edition, and publication date located on the front cover of the document. For online help content, include the product name, product version, help edition, and publication date located on the legal notices page.