

AOS-CX 10.14.xxxx Quality of Service Guide

832x, 9300, 10000 Switch Series



**Hewlett Packard
Enterprise**

Published: July 2024

Version: 1

Copyright Information

© Copyright 2024 Hewlett Packard Enterprise Development LP.

This product includes code licensed under certain open source licenses which require source compliance. The corresponding source for these components is available upon request. This offer is valid to anyone in receipt of this information and shall expire three years following the date of the final distribution of this product version by Hewlett Packard Enterprise Company. To obtain such source code, please check if the code is available in the HPE Software Center at <https://myenterpriselicense.hpe.com/cwp-ui/software> but, if not, send a written request for specific software version and product for which you want the open source code. Along with the request, please send a check or money order in the amount of US \$10.00 to:

Hewlett Packard Enterprise Company
Attn: General Counsel
WW Corporate Headquarters
1701 E Mossy Oaks Rd Spring, TX 77389
United States of America.



Notices

The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Confidential computer software. Valid license from Hewlett Packard Enterprise required for possession, use, or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

Links to third-party websites take you outside the Hewlett Packard Enterprise website. Hewlett Packard Enterprise has no control over and is not responsible for information outside the Hewlett Packard Enterprise website.

Acknowledgments

Intel®, Itanium®, Optane™, Pentium®, Xeon®, Intel Inside®, and the Intel Inside logo are trademarks of Intel Corporation in the U.S. and other countries.

Microsoft® and Windows® are either registered trademarks or trademarks of Microsoft Corporation in the United States and/or other countries.

Adobe® and Acrobat® are trademarks of Adobe Systems Incorporated.

Java® and Oracle® are registered trademarks of Oracle and/or its affiliates.

UNIX® is a registered trademark of The Open Group.

All third-party marks are property of their respective owners.

Contents

About this document	8
Applicable products	8
Latest version available online	8
Command syntax notation conventions	8
About the examples	9
Identifying switch ports and interfaces	9
QoS overview	11
End-to-end QoS behavior	11
Best effort service	11
Class of Service	11
Differentiated services	12
QoS on the switch	13
QoS trust	15
Dynamic QoS trust mode	16
Port rate limiting	16
Queue profiles	17
Schedule profiles	18
Egress queue shaping	18
Egress port shaping	19
Active Queue Management	19
Explicit Congestion Notification (ECN)	19
Weighted Random Early Detection (WRED)	19
Threshold profiles	19
Queue Monitoring	19
Queue depth	20
Aggregate statistics	20
Terms	20
QoS configuration	22
Configuring QoS	22
Configuring expedited forwarding for VoIP traffic	24
Configuring rate limiting	25
Configuring egress queue shaping	26
Configuring egress port shaping	27
Configuring threshold profiles	28
Monitoring queue operation	30
Data center bridging for storage and lossless Ethernet	31
Data center bridging components	31
PFC - Priority-based flow control	31
ETS - Enhanced transmission selection	31
QCN - Quantized congestion notification	32
DCBx - Data center bridging exchange protocol	32
DCBx guidelines	32
IP ECN	33
Host network interface cards	33
DCB layer 3 configuration task list	34
DCBx configuration	34

DCBx configuration considerations	35
Enabling DCBx	35
Priority-based flow control	36
PFC and the QoS pool	37
Asymmetric PFC	38
Configurable flow control buffer thresholds	38
XOFF thresholds	39
XON delta	39
Headroom size	39
Considerations and prerequisites	40
QoS queue profile	40
QoS queue profile configuration	42
QoS schedule profile - enhanced transmission selection	42
DWRR Calculation	43
QoS schedule profile configuration	43
Overriding the global schedule profile on an interface	44
IP explicit congestion notification	47
IP ECN configuration considerations	48
IP ECN configuration procedure	49
Lossless QoS pool	49
Considerations and prerequisites	50
QoS pool configuration	50
QoS trust	51
Layer 3 CoS-DSCP markings	52
Verifying trust settings	53
Troubleshooting data center bridging	53

QoS commands 57

apply qos	57
apply qos threshold-profile	59
dcbx application	60
dwrr queue	62
flow-control buffer cable-length	63
lldp dcbx (global)	64
lldp dcbx (per interface)	65
lldp dcbx disable	67
map queue	67
name queue	69
qos cos-map	70
qos dscp	71
qos dscp-map	73
qos pool	74
qos queue-profile	76
qos schedule-profile	77
qos shape	79
qos threshold-profile	80
qos trust	81
queue action	82
queue-monitor	85
rate-limit	85
show dcbx interface	88
show interface queues	92
show interface qos	93
show interface qos buffers	95
show qos cos-map	97
show qos dscp-map	98

show qos pool	99
show qos queue-profile	102
show qos schedule-profile	103
show qos threshold-profile	104
show qos trust	106
show queue-monitor	106
strict queue	109

Support and Other Resources	112
--	------------

Accessing HPE Aruba Networking Support	112
Accessing Updates	113
Warranty Information	113
Regulatory Information	113
Documentation Feedback	113

This document describes features of the AOS-CX network operating system. It is intended for administrators responsible for installing, configuring, and managing Aruba switches on a network.

Applicable products

This document applies to the following products:

- HPE Aruba Networking 8320 Switch Series (JL479A, JL579A, JL581A)
- HPE Aruba Networking 8325 Switch Series (JL624A, JL625A, JL626A, JL627A)
- HPE Aruba Networking 9300 Switch Series (R9A29A, R9A30A, R8Z96A, S0F81A, S0F82A, S0F83A, S0F84A, S0F85A, S0F86A, S0F87A, S0F88A, S0F95A, S0F96A)
- HPE Aruba Networking 10000 Switch Series (R8P13A, R8P14A)

Latest version available online

Updates to this document can occur after initial publication. For the latest versions of product documentation, see the links provided in [Support and Other Resources](#).

Command syntax notation conventions

Convention	Usage
<code>example-text</code>	Identifies commands and their options and operands, code examples, filenames, pathnames, and output displayed in a command window. Items that appear like the example text in the previous column are to be entered exactly as shown and are required unless enclosed in brackets ([]).
example-text	In code and screen examples, indicates text entered by a user.
Any of the following: ■ <code><example-text></code> ■ <code><example-text></code> ■ <i>example-text</i> ■ <i>example-text</i>	Identifies a placeholder—such as a parameter or a variable—that you must substitute with an actual value in a command or in code: ■ For output formats where italic text cannot be displayed, variables are enclosed in angle brackets (<code>< ></code>). Substitute the text—including the enclosing angle brackets—with an actual value. ■ For output formats where italic text can be displayed, variables might or might not be enclosed in angle brackets. Substitute the text including the enclosing angle brackets, if any, with an actual value.
	Vertical bar. A logical OR that separates multiple items from which you can choose only one. Any spaces that are on either side of the vertical bar are included for

Convention	Usage
	readability and are not a required part of the command syntax.
{ }	Braces. Indicates that at least one of the enclosed items is required.
[]	Brackets. Indicates that the enclosed item or items are optional.
... or ...	Ellipsis: ■ In code and screen examples, a vertical or horizontal ellipsis indicates an omission of information. ■ In syntax using brackets and braces, an ellipsis indicates items that can be repeated. When an item followed by ellipses is enclosed in brackets, zero or more items can be specified.

About the examples

Examples in this document are representative and might not match your particular switch or environment.

The slot and port numbers in this document are for illustration only and might be unavailable on your switch.

Understanding the CLI prompts

When illustrating the prompts in the command line interface (CLI), this document uses the generic term **switch**, instead of the host name of the switch. For example:

```
switch>
```

The CLI prompt indicates the current command context. For example:

```
switch>
```

Indicates the operator command context.

```
switch#
```

Indicates the manager command context.

switch(CONTEXT-NAME)#

Indicates the configuration context for a feature. For example:

```
switch(config-if) #
```

Identifies the **interface** context.

Variable information in CLI prompts

In certain configuration contexts, the prompt may include variable information. For example, when in the VLAN configuration context, a VLAN number appears in the prompt:

```
switch(config-vlan-100) #
```

When referring to this context, this document uses the syntax:

```
switch(config-vlan-<VLAN-ID>) #
```

Where <VLAN-ID> is a variable representing the VLAN number.

Identifying switch ports and interfaces

Physical ports on the switch and their corresponding logical software interfaces are identified using the format:

```
member/slot/port
```

On the Aruba 8xxx, 9300, and 10000 Switch Series

- *member*: Always 1. VSF is not supported on this switch.
- *slot*: Always 1. This is not a modular switch, so there are no slots.
- *port*: Physical number of a port on the switch.

For example, the logical interface **1/1/4** in software is associated with physical port 4 on the switch.



If using breakout cables, the port designation changes to x:y, where x is the physical port and y is the lane when split to 4 x 10G or 4 x 25G. For example, the logical interface 1/1/4:2 in software is associated with lane 2 on physical port 4 in slot 1 on member 1.

Quality of Service (QoS) enables network administrators to customize how different types of traffic are serviced on a network, taking into account the unique characteristics of each traffic type and its importance within an organization's infrastructure. QoS ensures uniform and efficient traffic handling, keeping the most important traffic moving at an acceptable speed, regardless of current bandwidth usage. It also provides methods for administrators to control the priority settings of inbound traffic arriving at each network device.

End-to-end QoS behavior

The QoS settings on each network device must be aligned to achieve the desired end-to-end QoS behavior for a network. Three service types can be used to categorize and prioritize network traffic:

- Best Effort Service
- Ethernet Class of Service (CoS)
- Internet Differentiated Services (DiffServ)

For a network as a whole, it is best to select one service type to use as the primary end-to-end behavior, and then use the other two service types as needed.

Best effort service

This is the simplest service type. All traffic is treated equally in a first-come, first-served manner. If the traffic load is low in relation to the capacity of the network links, then there is no need for the administrative complexity and costs of maintaining a more complex end-to-end policy. This is sometimes called over-provisioning, as all link speeds are much higher than peak loads on the network.

Class of Service

Class of Service (CoS) is a method for classifying network traffic at layer 2 by marking 802.1Q VLAN Ethernet frames with one of eight service classes.

CoS	Traffic type	Example protocols
7	Network Control	STP, PVST
6	Internetwork Control	BGP, OSPF, PIM
5	Voice (<10ms latency)	VoIP(UDP)
4	Video (<100ms latency)	RTP
3	Critical Applications	SQL RPC, SNMP

CoS	Traffic type	Example protocols
2	Excellent Effort	NFS, SMB
0	Best Effort	HTTP, TELNET
1	Background	SMTP, IMAP

CoS 1 is deliberately set as the lowest CoS. This enables a traffic service level below the default (best effort) traffic level to be specified.

The 3-bit Priority Code Point (PCP) field within the 16-bit Ethernet VLAN tag is used to mark the CoS.




Differentiated services

Differentiated services (DiffServ) is a method for classifying network traffic at layer 3 by marking packets with one of 64 different service classes. Services classes are identified by the Differentiated services Code Point (DSCP) value. Some common DSCP values are:

DSCP	Name	Service class	RFC
48	CS6	Network Control	2474
46	EF	Telephony	3246
40	CS5	Signaling	2474
34, 36, 38	AF41, AF42, AF43	Multimedia Conferencing	2597
32	CS4	Real-Time Interactive	2474
26, 28, 30	AF31, AF32, AF33	Multimedia Streaming	2597
24	CS3	Broadcast Video	2474
18, 20, 22	AF21, AF22, AF23	Low-Latency Data	2597
16	CS2	OAM	2474
00	CS0,BE,DF	Best Effort	2474
10, 12, 14	AF11, AF12, AF13	Bulk Data	2597
08	CS1	Low-Priority Data	3662

DSCP CS1 (08) CoS 1 is deliberately set as the lowest priority. This enables a traffic service level below the standard (best effort or default forwarding) level to be specified.

The DSCP value is carried within the IPv4 DSCP field or the upper 6-bits of the 8-bit IPv6 Traffic Class (TC) field.



QoS on the switch

There are five key stages a packet passes through when traversing a switch: ingress, prioritization, destination determination, egress queuing, and transmission. The following table provides an overview of each stage, and lists the commands that can be used to configure QoS settings.




Switches with at least 52 ports will experience negative performance if a flood occurs where at least 42 ports are members of the same VLAN and all 52 ports have QoS rules applied to them.

Ingress	Prioritization	Classification	Queueing	Transmission
Packets arrive at switch interface (port).	An initial local-priority value is assigned to the packet based on VLAN CoS or IP DSCP.	Packets local-priority and code points can be remarked, and rates policed, based on many packet header fields. (see ACL & Policy Guide)	Packets are queued based on the destination interface and local-priority of the packet.	A scheduler defines the order in which packets are selected from queues to be transmitted from the interface.
rate-limit Rate limiting can control ingress flow by packet type: broadcast, multicast, or unknown-unicast.	qos trust {cos dscp none} Assigns which packet values are used to determine initial local-priority either globally for all interfaces, or to specific interface(s). qos cos-map Defines how CoS values are mapped to local-priority values. qos dscp-map Defines how IP DSCP values are mapped to local-priority values. qos dscp Remarks egress DSCP value.	class {ip ipv6 mac} NAME match TERMS [count] ignore TERMS [count] Creates an ordered list of rules to identify packets to match the class. policy NAME class NAME ACTIONS Creates an ordered list of classes to which traffic is evaluated, and the actions to be taken on matching packets: remarking, policing etc. apply policy NAME Assigns the policy to evaluate all traffic inbound or outbound globally, or on specific interfaces or VLANs.	qos queue-profile Creates a profile that defines transmit queues. map queue Assigns a local-priority to a queue. Packets with a matching local-priority are placed in the queue. name queue Assigns a name to a queue. apply qos Assigns queue, schedule, or threshold profiles globally to all interfaces, or schedule or threshold profiles to a specific interface(s).	qos schedule-profile Creates a profile that defines the order in which packets are selected from the queues for transmission. strict queue Assigns the strict priority algorithm, with optional max-bandwidth, to a queue. dwrr queue Assigns the weighted DWRR algorithm to a queue.

Ingress	Prioritization	Classification	Queueing	Transmission
Packets arrive at switch interface (port).	An initial local-priority value is assigned to the packet based on VLAN CoS or IP DSCP.	Packets local-priority and code points can be remarked, and rates policed, based on many packet header fields. (see ACL & Policy Guide)	Packets are queued based on the destination interface and local-priority of the packet.	A scheduler defines the order in which packets are selected from queues to be transmitted from the interface.
	show int IFNAME qos Displays all the QoS settings configured on an interface: • Profiles • Rate-limits • DSCP remark • Trust show qos cos-map Displays the contents of the CoS Map table. show qos dscp-map Displays the content of the DSCP Map table. show qos trust Display the current default trust mode.	show class {ip ipv6 mac} Displays the ordered list of the match and ignore statements in the class. show policy Displays the order list of classes in the policy. show policy hitcounts Displays the ordered list of classes in the policy. For each class the number of times that each match or ignore statement in the class matched a packet.	show int IFNAME queues Displays the per-queue counts of packets transmitted, packets dropped, and bytes transmitted. show qos queue-profile Displays the definition of queue profiles.	show qos schedule-profile Displays the definition of schedule profiles.

The following diagram shows how different packets might traverse a switch. It also shows how QoS configuration settings apply at each stage.



QoS trust


Traffic priorities for networks can be carried in VLAN tags, using the CoS Priority Code Point (PCP), or in IP packet headers, using the Differentiated Services Code Point (DSCP). Whether these priorities affect how traffic is serviced, depends on how QoS trust mode is configured on the switch. QoS trust mode specifies how the switch assigns local priority values to ingress packets. Trust mode can be set globally for all interfaces, or individually for each interface. By default, trust mode is set to `none`, meaning that any QoS information in the packet (CoS or DSCP) is ignored, and local priority values are assigned from the CoS map value for code point 0. An exception to this can be configured, allowing a QoS remark to be applied to DSCP values when trust mode is `none`.

When trust mode is set to `CoS` or `DSCP`, the switch translates the QoS settings in VLAN tags (for `CoS`), or the DS field in an IP header (for `DSCP`), to local priority values on the switch. Translation is controlled by the CoS map or DSCP map tables.

For example:


CoS Map

When QoS trust mode = CoS



DSCP Map

When QoS trust mode = DSCP



When QoS trust mode = None



Dynamic QoS trust mode

The device profile feature can dynamically set the QoS trust mode on an interface based on the LLDP information exchanged with a link partner. The device profile's trust mode temporarily overrides the static trust mode configured for an interface. The override remains in place as long as that link partner is connected and its link state is **up**. Use command `show interface IFNAME qos` to view the current QoS trust mode for an interface.

Port rate limiting

Port rate limiting helps control undesirable traffic. Its purpose is to allow enough unicast, broadcast, multicast, and ICMP traffic for the network to function properly, while preventing flooding and traffic storms.

Some amount of each traffic type is required for normal network operation. For example, broadcast packets may include ARP and DHCP traffic. Video streams and certain types of network protocol packets are multicast traffic. Unknown-unicast packets may be intended for devices whose addresses have temporarily aged out of network forwarding caches. Configuring rate limits can help provide balance between necessary and flooded traffic.

The amount of permitted traffic (the rate limit) can be specified in packets-per-second (pps) or as a percentage of link bandwidth (percent).

Please note that rate limits are enforced separately on each individual member of a LAG, not on the LAG as a whole. When limits are configured in packets-per-second (pps), larger limits may be needed on high speed ports in order to allow normal network functionality. A limit of 100 packets-per-second might be reasonable on a 1 Gigabit port, however would be too restrictive on a port that is 10 Gigabits or faster. Conversely, limits configured as a percentage of link bandwidth may not be useful on high speed ports. A limit of 1 percent might be reasonable on a 1 Gigabit port, but too lax on a port that is 10 Gigabits or faster.

Queue profiles

A queue profile defines the queues that are associated with an interface to control the transmission of packets. Each profile supports up to eight queues, numbered 0 to 7. The larger the queue number, the higher its priority during transmission scheduling. Packets are assigned to a queue based on their local priority value (0 to 7). A queue profile must map all eight local priority values to whatever queues are being used on the switch, and a schedule profile must specify the configuration for those same queues. A queue without a local priority value assigned to it is not used to store packets.

The switch is automatically provisioned with an initial queue profile named `factory-default` which assigns each local priority to the queue of the same number. To see the default queue profile, use the command `show qos queue-profile factory-default`:

```
switch# show qos queue-profile factory-default
queue_num local_priorities name
-----
0          0                  Scavenger_and_backup_data
1          1
2          2
3          3
4          4
5          5
6          6
7          7
```

More than one local priority value can be assigned to the same queue. For example,

Local Priority	Queue
0	0
1	1
2	2
3	3
4	4
5	5
6	5
7	5

Commonly used commands for working with QoS queues are as follows:

- `qos queue-profile`: Creates an empty queue-profile and enters the profile configuration context.
- `name queue`: Assigns a descriptive name to a queue.
- `map queue`: Assigns a local-priority to a queue.
- `apply qos queue-profile`: Applies a queue-profile globally to all interfaces.

Schedule profiles

A schedule profile determines the order in which queues are selected for transmission, and the amount of service available for each queue. A schedule profile must be configured on every interface at all times. A schedule profile can be applied globally to all interfaces, or only to specific interfaces.

The following options are available:

- All queues use deficit weighted round robin queuing (DWRR)
- All queues use strict priority
- The highest priority queue uses strict priority, and all other queues use DWRR
- All queues use guaranteed minimum bandwidth
- The highest priority queue uses strict priority and all other queues use guaranteed minimum bandwidth

A weighted schedule profile assigns relative servicing for each queue. The amount of service per weight is relative to the underlying hardware implementation, and to the weights assigned to the other non-empty queues. Strict scheduling can be used to service queues purely on the basis of highest priority first (at the risk of starving lower-priority queues during high stress periods). A combination of strict and weighted scheduling offers more service to the highest priority queue when needed, while preserving scheduling between the remaining queues, thus decreasing the risk of starvation.

The switch is automatically provisioned with a schedule profile named **factory-default**, which assigns DWRR to all queues with a weight of 1. Use the command `show schedule-profile factory-default` to view the default schedule profile. (Do not use `show running-configuration`, as it only displays changes from the initial settings.)

```
switch# show qos schedule-profile default
queue_num algorithm weight
-----
0          dwrr       1
1          dwrr       1
2          dwrr       1
3          dwrr       1
4          dwrr       1
5          dwrr       1
6          dwrr       1
7          dwrr       1
```

Egress queue shaping

Egress queue shaping limits the amount of traffic transmitted per strict output and dwrr queues. The buffer associated with each egress queue stores excess traffic to absorb bursts and smooths the output rate. For example, an administrator might limit strict-priority or min-bandwidth queue traffic to prevent low-priority queue starvation in the event that a device inappropriately sends too many higher-priority packets.



Egress queue shaping can be configured on an Ethernet port or on a link aggregation group (LAG). To configure egress queue shaping, define a schedule profile with the strict priority algorithm assigned to each queue.

Egress port shaping

Egress port shaping limits the amount of aggregate traffic transmitted through a port. To be effective, the egress port-shaping rate must be less than the port's line rate. By default, the egress port-shaping rate is the same as the line-rate of the port. Buffers associated with each port store excess traffic. When both egress port-shaping and egress queue-shaping are configured on the same interface, the switch respects the minimum of both configurations.

Active Queue Management

Explicit Congestion Notification (ECN)

Explicit Congestion Notification (ECN) provides a mechanism for two end-points to exchange end-to-end notification of network congestion. ECN uses a 2-bit field in the IP header to indicate that the traffic load on network equipment in the path between an ECN-capable sender and receiver is causing packets to be buffered, as defined by IETF RFC 3168 (<https://tools.ietf.org/html/rfc3168>).

Weighted Random Early Detection (WRED)

WRED operates by random early-dropping packets, which can be helpful in signaling data path congestion to certain protocols. Protocols that respond to these drops slow their transmit rate in an effort to reduce network congestion. WRED drops are randomized in order to avoid potential synchronization between multiple streams using the same link. If drops occurred on all streams at the same time, multiple senders might respond by reducing their transmit rates and then increasing. Such synchronized behavior causes link utilization to fluctuate between high and low, wasting bandwidth. Random dropping ensures that only some streams detect drops, and that they detect them at different times. This results in better link utilization, as some senders continue to transmit at a higher rate while others reduce and ramp up again.

Threshold profiles

Threshold profiles configure individual queue utilization thresholds as triggers for taking action (i.e., ECN marking or WRED dropping) on a packet. A threshold profile is applied per-port and defines the thresholds and actions for each queue. Omitting configuration for a queue in a threshold profile means that queue will not be configured with a threshold value or action.

In an environment where responsive transport protocols are in use and congestion management features are required to reduce latency, ECN can be configured on queues carrying delay-sensitive traffic. The result is that queue utilization is actively managed, resulting in ECT packets being CE marked when queue utilization reaches or exceeds a configured threshold.

Queue Monitoring

The queue monitoring feature on 8325, 9300 and 10000 Switch series allows the switch to collect queue statistics at 10-second time intervals, then display this information as a *queue statistics history*. These aggregate statistics can be used to monitor network health, troubleshoot network issues, and identify normal network behavior patterns and performance anomalies. This feature is disabled by default, but

can be enabled on a maximum of 52 interfaces per switch. For more information on enabling this feature, see [queue-monitor](#).

Queue statistics history data can be displayed in the command-line interface as a list or a histogram reflecting the last eight hours of queue statistics history, or as a table showing the previous five minutes of queue statistics history. This information can include:

- A summary of congestion events across the switch
- Per-interface histograms of congestion data
- Per-interface detailed view of congestion data

As performance anomalies are identified, the interface queue experiencing the anomaly can further analyzed to identify the cause of the behavior. For example, traffic patterns for flows competing for network bandwidth can be examined to determine if corrective action is needed to address potential problems.

Queue depth

Queue history statistics can also be filtered by *queue depth*, the memory size required to buffer packets in a queue. The queue depth of an interface changes constantly based on the arrival rate of packets destined to egress a given queue (the queue fill rate) and the rate at which that queue is able to transmit packets (the queue drain rate).

If no queue depth is specified, the output of the **show interface queue-monitor** command displays statistics for queues with a minimum depth of one kbyte, which includes all queues which have reached a depth greater than zero.

Aggregate statistics

The queue monitoring feature protects switch memory resources by automatically aggregating more detailed data samples into one value that represents a larger window of time as the more detailed samples reach a maximum statistics age.

- As 10-second data samples reach an age of 8 hours, they are combined into a 2-minute sample.
- As 2-minute data samples reach an age of 3 days , they are combined into a 30-minute sample.
- As 30 minute data samples reach an age of 14 days, they are combined into a 12-hour sample.
- The 12-hour samples are retained for 45 days.

Queue statistics history information can be sent to an external collector if longer-term data review is required. For example, if you wanted to retain all 10-second data samples, you could configure an external collector to query the switch for 10-second data samples every six hours. Queue statistics history is not retained by the switch when it reboots.



Use the AOS-CX REST API to export and view data from both the 10-second samples and aggregate samples representing up to 12-hours of data. The command-line interface can display information for 10-second data samples only.

Terms

Class

For networking, a set of packets sharing a common characteristic. For example, all IPv4 packets.

Code point

The name of a packet header field, or the value carried within a packet header field:

- Example 1: Priority code point (PCP) is the name of a field in the IEEE 802.1Q VLAN tag.
- Example 2: Differentiated services code point (DSCP) is the name of a field carried within the DS field of an IP packet header.

Color

A metadata label associated with each packet within the switch. It has three values: green (0), yellow (1), or red (2). When packets encounter congestion for a resource (queue), the switch uses packet color to distinguish which packets must be dropped, and is mostly used for packets marked with Assured Forwarding (AF) DSCP values.

Not supported in this release.

Class of service (CoS)

A 3-bit value used to mark packets with one of eight classes (levels of priority). It is carried within the priority code point (PCP) field of the IEEE 802.1Q VLAN tag.

Differentiated services code point (DSCP)

A 6-bit value used to mark packets for different per-hop behavior as originally defined by IETF RFC 2474. It is carried within the differentiated services (DS) field of the IPv4 or IPv6 header.

Local priority

A meta-data label associated with a packet within the switch which is used to classify packets for different treatment (such as queue assignment). Eight local priorities are defined on the switch, numbered from 0 to 7. A queue profile must map all eight local priorities to whatever queues are in use on the switch, and a schedule profile must specify the configuration for these same queues.

Metadata

Information labels associated with each packet in the switch, separate from the packet headers and data. These labels are used by the switch in its handling of the packet. For example: arrival port, egress port, VLAN membership, and local priority.

Priority code point (PCP)

The name of a 3-bit field in the IEEE 802.1Q VLAN tag. It carries the CoS value to mark a packet with one of eight classes (priority levels).

Quality of service (QoS)

General term used when describing or measuring performance. For networking, it means how different classes of packets are treated when traversing a network or device.

Traffic class (TC)

General term for a set of packets sharing a common characteristic. It used to be the name of an 8-bit field in the IPv6 header originally defined by IETF RFC 2460. This field name was changed to differentiated services by IETF RFC 2474.

Type of service (ToS)

General term when there are different levels of treatment (fare class). It used to be the name of an 8-bit field in the IPv4 header originally defined by IETF RFC 791. This field name was changed to differentiated services by IETF RFC 2474.

Configuring QoS

Procedure

1. Configure how local priority values are assigned to ingress packets with the commands `qos cos-map`, `qos dscp-map`, and `qos trust`.
2. Optionally, add a rate limit for ingress traffic on one or more interfaces with the command `rate-limit`.
3. If you do not want to use the default QoS queue profile to map local priority to queue, create one or more custom queue profiles with the command `qos queue-profile`. For each queue in a custom queue profile:
 - a. Assign a local priority value with the command `map queue`.
 - b. Optionally, define a descriptive name with the command `name queue`. All local priorities (0 to 7) must be mapped to a queue, and the queues selected for use must be in contiguous order starting at 0.
4. If you do not want to use the default QoS schedule profile to determine the order in which queues are selected to transmit a packet, create one or more custom schedule profiles with the command `qos schedule-profile`. For each queue in a custom schedule queue profile, define scheduling priority with the commands `strict queue` and `dwrr queue`.
5. Optionally, on 8100 or 8360 Switch Series, create a threshold profile to limit throughput on one or more queues with the command `qos threshold-profile`. Assign threshold values to the queue with the command `queue`. Then, apply the profile with the command `apply qos threshold-profile`.
6. Optionally for strict queues, configure egress queue shaping to limit egress bandwidth on an interface to a value that is less than its line rate. Use the `max-bandwidth` parameter of the `strict queue` command.
7. Activate QoS settings with the command `apply qos`. This command lets you apply a queue profile and schedule profile globally to all interfaces, or a schedule profile override to individual interfaces.

When applying QoS settings to a port configured to support priority-based flow control, specific configuration settings must be respected when defining a CoS map and queue profile. See the command `flow-control` in the *Command-Line Interface Guide* for details.

8. View QoS configuration settings with the provided `show` commands.

Examples

This example creates the following configuration:

- Configures CoS to be used to assign local priority to ingress packets.
- Modifies the default CoS map to assign CoS 1 to local priority 1.

- Creates a queue profile named **Q1** and assigns local priorities as follows:

Queue	Local Priority
0	0
1	1
1	2
2	3
3	4
4	5
5	6
5	7

- Creates a schedule profile named **S1** and assigns DWRR to all queues in the schedule profile with the following weights:

Queue	Weight
0	5
1	10
2	15
3	20
4	25
5	50

- Applies **Q1** and **S1** to all interfaces that do not have a QoS override applied.

```

switch(config)# qos trust cos
switch(config)# qos cos-map 1 local-priority 1
switch(config)# qos queue-profile Q1
switch(config-queue)# map queue 0 local-priority 0
switch(config-queue)# map queue 1 local-priority 1
switch(config-queue)# map queue 1 local-priority 2
switch(config-queue)# map queue 2 local-priority 3
switch(config-queue)# map queue 3 local-priority 4
switch(config-queue)# map queue 4 local-priority 5
switch(config-queue)# map queue 5 local-priority 6
switch(config-queue)# map queue 5 local-priority 7
switch(config-queue)# qos schedule-profile S1
switch(config-schedule)# dwrr queue 0 weight 5
switch(config-schedule)# dwrr queue 1 weight 10
switch(config-schedule)# dwrr queue 2 weight 15
switch(config-schedule)# dwrr queue 3 weight 20
switch(config-schedule)# dwrr queue 4 weight 25

```

```
switch(config-schedule)# dwrr queue 5 weight 50
switch(config)# apply qos queue-profile Q1 schedule-profile S1
```

Configuring expedited forwarding for VoIP traffic

Voice over IP (VoIP) traffic is delay and jitter sensitive. For optimum transmission of VoIP traffic, dwell time in network devices must be kept to a minimum and all network devices in the data path must have identical per-hop behaviors. To configure a dedicated queue on the switch to handle VoIP traffic with priority service before all other queues, follow these steps.

Prerequisites

This scenario assumes that VoIP packets are uniquely identified using DiffServ code point 46, Expedited Forwarding (EF).

Procedure

1. Map DSCP EF packets exclusively to local priority 5. The default DSCP map has eight code points (40 through 47), that are mapped to local priority 5. To reserve local priority 5 for VoIP traffic, the other code points must be reassigned. In this scenario, local priority 6 is used for all reassignments, including for code point 40, Call Signaling protocol (CS5).

```
switch(config)# qos dscp-map 40 local-priority 6 name CS5
switch(config)# qos dscp-map 41 local-priority 6
switch(config)# qos dscp-map 42 local-priority 6
switch(config)# qos dscp-map 43 local-priority 6
switch(config)# qos dscp-map 44 local-priority 6
switch(config)# qos dscp-map 45 local-priority 6
switch(config)# qos dscp-map 47 local-priority 6
```

2. Queue 7 is the highest priority queue, so for best throughput, create a queue profile that maps local priority to queue 7.

```
switch(config)# qos queue-profile ef_priority
switch(config-queue)# name queue 7 Voice_Priority_Queue
switch(config-queue)# map queue 7 local-priority 5
switch(config-queue)# map queue 6 local-priority 7
switch(config-queue)# map queue 5 local-priority 6
switch(config-queue)# map queue 4 local-priority 4
switch(config-queue)# map queue 3 local-priority 3
switch(config-queue)# map queue 2 local-priority 2
switch(config-queue)# map queue 1 local-priority 1
switch(config-queue)# map queue 0 local-priority 0
```

3. Create a schedule profile that services queue 7 using strict priority (SP), and the remaining queues with DWRR. This scenario gives all DWRR queues equal weight.

```
switch(config)# qos schedule-profile voip
switch(config-schedule)# strict queue 7
switch(config-schedule)# dwrr queue 6 weight 1
switch(config-schedule)# dwrr queue 5 weight 1
```



```
switch(config-schedule)# dwrr queue 4 weight 1
switch(config-schedule)# dwrr queue 3 weight 1
switch(config-schedule)# dwrr queue 2 weight 1
switch(config-schedule)# dwrr queue 1 weight 1
switch(config-schedule)# dwrr queue 0 weight 1
switch(config-schedule)# exit
switch(config)#
```

4. Apply the profiles to all interfaces.


```
switch(config)# apply qos queue-profile ef_priority schedule-profile voip
```

5. Configure DSCP trust mode on all ports

```
switch(config)# qos trust dscp
```

Configuring rate limiting

This scenario illustrates how to use rate limiting to manage the traffic from various devices connected to a switch. The physical topology of the network looks like this:



A certain amount of broadcast traffic is necessary to maintain healthy network operation, particularly from routers and across service boundaries. In this scenario, both the service cloud and the router connections limit this traffic to 1 Gbps. The server has a smaller limit, as it does not require as much network protocol traffic as the service cloud and router.

A multicast server needs to be able to stream multicast traffic to clients, so a multicast rate limit may not be helpful. A computer, however, should not be generating large amounts of multicast traffic (it may be receiving streams, but typically not sending them). In this example, the computer is configured with a multicast rate limit to prevent malicious traffic from taking up network bandwidth.

Finally, while the service cloud and router may need to send traffic for unknown unicast addresses to resolve address forwarding, the server and computer should send very little of this type of traffic. Rate limiting unknown unicast traffic on those two devices enforces that.

Procedure

1. Configure broadcast and multicast rate limiting for the service cloud connection.

```
switch# config
switch(config)# interface 1/1/1
switch(config-if)# rate-limit broadcast 500 pps
switch(config-if)# rate-limit multicast 500 pps
switch(config-if)# exit
```

2. Configure broadcast rate limiting for the router connection.

```
switch(config-if)# interface 1/1/2
switch(config-if)# rate-limit broadcast 500 pps
switch(config-if)# exit
```

3. Configure broadcast rate limiting for the server connection.

```
switch(config-if)# interface 1/1/5
switch(config-if)# rate-limit broadcast 100 pps
switch(config-if)# exit
```

4. Configure broadcast, and multicast rate limiting for the computer connection.

```
switch(config-if)# interface 1/1/10
switch(config-if)# rate-limit broadcast 50 pps
switch(config-if)# rate-limit multicast 50 pps
```

Configuring egress queue shaping

This example shows how to apply egress queue shaping to an interface. First, a schedule profile is created that has per-queue bandwidth limits set on all queues with `strict` as the scheduling algorithm. Next, this profile is applied to an interface or LAG.

The following example creates a schedule profile named **EQSExample**, which services all queues using **strict** priority. This profile configures queues 1, 4, and 7 with a bandwidth limit of 10 Gbps, 20 Gbps, and 30 Gbps respectively. Queues 1 and 7 are also configured with a burst of 30 KB (burst configuration is only supported on the 8320, 8325, 9300, and 10000). The **apply qos schedule-profile** command is then used to apply the **EQSExample** profile to interface **1/1/1**. The actual burst and bandwidth configured on an interface can be found by using the **show interface <IF-NAME> qos** command.

```
switch(config)# qos schedule-profile EQSExample
switch(config-schedule)# strict queue 0
switch(config-schedule)# strict queue 1 max-bandwidth 10000000 kbps
switch(config-schedule)# strict queue 2
switch(config-schedule)# strict queue 3
switch(config-schedule)# strict queue 4 max-bandwidth 20000000 kbps
switch(config-schedule)# strict queue 5
switch(config-schedule)# strict queue 6
switch(config-schedule)# strict queue 7 max-bandwidth 30000000 kbps
switch(config-schedule)# exit
switch(config)# interface 1/1/1
switch(config-if)# apply qos schedule-profile EQSExample
```

The following example creates a schedule profile named **EQSEExample** which services all queues using **strict** priority. This profile configures queues 1, 4 and 7 with a bandwidth limit of 10%, 20%, and 30% of link bandwidth respectively. Queues 1 and 7 are also configured with a burst of 30 KB. The **apply qos schedule-profile** command is then used to apply the **EQSEExample** profile to the port. The actual burst and bandwidth configured on an interface can be found by using the **show interface <IF-NAME> qos** command.

```
(config)# qos schedule-profile EQSEExample
(config-schedule)# strict queue 0
(config-schedule)# strict queue 1 max-bandwidth 10 percent
(config-schedule)# strict queue 2
(config-schedule)# strict queue 3
(config-schedule)# strict queue 4 max-bandwidth 20 percent
(config-schedule)# strict queue 5
(config-schedule)# strict queue 6
(config-schedule)# strict queue 7 max-bandwidth 30 percent
```

Configuring egress port shaping

This example shows how to apply egress port shaping to an interface to limit the rate of egress traffic. Egress port shaping is configured by specifying the desired bandwidth rate in kilobits per second (kbps) or as a percentage of link bandwidth with an optional burst size in kilobytes (KB). In order to be effective, the configured rate must be less than the link rate of the port. If the configured egress port-shaping rate exceeds the port's link rate, then egress port shaping is effectively disabled. Egress port shapes are enforced separately on each individual member of a LAG, not on the LAG as a whole. Because shapes are in kilobits-per-second (kbps), larger shapes may be needed on high speed ports in order to allow normal network function. A shape of 100 kbps might be reasonable on a 1-Gigabit port, but too restrictive on a port that is 10-Gigabits or faster.

The configured egress rate and burst on a specific interface can be found by using the **show interface** and **show interface qos** commands.

The following example configures an egress rate of 100 Mbps and a burst of 30 KB:

```
switch(config)# interface 1/1/1
switch(config-if)# qos shape 100000 kbps burst 30
```

In the next example, both egress port shaping and egress queue shaping are configured on the same interface.

The example creates a schedule profile named **EQSEExample** with strict priority for all seven queues. Queue 7 is configured with a bandwidth limit of 300 Mbps and a burst of 30 KB. The profile is then applied to interface **1/1/1** with egress port shaping of 400 Mbps and a burst size of 30 KB. As egress queue shaping and egress port shaping are both configured on port **1/1/1**, egress queue shaping is subject to the lower port or queue shape rate. The effective bandwidth for the traffic egressing on queue 7 will be 300 Mbps and the effective burst will be 30 KB.

```
switch(config)# qos schedule-profile EQSEExample
switch(config-schedule)# strict queue 0
switch(config-schedule)# strict queue 1
switch(config-schedule)# strict queue 2
switch(config-schedule)# strict queue 3
switch(config-schedule)# strict queue 4
```

```
switch(config-schedule)# strict queue 5  
switch(config-schedule)# strict queue 6  
switch(config-schedule)# strict queue 7 max-bandwidth 300000 kbps burst 30  
switch(config-schedule)# exit  
switch(config)# interface 1/1/1  
switch(config-if)# apply qos schedule-profile EQSExample  
switch(config-if)# qos shape 400000 kbps burst 30
```

Configuring an egress port shape of 10 percent and a burst of 30 KB on port 1/1/1. The configured egress rate and burst on a specific interface can be found by using the **show interface** and **show interface qos** commands.

```
switch(config)# interface 1/1/1  
switch(config-if)# qos shape 10 percent burst 30
```

Configuring threshold profiles

The threshold profile is an optional configuration that specifies a per-packet action to be taken for each port queue when the queue utilization reaches or exceeds the configured threshold. The per-queue configuration within the threshold profile is optional, which means any number of queues may be configured. An unconfigured queue means that no threshold action behavior will occur. Also, an empty threshold profile is valid and can be applied.

Use the `apply qos threshold-profile <THRESHOLD-NAME>` command in the global context to configure all ports to use the profile, or in the interface context to configure the interface to use the profile. No threshold profiles are created or configured by default.

Each entry in a threshold profile can configure a WRED action for a particular queue and color combination. Each entry in a threshold profile can configure an ECN action or a WRED action, either of which can take effect depending on the contents of the packet being added to the queue. For WRED action, it is possible to configure separate thresholds for responsive protocol packets (for example, TCP) as well as non-responsive protocol packets (for example, UDP).

In an environment where congestion management features are required in order to reduce latency and responsive transport protocols are in use, ECN can be configured on queues carrying delay-sensitive traffic. The desired result is that a network device's port-queue utilization is actively managed, resulting in ECN-capable transport (ECT) packets being Congestion Encountered (CE) marked when queue utilization reaches or exceeds a configured threshold. Threshold profiles can be either strict or probabilistic. With a strict threshold, the action will always occur when the threshold is exceeded. For a probabilistic profile, minimum and maximum thresholds are set and the action starts probabilistically occurring when the minimum threshold is crossed. The action probability increases until the maximum threshold is crossed, after which it occurs on every packet.

Procedure

To configure one or more port queues with ECN or WRED, complete the following steps:

1. Create a threshold profile defining one or more queue threshold configurations with an ECN or WRED action.

(Option 1) Create a threshold profile with an ECN action on queue 7:

```
switch(config)# qos threshold-profile threshprofile
switch(config-threshold)# queue 7 action ecn all threshold 40 kbytes
```

(Option 2) Create a probabilistic threshold profile with an ECN and WRED action on queue 6:

```
switch(config)# qos threshold-profile threshprofile
switch(config-threshold)# queue 6 action ecn all min-threshold 40 kbytes
max-threshold 60 kbytes
switch(config-threshold)# queue 6 action wred-resp yellow min-threshold 80
kbytes max-threshold 500 kbytes max-prob 40 percent
switch(config-threshold)# queue 6 action wred-non-resp yellow min-threshold
100 kbytes max-threshold 550 kbytes max-prob 50 percent
```

2. Apply the threshold profile globally (all ports), OR apply the Threshold Profile as a port override to any ports that require ECN or WRED behavior.

(Option 1) Apply the threshold profile as a global default (all ports):

```
switch(config)# apply qos threshold-profile threshprofile
```

(Option 2) Apply the threshold profile to specific Ethernet or LAG interfaces as a port override:

```
switch(config)# int 1/1/1
switch(config-if)# apply qos threshold-profile threshprofile
switch(config)# int 1/1/2
switch(config-if)# apply qos threshold-profile threshprofile
switch(config)# int lag 10
switch(config-if)# apply qos threshold-profile threshprofile
```

To configure one or more port queues with WRED, complete the following steps:

1. Create a threshold profile defining one or more queue threshold configurations with a WRED action.

```
switch(config)# qos threshold-profile threshprofile
switch(config-threshold)# queue 6 action wred-resp yellow min-threshold 80
kbytes max-threshold 500 kbytes max-prob 40 percent
switch(config-threshold)# queue 6 action wred-non-resp yellow min-threshold
100 kbytes max-threshold 550 kbytes max-prob 50 percent
```

2. Apply the threshold profile globally (all ports), or apply the threshold profile as a port override to any ports that require WRED behavior.

(Option 1) Apply the threshold profile as a global default (all ports):

```
switch(config)# apply qos threshold-profile threshprofile
```

(Option 2) Apply the threshold profile to specific ethernet or LAG interfaces as a port override:

```

switch(config)# int 1/1/1
switch(config-if)# apply qos threshold-profile threshprofile
switch(config)# int 1/1/2
switch(config-if)# apply qos threshold-profile threshprofile
switch(config)# int lag 10
switch(config-if)# apply qos threshold-profile threshprofile

```

Monitoring queue operation

Use the show interface queues command to display the traffic transmitted per queue, and the number of packets dropped due to the queue being full. For example:

```

switch# show interface 1/1/1 queues
Interface 1/1/1 is (Administratively down)
Admin state is down
State information: admin_down

```

	Tx Packets	Tx Bytes	Tx Drops
Q0	100	8000	0
Q1	1234567	12345678908	5
Q2	0	0	0
Q3	0	0	0
Q4	0	0	0
Q5	0	0	0
Q6	0	0	0
Q7	0	0	0

- **Tx Bytes:** Total bytes transmitted. The byte count may include packet headers and internal metadata that are removed before the packet is transmitted. Packet headers added when the packet is transmitted may not be included.
- **Tx Packets:** Total packets transmitted.
- **Tx Drops:** The number of packets dropped by a queue before it was sent. When traffic cannot be transmitted out of an egress interface, it may be buffered in memory if space is available. The more servicing assigned to a queue by a schedule profile, the less likely traffic destined for that queue will back up and potentially dropped.

Data center bridging for storage and lossless Ethernet




Supported on the 8325, 9300, and 10000 Switch Series.

The term *lossless Ethernet* is intertwined with the Data Center Bridging protocol suite (DCB) to eliminate traffic loss within an Ethernet fabric. In a lossy fabric, traffic loss can occur due to queue buffer overflow. DCB standards define methods for avoiding packet drops, head of line blocking, and excessive latency, which can be applied to different priorities of traffic flowing on an interface. When these settings are applied uniformly across an Ethernet fabric, lossless Ethernet can be achieved.

DCB achieves this by managing bandwidth, priority, and flow control of selected traffic priorities when sharing the same Ethernet network link.

Data center bridging components

Data center bridging is comprised of the following key components:



IP ECN is an extension to IP supporting TCP and is covered by RFC 3168. It is a requirement to support lossless traffic flows over L3 networks.

PFC - Priority-based flow control

Priority-based Flow Control (PFC) is a link-level flow control mechanism similar to Ethernet PAUSE and standardized in IEEE 802.1Qbb.

The standard Ethernet PAUSE stops all traffic on an Ethernet link impacting all traffic flows. PFC PAUSE stops only the priorities specified by the data within the pause frame, allowing other traffic priorities to continue unaffected.

ETS - Enhanced transmission selection

Enhanced Transmission Selection (ETS) allows management of the link bandwidth by specifying bandwidth availability to each priority group. Priorities can be grouped according to bandwidth or bandwidth allocated to each individual priority.

QCN - Quantized congestion notification

Quantized Congestion Notification (QCN) provides congestion control at Layer 2. It was developed in parallel to the related effort on PFC.

QCN is designed to provide congestion awareness to an upstream link partner, allowing the link partner to reduce the traffic rate in order to avoid congestion on the downstream device.



This standard is not required if the PFC feature is available and is not supported on the CX switch platform.

DCBx - Data center bridging exchange protocol

Data Center Bridging Exchange protocol (DCBx) is a discovery and capability exchange protocol for communicating Data Center Bridging configuration information between link peers. DCBx is specified as part of IEEE 802.1Qaz-2011. DCBx uses LLDP as the underlying protocol for exchange of parameters with the peer.

The DCBx parameters are exchanged as LLDP TLVs. AOS-CX switches support both pre-standard CEE and IEEE standard DCBx.

DCBx supports VSX synchronization. For more information about enabling VSX synchronization, see the *Virtual Switching Extension (VSX) Guide* for your switch and software version.

DCBx is an extension of LLDP used by DCB devices to exchange configuration information with directly connected peers. Its applications include:

- Discovery of peer capabilities
- Detection of misconfigurations and identifying configuration mismatches
- Advising a DCB peer of which DCB configurations it should adopt

DCBx guidelines

- DCBx is disabled by default.
- LLDP must be enabled on the interfaces supporting DCBx.
- DCBx should be configured on all L2 interfaces.
- L3 hops which leverage PFC and ECN do not require DCBx, however it is recommended to enable it.
- DCBx is only supported on physical interfaces and not on management or logical interfaces, similar to how LLDP behaves.
- DCBx is not essential between host and switch but is highly desirable as to avoid manual PFC configurations of host devices.
- Supported protocols include:
 - DCBx Converged Enhanced Ethernet (CEE)
 - DCBx IEEE 802.1Qaz
- AOS-CX advertises DCBx with the willing bit set to 0 in all TLVs. This tells the peer that the switch is not willing to change its configuration to match the peer's configuration.
- When a peer device does not support the configured DCBx version (IEEE or CEE), a misconfiguration error will be displayed in the `show dcbx interface output`.

IP ECN

IP Explicit Congestion Notification (ECN) is not part of the DCB protocol standard and is covered separately by RFC 3168, however ECN plays a critical role in enhancing lossless Ethernet over a layer 3 IP fabric. IP ECN is a congestion notification protocol which throttles back traffic rather than pausing or dropping packets in the event of congestion. It is a slow acting mechanism and assists in avoiding pause frames generated by PFC by throttling back traffic before queue congestion occurs.

When devices in an IP ECN fabric experience congestion, the device marks congestion via ECN bits in the IP DiffServ field and forwards the IP packets on to either the next IP hop or to the end host destination. The receiving host is IP ECN aware and the congestion notification is handled by the upper layer protocols such as TCP. The IP ECN congestion bits are then echoed back to the transmitting host device which in turn reduces its transmission rate.

In order for IP ECN to work, the sender and receiver's upper layer protocol must be congestion-aware such as TCP or SCTP (Stream Control Transmission Protocol). In addition, all switch hops in the path between the sender and receiver are required to be ECN-aware and configured appropriately.

Host network interface cards

Network interface cards (NICs) on hosts that require lossless Ethernet can support the following protocols:

- PFC
- ETS
- DCBx*
- IP ECN



*Support for DCBx is not mandatory but highly desirable.

PFC support is a prerequisite for storage and lossless Ethernet networking. AOS-CX provides the following DCB features which are supported across specific listed CX platforms:

Feature		AOS-CX 10.10 or later
DCB Protocols	Priority-based Flow Control (PFC)	Yes (10000/9300/8400/8100/8360/8325)
	Enhanced Transmission Selection (ETS)	Yes (10000/9300/8400/8100/8360/8325)
	Data Center Bridging Exchange (DCBx)	Yes (10000/9300/8400/8100/8360/8325)
	Quantized Congestion Notification (QCN)	No
L3 Notification Protocol	IP Explicit Congestion Notification (ECN)	Yes (10000/9300/8400/8100/8360/8325)
Lossless PFC priorities per port		7 (10000/9300/8325) 2 (8100/8360) 1 (8400)
Lossless pools		3 (10000/9300/8325)




DCB protocols are only supported on the network underlay and are not supported across VXLAN overlay.



The Aruba 10000, 9300, and 8325 Switch Series support Asymmetric PFC.

DCB layer 3 configuration task list

1. Enable DCBx
2. Configure QoS pool(s)
3. Configure and apply QoS queue and schedule profiles
4. Optionally make any required configuration changes to the CoS and DSCP maps.
5. Configure global trust and override trust settings on individual interface settings
6. Enable PFC on the required priorities for each interface where it is necessary
7. Configure DCBx application TLVs
8. Configure ECN




DCBx configuration

DCBx is dependent on LLDP and uses LLDP TLVs to exchange configuration details. A DCBx-enabled device can advertise various parameters including ETS, PFC, and 802.1p CoS value.

A DCBx-enabled switch can request its peer to map certain application traffic to a priority using the DCBx TLV subtype (0x0c). Whether the host accepts the QoS parameters or rejects them is dependent on the host's local DCBx NIC willing "W" bit state:

- If set to 1 - Indicates that the adapter is willing to accept QoS parameters from the remote peer
- If set to 0 - Operational QoS parameters are always resolved from the local Host nic QoS parameters

An example of DCBx TLV advertisement where the switch advertises to the host to use TCP port 3260 for iSCSI and places it in queue 4 (802.1p CoS value 4):



DCBx configuration considerations

DCBx is expected to operate over a point-to-point link. If multiple LLDP neighbors are detected, then DCBx behaves if its peer DCBx TLVs are not present and behaves this way until the multiple LLDP neighbor condition is no longer present.

- DCBx can be enabled globally or within a specific interface and is disabled by default
- The IEEE DCBx IEEE 802.1Qaz is the default version when DCBx is enabled
- LLDP must be enabled prior to configuring DCBx
- DCBx is not essential between host and switch but is highly desirable as to avoid manual PFC configurations of host devices.
- DCBx should be configured on every switch hop interface.

Enabling DCBx

LLDP must be enabled globally prior to any DCBx configuration and can be enabled with the `lldp enable` command. DCBx can be enabled either globally or at the interface level and can specify either the IEEE or CEE version by using the `lldp dcbx` command..



The DCBx application TLV is part of the DCBx configuration but must be configured after the PFC configuration.

Enabling DCBx globally:

```
switch(config)# lldp dcbx
```

Enabling the IEEE version of DCBx globally:

```
switch(config)# lldp dcbx version ieee
```

Enabling the CEE version of DCBx globally:

```
switch(config)# lldp dcbx version cee
```

Enabling DCBx for an interface:

```
switch(config-if)# lldp dcbx
```

Enabling the IEEE version of DCBx for an interface:

```
switch(config-if)# lldp dcbx version ieee
```

Enabling the CEE version of DCBx for an interface:

```
switch(config-if)# lldp dcbx version cee
```

The `no` form of any of these commands will remove the configuration from the switch. In order to do this you must first enter the appropriate configuration context as indicated by the `switch(config)#` prompt for global configuration or `switch(config-if)#` for interface configuration. For example, removing the IEEE version of DCBx from a specific interface:

```
switch(config)# interface 1/1/1  
switch(config-if)# no lldp dcbx version ieee
```

Priority-based flow control

Priority-based Flow Control (PFC) provides an enhancement to the Ethernet flow control pause command. PFC operates at layer 2, supporting layer 2 and layer 3 network connectivity. The standard Ethernet pause frame will pause all traffic on an Ethernet link when an ingress interface buffer is full, impacting all traffic flows.

The 802.1Qbb standard for PFC allows up to 8 different traffic classes to separately be paused by a device. AOSCX switch devices support a varying number of PFC-enabled priorities on each port, depending on the model. However, a priority configured with PFC should be assigned to its own queue. Never combine more than one PFC priority in a queue because either priority being paused will result in both getting paused by the switch. The PFC receiving side sends a pause frame (XOFF) to the transmitting station when the buffer becomes full for a specific PFC-configured traffic class. This pause frame requests the transmitting station stop sending packets of that traffic class but allows all other traffic to continue.

It is recommended to configure PFC or link-level flow control RXTX mode and pools once on a running system. Changing the pool or flow control configuration on a running system will cause brief traffic disruption and packets will be lost in the data flow while the reconfiguration occurs. This may affect one, multiple, or all ports, depending on the existing configuration, and changes made to it.

Switch series	Priority-based flow control priorities
8325/9300/10000	Up to 7 PFC priorities per interface.


Switch series	Priority-based flow control priorities
8360	Up to 2 PFC priorities per interface.
8400	One PFC priority configured per interface.

PFC has the following characteristics:

- IEEE 802.1Qbb is based on Ethernet flow control
- Pause frames stop transmitters when downstream buffer is full
- Required to support lossless Ethernet (avoid dropping storage traffic)
- Provides mechanism with PAUSE per 802.1p priority
- Provides PAUSE of one or more traffic classes



Jumbo frame MTU size is required on network infrastructure to support storage networking traffic.



PFC and the QoS pool

The QoS pool option on the Aruba 8325, 9300, and 10000 Switch Series enables the creation of a packet buffer pool which is dedicated specially for lossless traffic. The `qos pool` command configures lossless pool size, headroom buffer associated with the pool, and the priorities that are mapped to that pool.

The lossless pool size is a percentage of the total available buffer memory on the device. The headroom pool memory is allocated from the lossless pool and is used for storing packets that arrive on a port after a PFC pause has been asserted.

The headroom pool is necessary because the upstream device may have transmitted frames which are in-flight on the wire when it receives the pause frame from the downstream switch. The downstream switch must have reserved buffer space for these in-flight frames in order to guarantee lossless delivery.

An example would be:

1. Packets are received on a switch interface from an upstream transmitting device where the receiving switch cannot transmit them as fast as they are arriving. These packets must be buffered by the switch in the shared lossless pool until they can be transmitted. If this condition

continues long enough, the switch interface allocation from the shared lossless pool becomes exhausted.

2. A PFC pause frame is sent from the switch interface to the upstream device.
3. Packets received by the switch after sending a pause to the upstream device are buffered in the headroom pool.

Asymmetric PFC

Symmetric PFC is when an interface has the same configuration in both directions. This is the same functionality as symmetric Ethernet flow control pause.

Asymmetric PFC allows the configuration of the pause function on an interface independently in each direction relative to the specific PFC priority. The interface can have different configurations for both the receive and transmit buffers.

- Symmetric PFC - Interface has same configuration in both directions
- Asymmetric PFC - Interface has different configuration in both directions
 - Provides the ability to separately enable taking action (stop transmit) upon receipt of a pause frame and transmitting a pause frame when the buffer limit is reached due to arriving packets.
 - Can be enabled in both directions, disabled in both directions, or enabled in one direction and disabled in the other direction

Receive (Rx) buffer	Transmit (Tx) buffer	Configured flow control state
On	Off	Interface doesn't request that the connected peer stop sending traffic. It will respond to PFC pause frames received from the connected peer but will not send pause frames upstream.
Off	On	Interface generates and sends pause frames but will not respond to received pause frames and will continue to transmit in the event that it receives one.
On	On	Interface both sends and responds to received pause frames applied to the specific configured PFC priority.

Configurable flow control buffer thresholds

The switch supports configuration of XOFF limit, XON delta and headroom limit buffer thresholds on flow control enabled interfaces. The purpose for each of these values is:

- XOFF limit is defined as the amount of shared buffer space flow-controlled packets are allowed to consume before pause is asserted to the link partner.
- XON delta is defined as the offset below XOFF limit when pause is de-asserted to the link partner.
- Headroom limit is defined as the amount of buffer space reserved to absorb any in-flight packets from the time that the XOFF message is transmitted by the local device until the time that the last packet is sent by the link partner and received by the local device.

By default when LLFC RXTX or PFC is enabled, the switch uses a dynamic buffer threshold for the XOFF limit, a delta value for XON and a headroom buffer limit based on 350m cable length. For many scenarios these defaults are sufficient to provide good performance without problems. However, in some situations it may be useful to tune the XOFF limit, XON delta and headroom parameters based on

the unique requirements of a particular deployment. In those situations, the buffer tuning capabilities described in this section can be useful to help optimize memory usage. Refer to the QoS Commands section for syntax of the CLI. QoS buffer configuration using **flow-control buffer** commands are only applicable to ports or priorities configured for flow-control. If specified for a port or priority not configured for flow-control, the configuration has no effect. QoS buffer configuration for PFC priorities are not applicable to LLFC RXTX ports and vice versa.

XOFF thresholds

The lossless buffer limits for traffic arriving on flow-control enabled ports or priorities can be optionally configured using one of two modes, dynamic or static. If no user-configured limit is specified, the default limit of dynamic alpha 2 is configured.

Dynamic XOFF thresholds

The limit scales as a function of available buffer resources and is configured by specifying an alpha value in the range between -7 and 3. The dynamic limit is based on the currently-available free space, multiplied by 2 raised to the power of the specified alpha value. A value less than zero means that the limit is a fraction of the pool free space, e.g. -1 corresponds to half. A value greater than zero means that the limit is a multiple of the free-space size, e.g. 1 corresponds to double.

Static XOFF thresholds

A static limit sets a fixed size for the XOFF threshold.

XON delta

The value is relative to the XOFF limit and applies to both dynamic and static XOFF limit configurations. When the buffer in-use amount from a source port (LLFC RXTX) or source port + priority (PFC) reaches the XOFF limit, pause is asserted to the link peer. At this point, pause will remain asserted until the shared buffer in-use amount by this port/priority decreases by the XON Delta amount. Once enough packets are transmitted to reduce In-Use by this amount, pause is de-asserted. It is important to note that the XON Delta value will impact the amount of flow-control pause frames sent to the link peer and can have a negative impact on performance. A value that is too small can result in many more pause frames sent to the link peer as the buffer in-use amount reaches the XOFF limit, drops slightly and meets the XON delta, then immediately jumps back to the XOFF limit. Conversely, a value that is too large can result in poor link bandwidth utilization due to insufficient buffered packets as the buffer drains too far before more packets arrive to utilize the link. The default XON Delta value is 12288 bytes.

Headroom size

This value can be explicitly configured by setting the buffer size or implicitly configured by setting the cable length. If both are specified, the buffer size value is used and the cable length is ignored.

Headroom limit

Headroom Limit for a LLFC RXTX port or PFC priority, sets a fixed size of headroom pool buffer space for in-flight packets that arrive after XOFF is asserted.

Cable length value

Indirectly adjusts the headroom buffer size on the physical interface. Cable length is one of the delay components that effects traffic round trip time (RTT). Based on IEEE std 802.1Qbb, RTT depends on several delay components such as PFC transmission delay, Interface Delay, Cable Delay and Higher layer delay. Cable length configuration changes the cable delay component of RTT and is used to compute the headroom buffer size for in-flight packets that arrive after XOFF is asserted. If both headroom buffer limit and cable length are configured, the headroom buffer limit is used and cable length is ignored.

Cable length auto

Auto mode will compute the headroom size using the length of the attached cable when known. If the cable length cannot be determined, headroom is computed using the configured cable length or the default length of 350 meters.

Cable length default

If none of the commands that adjust the headroom limit are configured, a default cable-length of 350 meters is used to compute the headroom pool limit.

Considerations and prerequisites

PFC configuration is dependent on the successful configuration of the appropriate QoS pool with the priorities required for the PFC interface configuration. PFC cannot be enabled on an interface if the PFC priority is not first mapped to a lossless pool. The exception is when RX-only flow control is allowed to be configured without any pools. This is to support semi-lossless behavior where the switch can be paused by the downstream device (LLFC or PFC depending on the config), but the switch will eventually start dropping packets if the queue (lossy buffer) becomes full. The number of unique priority-based flow control (PFC) combinations across all switch interfaces cannot be more than 3 for the 8325/10000 Switch Series and cannot be more than 7 for the 9300 Switch Series.

An example of QoS pool configuration:

```
switch(config)# qos pool 1 lossless size 30 percent headroom 2048 kbytes
priorities 1,2
switch(config)# qos pool 2 lossless size 30 percent headroom 2048 kbytes
priorities 3,4
switch(config)# qos pool 3 lossless size 30 percent headroom 2048 kbytes
priorities 5,6,7
```



PFC is required to be configured on all interfaces between the host initiator and the host target for both L2 and L3 connections.

QoS queue profile

The QoS queue profile maps local priorities to queues. The CoS-map defines the 802.1p priority code mappings to the local priorities. The relationship to the QoS queue profile for lossless Ethernet is the CoS map and the DSCP map. The factory defaults are applied for both maps at system start up.

Default CoS map

```
switch# show qos cos-map default
code_point local_priority color name
-----
0          1             green Best_Effort
1          0             green Background
2          2             green Excellent_Effort
3          3             green Critical_Applications
4          4             green Video
5          5             green Voice
6          6             green Internetwork_Control
7          7             green Network_Control
```



```
switch# show qos dscp-map
```

DSCP	code_point	local_priority	mpls_exp	color	name
000000	0	1	0	green	CS0
Codepoints 1-7 removed for brevity					
001000	8	0	1	green	CS1
Codepoints 9-15 removed for brevity					
010000	16	2	2	green	CS2
Codepoints 17-23 removed for brevity					
011000	24	3	3	green	CS3
Codepoints 25-31 removed for brevity					
100000	32	4	4	green	CS4
Codepoints 33-39 removed for brevity					
101000	40	5	5	green	CS5
Codepoints 40-47 removed for brevity					
110000	48	6	6	green	CS6
Codepoints 48-55 removed for brevity					
111000	56	7	7	green	CS7
Codepoints 57-63 removed for brevity					


Factory default QoS queue profile

A single factory-default queue profile is initially applied to every port:

```
switch# show qos queue-profile factory-default
```

queue_num	local_priorities	name
0	0	Scavenger_and_backup_data
1	1	
2	2	
3	3	
4	4	
5	5	
6	6	
7	7	

The default queue-profile has each local-priority mapped to a separate queue:



Best effort traffic will use Queue 1 using the default cos-map and qos-queue profile.

- Code point 0 (CoS value 0) is mapped to local priority 1 in the default cos-map for best effort traffic
- Code point 1 (CoS value 1) is mapped to local priority 0 in the default cos-map for less than best effort traffic
- Local priority 1 is mapped to queue 1 in the qos queue-profile factory-default
- The default CoS map and DSCP map can be leveraged noting the CoS value 0 (typically best effort traffic) will be mapped to queue 1 when using the default maps

A QoS queue profile is always applied to the switch and is necessary to specify internal priority to queue mapping. At all times either the factory-default queue profile or a user-specified queue profile is applied to the switch.

QoS queue profile configuration

The preferred QoS queue profile configuration will depend on the number PFC priorities as part of the lossless traffic class and the overall QoS requirements for all traffic classes.

- A local priority can only be associated with a single queue
- Any number of local priorities can be configured to a single queue
- It is essential that local priority 7 is given a dedicated queue with a minimum bandwidth for infrastructure control traffic to avoid buffer starvation from other priority queues



It is not recommended to assign multiple local priorities to the same queue if those local priorities are used for lossless traffic.

QoS queue profile configuration involves the following steps:

1. Identify local priorities for mapping to the QoS queue profile
2. Map local priorities to the appropriate queue

Create a QoS queue profile called `queue-profile-1`:

```
switch(config)# qos queue-profile queue-profile-1
```

Map queues in `queue-profile-1` to local priorities:

```
switch(config)# qos queue-profile queue-profile-1
switch(config-queue)# map queue 0 local-priority 0
switch(config-queue)# map queue 1 local-priority 1,2,5,6
switch(config-queue)# map queue 2 local-priority 3
switch(config-queue)# map queue 3 local-priority 4
switch(config-queue)# map queue 4 local-priority 7
```

The `no` form of the `qos queue-profile` command will remove the QoS queue profile:

```
switch(config)# no qos queue-profile queue-profile-1
```

QoS schedule profile - enhanced transmission selection

AOS-CX leverages the DWRR algorithm to provide different weights to each QoS queue which will support a single or a number of traffic classes in a profile. The profile is called the schedule-profile and by default all interfaces will have the default-schedule profile applied.



The configuration examples below apply to the 8100, 8325, 8360, 9300, and 10000 Switch Series. Configuration on the 8400 Switch Series is similar with the only difference being it uses the WFQ algorithm instead of DWRR.

The factory default schedule profile :

```
switch# show qos schedule-profile factory-default
queue_num algorithm      weight max-bandwidth_kbps burst_KB
-----
0          dwrr          1
1          dwrr          1
2          dwrr          1
3          dwrr          1
4          dwrr          1
5          dwrr          1
6          dwrr          1
7          dwrr          1
```

DWRR Calculation

To calculate the bandwidth relative to the DWRR weight, the queue weight is divisible by the sum of the weights for all queues. The default schedule profile has 8 queues with each queue having a weight of 1. To calculate the bandwidth used for each queue, the individual queue weight is divided by the sum of weights of all queues. An example using the factory default schedule profile which represents the minimum bandwidth for each queue:

- Queue 0 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 1 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 2 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 3 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 4 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 5 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 6 dwrr weight $1 - 1(8) = 12.5\%$
- Queue 7 dwrr weight $1 - 1(8) = 12.5\%$



The total percentage should add up to 100% for all weights.

For lossless Ethernet traffic flows, a custom schedule profile must provide the appropriate amount of bandwidth to the lossless queues as well as the lossy queues configured on the switch system. The queue-profile and the schedule-profile must specify the same queues. If the custom schedule-profile uses the same queues as the factory-default queue-profile, then no custom queue-profile is required.

QoS schedule profile configuration

The following considerations should be taken into account when configuring a QoS schedule profile:

- Assign DWRR weights based on lossless flows, traffic flow volume, and a consideration to other traffic classes.

- Ensure that the highest numbered queue specified in the applied queue-profile (assigned to critical networking traffic) has a minimum amount of bandwidth to avoid queue starvation.
- Number of queues
 - On the 9300 Switch Series the number of queues in the queue profile must align to the number of queues with allocated weights in the schedule profile.
 - On the 8325/10000 Switch Series it is possible to configure any of the 8 queues (0-7). The number of queues as well as the queue numbers themselves must be the same between the applied queue profile and schedule profile.
- It is recommended to use weights totaling 100 for easy reference to bandwidth percentage reservations.

QoS schedule profile configuration involves the following steps:

1. Configure QoS queue profile
2. Create the schedule profile
3. Apply the QoS queue profile and schedule profile

Creating a QoS schedule profile:

```
switch (config)# qos schedule-profile schedule-profile1
```

The below example shows an ETS configuration within a 2-queue environment. The settings use a weight to set the amount of available bandwidth for each queue. The settings in the example below would ensure that 50% of bandwidth will be applied to both queue 0 and queue 1:

```
switch(config)# qos schedule-profile schedule-profile1
switch(config-schedule)# dwrr queue 0 weight 15
switch(config-schedule)# dwrr queue 1 weight 15
```

Applying the schedule profile and queue profile:

```
switch (config)# apply qos queue-profile profile1 schedule-profile schedule-profile1
```

The global configuration will be applied to all interfaces.

Overriding the global schedule profile on an interface

A different schedule profile can be applied directly to interfaces which will override the global schedule profile. The caveat is that the schedule profile queues must align to the same queues as defined in the applied QoS queue profile. The procedure for applying a different schedule profile to a specific interface is as follows:

1. Create a schedule profile that specifies the same queues used in the globally-applied queue profile
2. Customize the scheduling algorithm and weight (depending on algorithm) for each queue
3. Apply the override schedule profile to the desired interface(s)

Creating the interface override schedule profile called `schedule-profile2`:

```

switch (config)# qos schedule-profile schedule-profile2
switch (config-schedule)# dwrr queue 0 weight 5
switch (config-schedule)# dwrr queue 1 weight 50
switch (config-schedule)# dwrr queue 2 weight 20
switch (config-schedule)# dwrr queue 3 weight 15
switch (config-schedule)# dwrr queue 4 weight 10

```

Applying schedule profile `schedule-profile2` mapped with queue profile `profile1` to the global config:

```

switch(config)# apply qos queue-profile profile1 schedule-profile schedule-profile2

```

Verifying that that the interface override has occurred:

```

switch# show qos schedule-profile
profile_status profile_name
-----
complete      factory-default
applied       schedule-profile2
complete      strict
incomplete     test1

```

Validating that schedule profile `schedule-profile2` is applied to an interface:

```

switch# show interface 1/1/1 qos
Interface 1/1/1 is up
Admin state is up
qos trust none (global)
qos queue-profile profile1 (global)
qos schedule-profile schedule-profile2 (global)

```

Creating the schedule profile called `profile1` that will become the global schedule profile:

```

switch (config)# qos schedule-profile schedule-profile1
switch (config-schedule)# dwrr queue 0 weight 5
switch (config-schedule)# dwrr queue 1 weight 25
switch (config-schedule)# dwrr queue 2 weight 35
switch (config-schedule)# dwrr queue 3 weight 25
switch (config-schedule)# dwrr queue 4 weight 10

```

Applying schedule profile `profile1` to queue profile `profile1`:

```

switch (config)# apply qos queue-profile profile1 schedule-profile schedule-profile1

```

Validating that the schedule profile `profile1` is applied globally and has overwritten the global setting of schedule profile `profile2`:

```

switch# show qos schedule-profile
profile_status profile_name

```

```

-----
complete      factory-default
applied      schedule-profile1
complete      schedule-profile2
complete      strict

```

Validating the global the schedule profile `profile1` has been applied to all interfaces:

```

switch# show interface qos
Interface 1/1/1 is up
  Admin state is up
  qos trust none (global)
  qos queue-profile profile1 (global)
  qos schedule-profile schedule-profile1 (global)

output removed for brevity

Interface 1/1/56 is up
  Admin state is up
  qos trust none (global)
  qos queue-profile profile1 (global)
  qos schedule-profile schedule-profile1 (global)

```

Applying schedule profile `profile2` to a specific interface to override the global schedule profile:

```

switch (config)# interface 1/1/1
switch (config-if)# apply qos schedule-profile schedule-profile2

```

Verifying that both schedule profiles are applied:

```

switch# show qos schedule-profile
profile_status  profile_name
-----
complete      factory-default
applied      schedule-profile1
applied      schedule-profile2
complete      strict

```

Verifying that schedule profile `profile2` has overridden `profile1` on interface 1/1/1 and `profile1` has been applied to all other interfaces:

```

switch# show interface qos
Interface 1/1/1 is up
  Admin state is up
  qos trust none (global)
  qos queue-profile profile1 (global)
  qos schedule-profile schedule-profile2 (override)

interfaces removed for brevity

Interface 1/1/56 is up
  Admin state is up
  qos trust none (global)
  qos queue-profile profile1 (global)
  qos schedule-profile schedule-profile1 (global)

```

Schedule profile `profile2` has been applied only to interface 1/1/1. All other interfaces retain the global schedule profile setting of `profile1`.

IP explicit congestion notification

TCP/IP networks signal congestion by dropping packets. When packets are dropped due to congestion, TCP/IP easily detects this but performance suffers. For lossless Ethernet fabric, dropping packets is not allowed, so a different action is required to inform endpoints of congestion. Instead of dropping packets when congestion is experienced, ECN marks the IP packet as congestion experienced (CE).

ECN is used in conjunction with DCBx and used to avoid hitting the PFC limits over L3 network hops. The ECN bits of a packet are conditionally marked by a switch based on a configurable ECN threshold which allows TCP to function normally without dropping IP packets. ECN has the following characteristics:

- Uses the DiffServ field in the IP header to mark the congestion status along the packet transmission path.
- End-to-end communication protocol which requires support by each end device as well as all intermediate switch hops.
- Applied on a queue basis and can be applied to 1 or more queues concurrently up to the 8 queue limit.

ECN uses the two "least significant" bits in the Traffic Class (DiffServ) field of the IP header (the bits that are farthest to the right) to signal the current state of network congestion. The possible markings are:


Bits	Congestion status	Abbreviation
00	Non ECN-Capable Transport	Non-ECT
01	ECN Capable Transport	ECT(1)
10	ECN Capable Transport	ECT(0)
11	Congestion Encountered	CE



ECT(1) and ECT(0) are considered to be functionally the same by the switch. If congestion is encountered, the zero bit will be flipped to indicate congestion.

The benefit of ECN is that it reduces the likelihood of dropping packets in a lossy fabric or reduces the usage of flow-control pause in a lossless fabric. ECN accomplishes this via the following actions:

1. IP packets flow normally on an IP network from an ECN-enabled device with the ECN field marked ECT(0) or ECT(1) to indicate support for ECN
2. If an ECN-enabled host device receives a packet marked CE; it acknowledges the marked packet with its upper layer protocols (TCP)
3. The host receiving the CE marked packet sends the congestion notification back to the transmitting host
4. The transmitting host device handles ECE marking in the received acknowledgment frame by reducing its transmit rate
5. The reverse congestion notification occurs in the TCP header, not the IP header, using the ECE flag of the acknowledgment frame.



IP ECN configuration considerations



Only supported by the Aruba 8325, 9300, and 10000 Switch Series.

The Aruba 8325, 9300, and 10000 Switch Series have a dual threshold profile with a minimum and maximum threshold in bytes. Between the minimum and maximum thresholds there is a probability curve. As the buffer utilization increases past the minimum threshold, there is a higher probability of packets being marked by ECN due to congestion. When buffer utilization reaches the maximum threshold all packets are marked for ECN congestion.

An ECN threshold profile can be applied either globally or at the interface level to override the global threshold setting. Only one ECN threshold profile can be applied to a port.

The following factors should always be taken into account when configuring IP ECN:

- ECN is configured on queues carrying delay-sensitive traffic.
- Any number of queues maybe configured in a profile.
- An unconfigured queue means no threshold action behavior will occur.
- ECN can be applied globally or on individual interfaces.
- No threshold policies are applied by default.

- All switches in the transmission path between two ECN-enabled endpoints must have ECN enabled and configured for the appropriate queue.
- General guidance is to allocate 512 kilobytes per flow (.5MB) to derive the minimum threshold and add between 512 kilobytes and 1024 kilobytes (1MB) to the total for the maximum threshold. An example when configuring 4 ingress flows:
 - Minimum threshold = 2048 kilobytes (2MB)
 - Maximum threshold = 2560 kilobytes (2.5MB)
- Testing is recommended to ensure ECN thresholds are correctly applied.

IP ECN configuration procedure

The procedure for configuring IP ECN is as follows:

- Create a threshold profile using the [qos threshold-profile](#) command and apply threshold values
- Optionally apply a QoS threshold profile globally
- Optionally apply a QoS threshold profile to the desired interface(s) to override the global threshold profile setting

Creating a QoS threshold profile called **ECN** applied globally to all interfaces:

```
switch(config)# qos threshold-profile ECN
switch(config-threshold)# queue 1 action ecn all min-threshold 512 kbytes max-
threshold 1520 kbytes
switch(config)# apply qos threshold-profile ECN
```

Creating a QoS threshold profile called **ECN2** :

```
switch(config)# qos threshold-profile ECN2
switch(config-threshold)# queue 1 action ecn all min-threshold 3072 kbytes max-
threshold 4096 kbytes
```

Applying QoS threshold profile **ECN2** to a specific interface:

```
switch(config)# interface 1/1/1
switch(config-if)# apply qos threshold-profile ECN2
```

Lossless QoS pool

The QoS pool is a packet buffer pool. The Aruba 8325, 9300, and 10000 Switch Series have a configurable QoS pool with an associated headroom buffer space, while the 8360 has a single fixed QoS pool which is not exposed to the CLI for configuration. The 8400 Switch Series assigns reserved buffer space per-interface per-PFC-priority, there is no fixed lossless pool.

Switch series	Lossless QoS pools
8325/9300/10000	Three configurable lossless pools
8360/8400	Single fixed lossless pool

The QoS pool is a packet buffer pool resident in memory in the switch. Only the QoS pools on the Aruba 8325, 9300, and 10000 Switch Series are configurable and have the following attributes:

- The QoS pool option enables the creation of a packet buffer pool which is dedicated to lossless traffic
- The QoS pool command allows the creation of a lossless pool size, headroom buffer associated with the pool, and the priorities that are mapped to that pool.
- The lossless pool size is a percentage of the total available buffer memory on the device.
- The headroom pool memory is allocated from the lossless pool and is used for storing packets that arrive on a port after a pause has been asserted.

Considerations and prerequisites

There are no prerequisite configuration activities required to configure QoS pools. However, a PFC priority or any related PFC commands cannot be enabled on an interface until the creation of the QoS pool and the appropriate PFC priority mapping to that pool. The following considerations should be taken into account when configuring QoS pools:

- Maximum buffer size in memory
 - 8325/10000: 32Mb
 - 9300: 66 MB
- Maximum configurable headroom memory size: 10240 kbytes (10Mb)
- Factory-default headroom size: 3072 kbytes (3Mb)
- The sum of pool sizes cannot be larger than 90% of the available Kbyte memory
- The headroom size cannot be greater than half of the lossless pool buffer size.



Best practice for QoS pools on all platforms is to not assign the best effort traffic class (default local priority 1) and the critical control traffic (local priority 7) to a lossless pool.

QoS pool configuration

QoS pools are enabled using the QoS pool command in the global configuration context.



For additional information on QoS pool configuration please refer to the [qos pool](#) command.



The following examples are applicable to the 8325, 9300, and 10000 Switch Series.

Configuration options for QoS pool size:

```
switch(config)# qos pool (1-3) lossless size ?
<10-90>          The percentage of packet buffer memory to be allocated to
                  this pool in integer format
<10.00-90.00>    The percentage of packet buffer memory to be allocated to
                  this pool in decimal format
factory-default  Configure the pool size to the default value
```

Configuration options for headroom size:

```
switch(config)# qos pool 2 lossless size factory-default percent headroom ?  
<0-10240>      The headroom buffer size in kilobytes  
factory-default Configure the default headroom buffer size in kilobytes
```

Configuring a QoS pool size of 60% of the total available buffer memory and a headroom size of 2048 kB with priorities 1, 2, and 3 mapped to the QoS pool:

```
switch(config)# qos pool 2 lossless size 60 percent headroom 2048 kbytes  
priorities 1,2,3
```

The `no qos pool` command will remove the configuration from the switch. For example, removing QoS pool 2:

```
switch (config)# no qos pool 2
```

PFC flow-control priorities must be removed on all interfaces prior to removing the QoS pool otherwise a warning message will be displayed and the QoS pool will not be removed from the configuration. The warning message reads as follows:

```
'flow-control priority' must be unconfigured for all interfaces before deleting  
pool 2.
```

Use the `no flow-control priority` command within the appropriate interface context:

```
switch(config-if)# no flow-control priority  
<0-7>      The packet priority to flow control  
rx         Honor received PAUSE requests for this priority  
rxtx       Send and receive PAUSE requests for this priority  
tx         Send PAUSE requests for this priority
```


Reapply the `no qos pool` command:

```
switch (config)# no qos pool 2
```

QoS trust

QoS trust can be applied either globally or at the interface level which will override the global setting. The following considerations should be taken into account when configuring QoS trust:

- L2 communications between the switch and the host must leverage the 802.1Q tag meaning the interface between the switch and host must be a trunk link.
- L2 switch fabric 802.1Q tags must be configured between target and initiator across the layer 2 fabric.
- L3 switch fabric must sync the appropriate CoS map to local preferences with the DCSP map local preferences. This is to ensure appropriate mapping for CoS and DSCP markings across the network.




Set a global trust setting or leave global QoS trust the default setting. Override the global trust setting on specific interfaces as desired.



QoS trust is set to none by default.

Layer 3 CoS-DSCP markings

A remark policy may be advantageous in order to ensure that the local priority derived from the "trusted" CoS value is remarked to the desired DSCP code point for transmission across the IP fabric. The diagram illustrates an example of layer 3 switch configuration with remark policy:



Configuration steps:

1. Switch A

- A classifier is applied to match traffic with a priority code point (PCP) of 4
- A policy is applied to remark the classified traffic to the desired DSCP code point (CS4)
- The policy is applied to the switch egress interface

2. Switch C

- A classifier is applied to match traffic with a DSCP code point of CS4
- A policy is applied to remark the classified traffic to the local-priority of 4 (CoS 4)
- The policy is applied to the ingress interface (facing the initiator host)

3. Switch C

- A classifier is applied to match traffic with a PCP of 4
- A policy is applied to remark the classified traffic to the desired DSCP code point (CS4 in this example)
- The policy is applied to the switch egress interface

4. Switch A

- A classifier is applied to match traffic with a DSCP code point of CS4
- A policy is applied to remark the classified traffic to the local-priority of 4 (CoS 4)
- The policy is applied to the routed ingress interface (facing the target host)



The policies applied ensure that the DSCP value is consistent across the L3 network. At the final hop, the DSCP value is mapped to the correct CoS value for transmission to the target host. The policy is applied to the reverse flow.

Verifying trust settings

Global trust verification


```
switch# show qos trust
qos trust cos
```

Interface trust verification

```
switch# show interface 1/1/1 qos
Interface 1/1/1 is up
Admin state is up
qos trust cos (override)
qos queue-profile factory-default (global)
qos schedule-profile factory-default (global)
```

Troubleshooting data center bridging

It is assumed that prior to engaging in troubleshooting that host targets and initiators have network connectivity, are using DCBx, and that hosts are supporting appropriate DCB protocols like PFC. Checks should be applied to every switch hop. The following flow chart represents the Aruba-recommended data center bridging troubleshooting flow supporting common DCB implementation:



1. Verify that host target/initiator configurations are using 801.Q tagged frames to support the 802.1p field.

Check that both host and initiator are using 802.1Q tagged frames between host and switch. This requires the switch interface to be configured as a trunk interface supporting the appropriate VLANs. Check that the appropriate `qos trust` setting is applied to the interface. This can either globally or via the interface context with the `qos trust` and `show interface 1/1/1 qos` commands. The `qos trust` setting should be set to `qos trust cos` if the host is supporting PFC as well as DCBx and is configured to use the appropriate CoS setting for a specific lossless traffic flow.

2. Validate DCBx application and QoS Pools



Only applicable to the Aruba 8325, 9300, and 10000 Switch Series

Validate that DCBx is configured on the correct interfaces, that the adjacent neighbor is using the correct DCBx version, and using the same PFC priority. If DCBx is configured and the PFC priority matches at either end for the link, the following output is displayed:

```

switch# show dcbx interface 1/1/1
DCBx admin state      : enabled
DCBx operational state : active
DCBx version          : local = ieee, remote = ieee
PFC operational state  : active

Mismatch  Advertisement      Local    Peer
-----  -
Priority Flow Control (PFC)
Willing:                No        No
MACsec Bypass Capability: No        No
Max PFC Traffic Classes: 2          2
Priority 0:              False     False
  
```

Priority 1:	False	False
Priority 2:	False	False
Priority 3:	False	False
Priority 4:	True	True
Priority 5:	False	False
Priority 6:	False	False
Priority 7:	False	False

Output omitted for brevity

```
switch(config)# show qos pool config
```

Global Packet-Buffer Configuration and Status (in Kbytes)

Configured Priority Packet Buffer Pools Size Mapping	Sum of Limits	Applied Port Mapping	Priority

--			
Lossy Pool 9546	--	0,1,2,5,6,7	
0,1,2,5,6,7			
Lossless Pool 1 7546 (30.00%)	--	4	
4			
Headroom Pool 1 2000	0	--	
--			
Lossless Pool 2 9728 (40.00%)	--	3	
3			
Headroom Pool 2 3000	0	--	
--			
Reserved Memory 948	--	--	
--			

--			
TOTAL:	32768		

3. Verify qos-queue profile and qos schedule-profile

Verify that the `qos-queue profile` and the global `qos schedule-profile` have been applied. A single qos queue-profile is applied to the system. A schedule is always applied to all interfaces (either the globally configured schedule profile, an interface-override schedule profile, or the factory-default schedule-profile if the configuration cannot be applied). Each port can have its own schedule profile. Verify that the correct QoS queue profile is applied. An example of the QoS queue profile **4qp** applied to the system:

```
switch# show qos queue-profile
profile_status profile_name
-----
complete      3qp
applied       4qp
complete      SMB
complete      factory-default
```

Verify that the correct QoS schedule profile is applied to an interface. An example of the QoS schedule profile **4sp** applied to the system:

```
switch# show interface 1/1/1 qos  
Interface 1/1/1 is up  
Admin state is up  
qos trust cos (override)  
qos queue-profile 4qp (global)  
qos schedule-profile 4sp (global)
```

4. Check DCBx TLV and PFC Configurations

Verify that the correct DCBx application TLV is configured and has the correct priority using the `show dcbx interface` command.

Verify that appropriate interfaces have PFC configured:

```
switch# show interface 1/1/2  
interface 1/1/2  
  no shutdown  
  flow-control priority rxtx 4  
  no routing  
  vlan trunk native 1  
  vlan trunk allowed 1,110  
  qos trust cos
```

Output omitted for brevity

apply qos

```
apply qos [queue-profile <QUEUE-NAME>] schedule-profile <SCHEDULE-NAME>
no apply qos [queue-profile <QUEUE-NAME>] schedule-profile <SCHEDULE-NAME>
```

Description

Applies a queue profile and schedule profile globally to all Ethernet and LAG interfaces on the switch, or applies a schedule profile to a specific interface. When applied globally, the specified schedule profile is configured only on Ethernet interfaces and LAGs that do not already have their own schedule profile.

The same profile can be applied both globally and locally to an interface. This guarantees that an interface always uses the specified profile, even if the global profile is changed.

The **no** form of this command removes the specified schedule profile from an interface and the interface uses the global schedule profile. This is the only way to remove a schedule profile override from the interface.



When applying QoS settings to a port configured to support priority-based flow control, specific configuration settings must be respected when defining a CoS map and queue profile. See the command **flow-control** in the *Command-Line Interface Guide* for details.



Interfaces may shut down briefly during reconfiguration.

Parameter	Description
<code>queue-profile <QUEUE-NAME></code>	Specifies the name of the queue profile to apply. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-). This parameter is not supported in the config-if context.
<code>schedule-profile <SCHEDULE-NAME></code>	Specifies the name of the schedule profile to apply. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Usage

- The switch must always have a globally-applied queue and schedule profile. To stop using a given profile, apply a different profile.
- For a queue profile to be complete and ready to be applied, all eight local priorities must be mapped to a queue.

- For a schedule profile to be complete and ready to be applied, it must define all queues specified in the queue profile. All queues must use the same algorithm, except for the highest numbered queue, which can be **strict**.
- Both the queue profile and the schedule profile must specify the same number of queues.
- Schedule profiles can be modified while applied, but only in ways where a single command will not result in the profile becoming invalid. For example, queue 7 can have the algorithm changed, and weighted queues can have their weights changed.

If there are interfaces running with priority-based flow control (PFC) and a new queue profile to be applied maps a local-priority used by the PFC traffic to another queue, the following steps are required:

1. Shutdown all interfaces
2. Apply the new queue profile
3. Save the configuration to startup
4. Reboot the switch

If the new queue profile was applied before shutting down the PFC interfaces then PFC traffic will still use the same queue from the previous profile until the switch is rebooted.

If the number of queues was changed from the previous queue profile to the new one, any Ethernet or LAG interfaces with locally applied schedule profiles will program the newly applied global schedule-profile. The *show running-config interface* command will list the existing *apply qos schedule-profile* command with a comment describing the actual profile applied:

```
apply qos schedule-profile Old_Schedule
!actual schedule-profile New_Schedule
```

Examples

The following commands illustrate a valid configuration where every local priority value is assigned to a queue and all assigned queues are defined:

```
switch(config)# qos cos-map 1 local-priority 1
switch(config)# qos queue-profile Q1
switch(config)# map queue 0 local-priority 0
switch(config)# map queue 1 local-priority 1
switch(config)# map queue 2 local-priority 2
switch(config)# map queue 3 local-priority 3
switch(config)# map queue 4 local-priority 4
switch(config)# map queue 5 local-priority 5
switch(config)# map queue 6 local-priority 6
switch(config)# map queue 7 local-priority 7
switch(config)# qos schedule-profile S1
switch(config)# dwrr queue 0 weight 5
switch(config)# dwrr queue 1 weight 10
switch(config)# dwrr queue 2 weight 15
switch(config)# dwrr queue 3 weight 20
switch(config)# dwrr queue 4 weight 25
switch(config)# dwrr queue 5 weight 50
```

The following commands illustrate an invalid configuration because local priority 2 is not assigned to a queue:

```

switch(config)# qos cos-map 1 local-priority 1
switch(config)# qos queue-profile Q1
switch(config)# map queue 0 local-priority 0
switch(config)# map queue 1 local-priority 1
switch(config)# map queue 3 local-priority 3
switch(config)# map queue 4 local-priority 4
switch(config)# map queue 5 local-priority 5
switch(config)# map queue 5 local-priority 6
switch(config)# map queue 5 local-priority 7
switch(config)# qos schedule-profile S1
switch(config)# dwrr queue 0 weight 5
switch(config)# dwrr queue 1 weight 10
switch(config)# dwrr queue 3 weight 15
switch(config)# dwrr queue 4 weight 25
switch(config)# dwrr queue 5 weight 50

```

Applying the QoS profile **Q1** and the schedule profile **S1** to all interfaces that do not have an applied interface-specific schedule profile:

```

switch(config)# apply qos queue-profile Q1 schedule-profile S1

```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config config-if config-lag-if	Administrators or local user group members with execution rights for this command.

apply qos threshold-profile



Only supported on the 8325, 9300, and 10000 Switch Series.

```

apply qos threshold-profile <THRESHOLD-NAME>
no apply qos threshold-profile

```

Description

Applies a threshold profile globally to all Ethernet and LAG interfaces on the switch, or to a specific interface. When applied globally, the specified threshold profile is configured only on Ethernet interfaces and LAGs that do not already have their own schedule profile.

The same profile can be applied both globally and locally to an interface. This guarantees that an interface always uses the specified threshold profile, even if the global profile is changed.

The **no** form of this command removes the specified threshold profile from an interface, and causes it to use the global threshold profile. This is the only way to remove a threshold profile override from an interface. A profile can only be deleted once it is no longer applied to any interface.

Parameter	Description
<code><THRESHOLD-NAME></code>	Specifies the name of the threshold profile to apply. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Example

Applying the threshold profile **mythreshold** to all interfaces that do not have an applied profile:

```
switch(config)# apply qos threshold-profile mythreshold
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8325	config	Administrators or local user group members with execution rights for this command.
9300	config-if	
10000	config-lag-if	

dcbx application

```
dcbx application {iscsi | tcp-sctp <PORT-NUM> | tcp-sctp-udp <PORT-NUM> | tcp-udp <PORT-NUM>
| udp <PORT-NUM> | ether <ETHERTYPE>} priority <PRIORITY>
no dcbx application
```

Description

Configures application to priority map that gets advertised in DCBx application priority messages. This tells the DCBx peer to send the application traffic with the configured priority so that the traffic is treated as lossless. Multiple applications can be configured in this manner. PFC lossless priority configured on the switch should be the same as this priority.

The **no** form of this command removes the existing configuration.

Parameter	Description
<code>iscsi</code>	Specifies a physical port on the switch. TCP ports 860 and 3260.
<code>tcp-sctp <PORT-NUM></code>	Specifies the traffic for a specified TCP or SCTP port. Range: 1 to 65535.
<code>tcp-sctp-udp <PORT-NUM></code>	Specifies the traffic for a specified TCP or SCTP or UDP port. Range: 1 to 65535.

Parameter	Description
<code>tcp-udp <PORT-NUM></code>	Specifies the traffic for a specified TCP or UDP port. Range: 1 to 65535.
<code>udp <PORT-NUM></code>	Specifies the traffic for a specified UDP port. Range: 1 to 65535.
<code><ETHERTYPE></code>	Specifies the traffic for a specific Ethernet type. Range: 1536 to 65535.
<code><PRIORITY></code>	Specifies the application priority. Range: 0 to 7.

Usage

- In CEE DCBx version, the following traffic type configurations are sent using application TLVs:
 - Ethertype
 - iSCSI
 - TCP-UDP
- In IEEE DCBx version, the following traffic type configurations are sent using application TLVs:
 - Ethertype
 - iSCSI
 - TCP-SCTP
 - TCP-SCTP-UDP
 - UDP

Examples

Mapping iSCSI traffic to priority 5.

```
switch(config)# dcbx application iscsi priority 5
```

Mapping TCP or SCTP traffic with port 860 to priority 3.

```
switch(config)# dcbx application tcp-sctp 860 priority 3
```

Command History

Release	Modification
10.08	Added a parameter option tcp-udp .
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325	<code>config</code>	Administrators or local user group members with execution rights for this command.

Platforms	Command context	Authority
9300 10000		

dwrr queue

```
dwrr queue <QUEUE-NUMBER> weight <WEIGHT> [max-bandwidth <RATE> {kbps|percent} [burst
<SIZE>]]
no dwrr queue <QUEUE-NUMBER> [weight <WEIGHT> [max-bandwidth <RATE> {kbps|percent} [burst
<SIZE>]]
```

Description

Assigns the deficit weighted round robin (DWRR) algorithm and its weight to a queue in a schedule profile. DWRR allocates available bandwidth among all non-empty queues in relation to the queue weights.

Egress queue shaping can be configured using the **max-bandwidth** option to limit the amount of traffic transmitted per output queue at all times, even when there is leftover bandwidth available on the port. The buffer associated with each egress queue stores the excess traffic to smooth the output rate. Sustained rates of traffic above the maximum bandwidth will eventually fill the output queue, causing tail drops. Use **show interface <IF-NAME> queues** to determine if any tail-drop errors have occurred.

Use **show qos schedule-profile <NAME>** to view the settings of a specific schedule profile.

To remove only egress queue shaping, re-enter the dwrr queue command without the **max-bandwidth** parameter.

The **no** form of this command removes the DWRR algorithm from a queue in a schedule profile.

Parameter	Description
<QUEUE-NUMBER>	Specifies the queue number. Range: 0 to 7.
weight <WEIGHT>	Specifies the scheduling weight. Range: 1 to 127.
max-bandwidth <RATE>	Specifies the maximum bandwidth rate allowed on the queue in kbps. Range: 64 to 100000000. Alternatively, the maximum bandwidth rate can be configured on the queue as a percentage of the port shape or link bandwidth if a port shape is not configured. The allowed range is 1-100.
burst <SIZE>	Specifies the maximum burst size allowed on the queue in kbytes. Range: 1 to 1073.* Default: 32.

Examples

Assigning DWRR with a weight of **17** to queue **2** in the schedule profile **MySchedule**:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)# dwrr queue 2 weight 17
```

Deleting DWRR for queue **2** from the schedule profile **MySchedule**:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)# no dwrr queue 2
```

Command History

Release	Modification
10.13	Added max-bandwidth and burst parameters.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config-schedule-<NAME>	Administrators or local user group members with execution rights for this command.

flow-control buffer cable-length

```
flow-control buffer cable-length {auto|<VALUE>}
no flow-control buffer cable-length {auto|<VALUE>}
```

Description

Configures the cable length value in meters used to calculate the headroom buffer size if an explicit headroom buffer limit is not configured and flow-control is enabled. If auto is used then the L1 reported cable length is used to calculate the headroom buffer size.

If no headroom limit or cable length configuration is specified (or cable-length auto is specified but hardware is not reporting a cable length) then a default cable length of 350 meters is used to compute the headroom limit.

The **no** form of this command removes the existing cable length configuration.

Parameter	Description
auto	Specifies the interface-detected cable length.
<VALUE>	Specifies cable length value in meters.

Examples

Configuring a cable length of **400** meters on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# flow-control buffer cable-length 400
```

Removing the configured cable length value of **400** meters on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# no flow-control buffer cable-length 400
```

Configuring interface-detected cable length on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# flow-control buffer cable-length auto
```

Removing the configuration for interface-detected cable length on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# no flow-control buffer cable-length auto
```

Command History

Release	Modification
10.13	Command introduced.

Command Information

Platforms	Command context	Authority
8325 9300 10000	config-if	Administrators or local user group members with execution rights for this command.

lldp dcbx (global)

```
lldp dcbx [version {cee|ieee}]
no lldp dcbx [version {cee|ieee}]
```

Description

Enables advertisement of the DCBx TLVs in LLDP packets either globally or per interface. By default, DCBx is disabled in the switch.

The **no** form of this command disables DCBx advertisement.

Parameter	Description
version {cee ieee}	Configures the DCBx version in either CEE (Converged Enhanced Ethernet) or IEEE (IEEE 802.1Qaz). Default version: IEEE

Examples

Enabling DCBx globally with default version:

```
switch(config)# lldp dcbx
```

Disabling DCBx globally:


```
switch(config)# no lldp dcbx
```

Enabling DCBx globally with the CEE version:

```
switch(config)# lldp dcbx version cee
```

Enabling DCBx on an interface with default version:

```
switch(config-if)# lldp dcbx
```

Disabling DCBx on an interface:

```
switch(config-if)# no lldp dcbx
```

Enabling DCBx on an interface with the CEE version:

```
switch(config-if)# lldp dcbx version cee
```

Command History

Release	Modification
10.08	Added optional CEE and IEEE parameters.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config config-if	Administrators or local user group members with execution rights for this command.

lldp dcbx (per interface)

```
lldp dcbx [version {cee|ieee}]  
no lldp dcbx [version {cee|ieee}]
```

Description

Enables DCBx on an interface. By default, an interface follows the global DCBx configuration. DCBx must be enabled globally for the interface configuration to take effect.

The **no** form of this command removes the port-override DCBx configuration and reverts the port behavior to the global DCBx configuration.

Parameter	Description
version { cee ieee }	Configures the DCBx version in either CEE (Converged Enhanced Ethernet) or IEEE (IEEE 802.1Qaz). Default version: IEEE

Usage

If the interface command specifies a different version than the global configuration, it overrides the globally configured DCBx version. If the command is executed without specifying a version, the IEEE version is configured.

Examples

Enabling DCBx on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# lldp dcbx
```

Disabling DCBx on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# no lldp dcbx
```

Enabling DCBx on interface 1/1/1 with the CEE version:

```
switch(config)# interface 1/1/1
switch(config-if)# lldp dcbx version cee
```

Enabling DCBx and configuring PFC for priority 4 on interface 1/1/1.



Priority Flow Control (PFC) commands are only supported on the 8100, 8325, 9300, and 8360 Switch Series.

```
switch(config)# interface 1/1/1
switch(config-if)# lldp dcbx
switch(config-if)# flow-control priority 4
```

Command History

Release	Modification
10.08	Added a parameter 'version'
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325	config-if	Administrators or local user group members with execution rights for this command.

Platforms	Command context	Authority
9300 10000		

lldp dcbx disable

```
lldp dcbx disable
no lldp dcbx disable
```

Description

Disables DCBx on an interface.

The **no** form of this command removes the DCBx configuration from the interface, which will then default to the global configuration.

Usage

If the interface command specifies a different version than the global configuration, it overrides the globally configured DCBx version. If the command is executed without specifying a version, the IEEE version is configured.

Examples

Disabling DCBx on interface 1/1/1:

```
switch(config)# interface 1/1/1
switch(config-if)# lldp dcbx disable
```

Reverting interface 1/1/1 to the global DCBx configuration:

```
switch(config)# interface 1/1/1
switch(config-if)# no lldp dcbx
```

Command History

Release	Modification
10.08	Command introduced.

Command Information

Platforms	Command context	Authority
8320 8325 10000	config-if	Administrators or local user group members with execution rights for this command.

map queue

```
map queue <QUEUE-NUMBER> local-priority <PRIORITY-NUMBER>
no map queue <QUEUE-NUMBER> [local-priority <PRIORITY-NUMBER>]
```

Description

Assigns a local priority to a queue in a queue profile. By default, the larger the queue number the higher its priority. A queue without a local priority value assigned to it is not used to store packets. The same queue can be assigned multiple local priorities.

The **no** form of this command removes the specified local priority from a specific queue. If no local priority number is specified, then all local priorities are removed from the queue.

Parameter	Description
<QUEUE-NUMBER>	Specifies the queue number. Range: 0 to 7.
<PRIORITY-NUMBER>	Specifies the local priority. Range: 0 to 7, where 0 is the lowest priority and 7 is the highest.

Usage

For a queue profile to be complete and ready to be applied, all eight local priorities must be mapped to a queue. Any local priority used by interface Priority-based Flow Control (PFC) must be the only local priority mapped to its queue. In order for PFC pausing to work as intended, no other local priorities should be mapped to that same queue. This queue mapping should be configured during initial switch provisioning and only changed during maintenance periods where all ports are disabled.



NOTE: PFC is only supported on the 8325, 9300, and 10000 Series Switch

The following commands illustrate a valid configuration, where every local priority value is assigned to a queue:

```
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 1 local-priority 2
map queue 3 local-priority 3
map queue 4 local-priority 4
map queue 5 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
```

The following commands illustrate an invalid configuration, because local priority 2 is not assigned to a queue:

```
map queue 0 local-priority 0
map queue 1 local-priority 1
map queue 2 local-priority 3
map queue 3 local-priority 4
map queue 4 local-priority 5
map queue 5 local-priority 6
map queue 5 local-priority 7
```

Examples

Assigning priority **7** to queue **7** in profile **myprofile**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# map queue 7 local-priority 7
```

Removing priority **7** from queue **7** in profile **myprofile**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# no map queue 7 local-priority 7
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config-queue	Administrators or local user group members with execution rights for this command.

name queue

```
name queue <QUEUE-NUMBER> <DESCRIPTION>
no name queue <QUEUE-NUMBER>
```

Description

Assigns a description to a queue in a queue profile. This is for identification purposes and has no effect on configuration.

The **no** form of this command removes the description associated with a queue.

Parameter	Description
<QUEUE-NUMBER>	Specifies the queue number. Range: 0 to 7.
<DESCRIPTION>	Specifies a queue description for identification purposes. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Examples

Assigning the description **priority-traffic** to queue **7**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# name queue 7 priority-traffic
```

Removing the description from queue **7**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)# no name queue 7
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config-queue	Administrators or local user group members with execution rights for this command.

qos cos-map

```
qos cos-map <CODE-POINT> local-priority <PRIORITY-NUMBER> [color <COLOR>] [name <DESCRIPTION>]
no qos cos-map <CODE-POINT>
```

Description

Defines the local priority assigned to incoming packets for a specific 802.1 VLAN priority code point (CoS) value. The CoS map values are used to mark incoming packets when QoS trust mode is set to **cos**. In **trust none** mode, CoS map entry 0 is used to set the port default local priority and color.

To see the default CoS map settings, use the following command:

```
switch# show qos cos-map default
code_point  local_priority  color    name
-----
0           1             green    Best_Effort
1           0             green    Background
2           2             green    Excellent_Effort
3           3             green    Critical_Applications
4           4             green    Video
5           5             green    Voice
6           6             green    Internetwork_Control
7           7             green    Network_Control
```

The **no** form of this command restores the assignments for a CoS map value to the default setting.

Parameter	Description
<CODE-POINT>	Specifies an 802.1 VLAN priority CoS value. Range: 0 to 7. Default 0.
local-priority <PRIORITY-NUMBER>	Specifies a local priority value to associate with the CODE-POINT value. Range: 0 to 7. Default: 0.
color <COLOR>	Reserved for future use.

Parameter	Description
name <DESCRIPTION>	Specifies a description for the CoS setting. The name is for identification only, and has no effect on queue configuration. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Usage

Any code point configured for use by interface Priority-based Flow Control (PFC) must be assigned a unique local priority in the CoS map. No other code point can be assigned that same local priority. This should be configured during initial switch provisioning and only changed during maintenance periods where all ports are disabled.



NOTE: PFC is only supported on the 8325, 9300, and 10000 Series Switch

Examples

Mapping CoS value **1** to a local priority of **2**:

```
switch(config)# qos cos-map 1 local-priority 2
```

Mapping CoS value **1** to the default local priority value:

```
switch(config)# no qos cos-map 1
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config	Administrators or local user group members with execution rights for this command.

qos dscp

```
qos dscp <CODE-POINT>
no qos dscp
```

Description

Configures a differentiated services code point (DSCP) remark for an Ethernet or LAG interface. IPV4 and IPV6 packets that ingress on the interface are remarked at egress using the configured DSCP value.

The remark only occurs when QoS trust mode on the interface is set to **none**. If a DSCP remark is configured and then trust mode is subsequently set to **cos** or **dscp**, then the DSCP remark is ignored.

The following commands will show the remark status as *not applied* (incompatible QoS global or port trust configuration):

- **show running-configuration**
- **show interface <INTERFACE-NAME>**
- **show interface <INTERFACE-NAME> qos**

The **no** form of this command removes a CoS remark on an interface.

Parameter	Description
<code><CODE-POINT></code>	Specifies an IP differentiated services code point value. Range: 0 to 63.

Usage

Order of operation for arriving IPv4 or IPv6 packets:

1. Trust none is applied with initial local-priority and color metadata assigned from the CoS Map entry index 0.
2. The local-priority value and the queue profile are then used to determine the queue for the packet.
3. The remark of the packet's DSCP metadata field is performed. When the packet is transmitted, its IPv4 or IPv6 DS header is remarked with the DSCP metadata.

For arriving non-IP packets:

Trust none is applied with initial local-priority and color metadata assigned from the CoS Map entry index 0. This selects the queue for packet scheduling. The PCP of any tagged non-IP packets is unchanged.

Examples

Configuring a DSCP remark of **43** on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# qos trust none
switch(config-if)# qos dscp 43
```

Deleting a DSCP remark of **43** on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# no dscp 43
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config-if config-lag-if	Administrators or local user group members with execution rights for this command.

qos dscp-map

```
qos dscp-map <CODE-POINT> local-priority <PRIORITY-NUMBER> [color <COLOR>] [name <DESCRIPTION>]
no qos dscp-map <CODE-POINT>
```

Description

Defines the local priority color assigned to incoming packets for a specific IP differentiated services code point (DSCP) value. The DSCP map values are used to prioritize incoming packets when QoS trust mode is set to **dscp**.

The **no** form of this command restores the assignments for a code point to the default setting.

Use **show qos dscp-map** to view the current settings. To see the default DSCP map settings, use the following command:

```
switch# show qos dscp-map default
code_point local_priority color name
-----
0          1             green CS0
1          1             green
2          1             green
3          1             green
4          1             green
5          1             green
...
45         5             green
46         5             green EF
47         5             green
48         6             green CS6
...
61         7             green
62         7             green
63         7             green
```

Parameter	Description
<code><CODE-POINT></code>	Specifies an IP differentiated services code point. Range: 0 to 63. Default: 0.
<code>local-priority <PRIORITY-NUMBER></code>	Specifies a local priority value to associate with the CODE-POINT value. Range: 0 to 7. Default: 0.
<code>color <COLOR></code>	Configures the QoS CoS map color. The supported colors are green, red, and yellow. The default color is green.
<code>cos <PCP-VALUE></code>	Specifies an optional 802.1p VLAN Priority Code Point remark value. Range: 0 to 7. Default: No remark.

Parameter	Description
name <DESCRIPTION>	Specifies a description for the DSCP setting. The name is used for identification only, and has no effect on queue configuration. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Examples

Setting code point **1** to a local priority of **2** and a CoS of **0**:

```
switch(config)# qos dscp-map 1 local-priority 2 cos 0
```

Setting code point **1** to the default value:

```
switch(config)# no qos dscp-map 1
```

Setting code point **1, 3-5** to a local priority **2**, and CoS color as **green** with the description **EntryName**:

```
switch(config)# qos dscp-map 1,3-5 local-priority 2 color green name EntryName
```

Command History

Release	Modification
10.13	Added <COLOR> parameters.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config	Administrators or local user group members with execution rights for this command.

qos pool

```
qos pool <INDEX> lossless size {factory-default|<PERCENT>} percent headroom
{factory-default|<HSIZE>} kbytes [priorities <PRIORITY>]
no qos pool <INDEX> lossless size {factory-default|<PERCENT>} percent headroom
{factory-default|<HSIZE>} kbytes [priorities <PRIORITY>]
```



This command is not supported on the 8320 Switch Series.



It is not recommended to configure PFC on priority 7. This priority is used by protocol and control traffic, and device operation can be negatively impacted by pausing of traffic.

Description

Creates or modifies a packet buffer pool for lossless traffic by configuring lossless pool size, headroom buffer associated with the pool, and the priorities that are mapped to the pool.

The lossless pool size is a percentage of the total available buffer memory on the device. The headroom pool memory is allocated from the lossless pool and is used for storing packets that arrive on a port after pause has been asserted.

The **no** form of this command removes the QoS pool and returns the memory back to the default lossy pool.

Parameter	Description
<INDEX>	Specifies the index of the lossless pool.
factory-default	Specifies the default lossless pool or headroom buffer sizes.
<PERCENT>	Specifies the lossless pool size in percent. Default: 40.6. Range: 10 to 90.
<HSIZE>	Specifies the headroom buffer size in kilobytes. Factory-default size: 3072 kilobytes. Range: 0 to 10240 kilobytes or half the size of the lossless pool, whichever is lower.
<PRIORITY>	Specifies the PFC priorities to be mapped to the pool. Range: 0 to 7. To map multiple priorities, use a comma-separated list. Ex: 1,3,6.

Usage

- If the switch is configured with PFC and upgraded to 10.08 or a later version, the default lossless pool is created for existing priorities.
- It is recommended to configure PFC or link-level flow control RXTX mode and pools only once on a running system. Changing the pool or flow control configuration on a running system will cause a brief disruption in traffic. The reconfiguration also results in packet loss in the data flow. This may affect one, multiple, or all ports depending on the existing configuration and changes made to it.

Examples

Configuring the QoS pool using lossless pool size and headroom buffer size with one or two PFC assigned to it:

```
switch(config)# qos pool 1 lossless size 60 percent headroom 2048 kbytes
priorities 4
switch(config)# qos pool 1 lossless size 60 percent headroom 2048 kbytes
priorities 3,4
switch(config)# qos pool 1 lossless size factory-default percent headroom factory-
default kbytes priorities 3,4
```

Removing the QoS pool configuration:

```
switch(config)# no qos pool 1 lossless size 60 percent headroom 2048 kbytes
priorities 4
```

Removing a priority from the existing QoS pool configuration:

```
switch(config)# qos pool 2 lossless size 70 percent headroom 3072 kbytes
priorities 5,6
switch(config)# qos pool 2 lossless size 70 percent headroom 3072 kbytes
priorities 5
```

Command History

Release	Modification
10.10.1000	Support for 9300 series switches added.
10.10	Support for 3 lossless pools added.
10.09	Support for 10000 series switches added.
10.08	Command introduced.

Command Information

Platforms	Command context	Authority
8325 9300 10000	config	Administrators or local user group members with execution rights for this command.

qos queue-profile

```
qos queue-profile <NAME>
no qos queue-profile <NAME>
```

Description

Creates a new QoS queue profile and switches to the **config-queue** context for the profile. Or, if the specified QoS queue profile exists, this command switches to the **config-queue** context for the profile. A queue profile maps queues to local-priority values. Each profile has one to eight queues numbered 0 to 7. The larger the queue number, the higher its priority during transmission scheduling.

The **no** form of this command removes the specified QoS queue profile. Only profiles that are not currently applied can be removed.

Parameter	Description
<NAME>	Specifies the name of the QoS queue profile to create or configure. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Examples

Creating the profile **myprofile**:

```
switch(config)# qos queue-profile myprofile
switch(config-queue)#
```

Deleting the profile **myprofile**:

```
switch(config)# no qos queue-profile myprofile
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config	Administrators or local user group members with execution rights for this command.

qos schedule-profile

```
qos schedule-profile <NAME>
no qos schedule-profile <NAME>
```

Description

Creates a QoS schedule profile and switches to the **config-schedule** context for the profile. If the specified schedule profile exists, this command switches to the **config-schedule** context for the profile. The schedule profile determines the order in which queues are selected to transmit a packet, and the amount of service defined for each queue.

Parameter	Description
<NAME>	Specifies the name of the QoS queue profile to create or configure. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Usage

Queues in a schedule profile are numbered consecutively starting from zero. Queue zero is the lowest priority queue. The larger the queue number, the higher priority the queue has in scheduling algorithms.

A profile named **factory-default** is defined by default and applied to all interfaces. It cannot be edited or deleted. To see its settings, use the command:

```
switch# show qos schedule-profile factory-default
queue_num algorithm weight
-----
0          dwrr       1
```

1	dwrr	1
2	dwrr	1
3	dwrr	1
4	dwrr	1
5	dwrr	1
6	dwrr	1
7	dwrr	1

A profile named **strict** is predefined and cannot be edited or deleted. The strict profile services all queues of the queue profile to which it is applied, using the strict priority algorithm.

A schedule profile must be defined on all interfaces at all times.

There are two permitted configurations for a schedule profile:

1. All queues use the same scheduling algorithm (for example, DWRR).
2. The highest queue number uses strict priority, and all remaining (lower) queues use the same algorithm (for example, DWRR). This supports priority scheduling behavior necessary for the IETF RFC 3246 Expedited Forwarding specification (<https://tools.ietf.org/html/rfc3246>).

Only limited changes can be made to an applied schedule profile:

- The weight of a dwrr queue.
- The bandwidth of a strict queue or a min-bandwidth queue.
- The algorithm of the highest numbered queue can be swapped between dwrr and strict, and vice versa.

Applicable to REST: Any other changes will result in an unusable schedule profile, and the switch will revert to the **factory-default** profile until the profile is corrected.

The **no** form of this command removes the specified QoS schedule profile when it is not applied. Only profiles that are not currently applied to an interface can be removed.

Examples

Creating the schedule profile **MySchedule**:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)#
```

Deleting the schedule profile **MySchedule**:

```
switch(config)# no qos schedule-profile MySchedule
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config	Administrators or local user group members with execution rights for this command.

qos shape

```
qos shape <RATE> [kbps|percent] [burst <SIZE>]
```

```
no qos shape
```

Description

Limits the egress bandwidth on an interface to a value that is lower than its line rate. An optional burst value may also be specified.

Errors will be generated in the following events:

- A user configures a port-shaping value that is greater than 100 percent
- A user configures a kbps port-shaping value that is less than the supported minimum kbps value. The supported minimum kbps shaping value can be retrieved using the **show capacities** command.

The **no** form of this command removes shaping from an interface.

Parameter	Description
<RATE>	Specifies the maximum traffic rate in kbps. Range: 1 to 100000000. 1 to 400000000 for the Aruba 9300 Switch Series. Alternatively, the bandwidth can also be configured as a percentage of link bandwidth. The supported range is 1-100. Default units are kilobits per second.
<SIZE>	Specifies the maximum burst size in kilobytes. Range: 1 to 128. Default: 32.

Usage

When the traffic rate destined for the port exceeds the configured egress bandwidth, the switch will buffer the excess up to the limit of the queues. Rates larger than the interface's link rate will have no effect. When set on a LAG, each member Ethernet port independently shapes its egress bandwidth to the specified rate.

Examples

Configuring an egress port shaping rate of 400 Mbps on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# qos shape 400000 kbps
```

Configuring an egress port-shaping rate of 40% on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# qos shape 40 percent
```

Deleting egress port shaping on interface **1/1/1**:

```
switch(config)# interface 1/1/1
switch(config-if)# no qos shape
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config-if	Administrators or local user group members with execution rights for this command.

qos threshold-profile

```
qos threshold-profile <NAME>
no qos threshold-profile <NAME>
```

Description

Creates a QoS threshold profile and switches to the **config-threshold** context for the profile. If the specified threshold profile exists, this command switches to the **config-threshold** context for the existing profile. The threshold profile determines the action to take when a threshold is exceeded for each queue.

A threshold profile is composed of up to 8 queues, numbered from 0 to 7. Each queue defines the action to take when buffer utilization exceeds a specific threshold.

Configure queues with the command **queue**.

The **no** form of this command removes the specified QoS threshold profile. Only profiles that are not currently applied to an interface can be removed.

Parameter	Description
<NAME>	Specifies the name of the QoS threshold profile to create or configure. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).

Usage

Queues in a threshold profile can be any valid queue number, although it is also valid to create a threshold profile with no queues specified. Queue zero is the minimum allowed queue number. The maximum allowed queue number may vary by product. For products supporting eight queues, the largest queue number is seven. If an applied threshold profile specifies configuration of a queue number that is not in use based on the configured queue profile, the threshold configuration of that unused queue is ignored.

Examples

Creating the threshold profile **mythreshold**:

```
switch(config)# qos threshold-profile mythreshold
switch(config-threshold)#
```

Deleting the threshold profile **MySchedule**:

```
switch(config)# no qos threshold-profile mythreshold
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8325 9300 10000	config	Administrators or local user group members with execution rights for this command.

qos trust

```
qos trust {none|cos|dscp}
no qos trust
```

Description

In the **config** context:

- This command sets the trust mode that is globally applied to all interfaces that do not have a trust mode configured.
- The **no** form of this command restores all interfaces that do not currently have a trust mode configured to the default setting.

In the **config-if** context:

- This command sets the trust mode override for a specific interface.
- The **no** form of this command clears a trust mode override. The interface then uses the global setting. This is the only way to remove a trust mode override.

Parameter	Description
none	Ignores all packet headers. Ingress packets are assigned the local priority and color values configured for CoS map entry 0. Default.
cos	For 802.1 VLAN tagged packets, use the priority code point field from the outermost VLAN header as the index into the CoS map to obtain the local priority and color values for the packet. If the

Parameter	Description
	packet is untagged, use the local priority and color values configured for CoS map entry 0.
dscp	For IP packets, use the DSCP field from the IP header as the index into the DSCP Map. For non-IP packets with 802.1 VLAN tags, use the priority code point field from the outermost VLAN header as the index into the CoS map to obtain the local priority and color values for the packet. For untagged, non-IP packets, use the local priority and color values configured for CoS map entry 0.

Example

Setting the global trust mode to **dscp**, which is applied to all interfaces that do not already have an individual trust mode configured. An override is then applied to interface **2/2/2**, and LAG 100, setting trust mode to **cos**:

```
switch(config)# qos trust dscp
switch(config)# interface 2/2/2
switch(config-if)# qos trust cos
switch(config-if)# interface lag 100
switch(config-if)# qos trust cos
```



WARNING: QoS port remark configurations are not applied when the QoS trust mode is *mode*. This warning message is seen if a port trust command other than *trust none* is attempted when there is already a remark configuration on the port. To restore the old remark configuration, configure the port trust mode to *none*.



WARNING: QoS port remark configurations are not applied when the global QoS trust mode is *mode*. This warning message is seen if a port *no qos trust* command is attempted when there is already a remark configuration on the port and the global trust mode is not *none*. To re-apply the remark configuration, set the port trust mode to *none*.

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config config-if config-lag-if	Administrators or local user group members with execution rights for this command.

queue action

WRED

Aruba 8320,8325,9300,10000 Switch Series

```
queue <0-7> action wred-resp { green | yellow | red } min-threshold <WRED-MIN-LIMIT>
kbytes max-threshold <WRED-MAX-LIMIT> kbytes max-prob <WRED-MAX-PROB> percent
queue <0-7> action wred-non-resp { green | yellow | red } min-threshold <WRED-MIN-LIMIT>
kbytes max-threshold <WRED-MAX-LIMIT> kbytes max-prob <WRED-MAX-PROB> percent
```

ECN

Aruba 8325, 9300, 10000 Switch Series

```
queue <0-7> action ecn all threshold <LIMIT> kbytes
queue <0-7> action ecn all min-threshold <MIN-LIMIT> kbytes max-threshold <MAX-LIMIT>
kbytes
no queue <0-7>
```

Description

Defines the threshold settings and action for a specified queue in a threshold-profile. For ECN, when queue utilization exceeds the threshold value, ECT (ECN-Capable Transport) packets will be CE (Congestion Encountered) marked when transmitted.

For WRED, when queue utilization exceeds the threshold value, WRED action will randomly early-drop packets to signal congestion. More than one WRED action can be configured on a single queue for different packet colors.

When supported, dual threshold configures an ECN action based on the following queue thresholds:

- When the queue exceeds the minimum threshold limit, the ECN action marks ECT packets CE with increasing probability.
- When the queue exceeds the maximum threshold limit, the ECN action marks all ECT packets CE.

The **no** form of this command removes the settings for a queue.

Parameter	Description
<0-7>	Specifies the queue number. Range: 0 to 7.
<LIMIT>, <MIN-LIMIT>, <MAX-LIMIT>	Specifies the threshold value in kilobytes. Range: 0 to 6000 kilobytes.
<WRED-MIN-LIMIT>	Specifies the queue minimum utilization threshold value for WRED to probabilistically start dropping packets.
<WRED-MAX-LIMIT>	Specifies the queue maximum utilization threshold value for WRED, after which every packet is dropped.
<WRED-MAX-PROB>	Specifies the maximum WRED probability of dropping a packet for the specified queue. NOTE: Applicable only on the Aruba 6400v2, 8100, 8360v2, 8320, 8325, 8400, 9300, and 10000 Switch Series.

Examples

Assigning a threshold to queue **7** in profile **mythreshold**:

```
switch(config)# qos threshold-profile mythreshold
switch(config-threshold)# queue 7 action ecn all threshold 4000 kbytes
```

Assigning a minimum and maximum threshold in kilobytes to queue **5** in profile **mythreshold**:

```
switch(config)# qos threshold-profile mythreshold
switch(config-threshold)# queue 5 action ecn min-threshold 2000 kbytes max-
threshold 4000 kbytes
```

Configuring a responsive WRED action on queue 2 for red-, yellow-, and green-colored packets:

```
switch(config)# qos threshold-profile threshprofile
switch(config-threshold)# queue 2 action wred-resp green min-threshold 100 kbytes
max-threshold 2000 kbytes max-prob 70 percent
switch(config-threshold)# queue 2 action wred-resp yellow min-threshold 80 kbytes
max-threshold 1000 kbytes max-prob 85 percent
switch(config-threshold)# queue 2 action wred-resp red min-threshold 50 kbytes
max-threshold 900 kbytes max-prob 90 percent
```

Configuring a non-responsive WRED action on queue 4 for red-, yellow-, and green-colored packets:

```
switch(config)# qos threshold-profile threshprofile
switch(config-threshold)# queue 4 action wred-non-resp green min-threshold 500
kbytes max-threshold 2100 kbytes max-prob 71 percent
switch(config-threshold)# queue 4 action wred-non-resp yellow min-threshold 250
kbytes max-threshold 1000 kbytes max-prob 82 percent
switch(config-threshold)# queue 4 action wred-non-resp red min-threshold 50 kbytes
max-threshold 900 kbytes max-prob 95 percent
```

Removing a threshold from queue **7** in profile **mythreshold**:

```
switch(config)# qos threshold-profile mythreshold
switch(config-threshold)# no queue 7
```

Command History

Release	Modification
10.11	WRED syntax support introduced in 6200, 6300, 6400, 8320, 8325, 8360, 8400, 9300, and 10000 Switch Series.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	config-threshold	Administrators or local user group members with execution rights for this command.

queue-monitor

[no]queue-monitor

Description

This command enables the queue monitoring feature, allowing the switch to collect queue statistics at 10-second time intervals. then display this information as a queue statistics history. These aggregate statistics can be used to monitor network health, troubleshoot network issues, and identify normal network behavior patterns and performance anomalies. This feature is disabled by default, but can be enabled on a maximum of 52 interfaces per switch.

Examples

Enabling the queue monitoring feature on interface 1/1/1:

```
(config)# interface 1/1/1
(config-if)# queue-monitor
```

Disabling the queue monitoring feature on interface 1/1/1:

```
(config)# interface 1/1/1
(config-if)# no queue-monitor
```

Related Commands

Command	Description
show queue-monitor	This command is used to view the global state of the queue monitor feature. Output includes the memory consumed by the feature, the statistics being monitored for collection, and the interfaces currently enabled for monitoring.

Command History

Release	Modification
10.14.1000	Command introduced on 9300 Switch series
10.14	Command introduced on 8325 and 10000 Switch series

Command Information

Platforms	Command context	Authority
8325 9300 10000	config-if	Administrators or local user group members with execution rights for this command.

rate-limit

```
rate-limit {broadcast|multicast|unknown-unicast} <RATE> {percent|pps}
no rate-limit {broadcast|multicast|unknown-unicast}
```

Description

Sets the amount of traffic of a specific type that can ingress on an Ethernet port, or on each port of a LAG interface. Rate limits are enforced separately on each individual member of a LAG, not on the LAG as a whole.

The Aruba 832x, 9300, and 10000 Switch Series do not allow mixing percentage-based and pps-based rate limits on the same interface at the same time. While the CLI prevents this, REST users who notice unapplied rate limits should check message logs to determine whether the failure is due to invalid configuration.

If an interface's overall rate-limiting configuration becomes invalid, any prior rate-limit settings remain in hardware until explicitly removed. This is to prevent unexpected excess traffic, in the event of accidental misconfiguration. Apart from deletions, other rate-limiting configuration changes on that interface are ignored until conflicts between the hardware unit-types are resolved (i.e., until all of the interface's rate-limit settings use only one of the two unit-types).

The **no** form of this command removes the traffic limit for the specified traffic type.

Parameter	Description
<code>{broadcast multicast unknown-unicast}</code>	Specifies the type of ingress traffic to which the rate limit applies: broadcast , multicast , or unknown-unicast . The multicast rate limit affects multicast and broadcast traffic. The broadcast rate limit only affects broadcast traffic. When both types are applied to the same interface, broadcast packets are limited to the lower of the two rate values. Layer 2 BPDU packets, like spanning tree, are also included in the multicast rate limit.
<code><RATE> {percent pps}</code>	Specifies the rate limit in packets per second or as a percentage of link bandwidth. Range: 64 to 209090910 pps, or 1-100 percent.* For percent mode, the actual rate limit will be approximately equivalent to the minimum of the two 64-kbps step values that are closest to the kbps-converted percentage rate. The actual applied rate limit can be verified using the show interface <IF-NAME> qos command. For percentage mode, rate-limits may be shown as "not applied" until after link-up has occurred on the configured port or LAG.



The Aruba 9300 Switch Series has a range of 64 to 836363640 pps, or 1-100 percent.

Examples

Limiting **multicast** traffic to **4000 pps** on interface **1/1/3**:

```
switch(config)# interface 1/1/3
switch(config-if)# rate-limit multicast 4000 pps
```

Viewing the results of the configuration setting:

```
switch# show interface 1/1/3 qos
Interface 1/3/3 is down (Administratively down)
```

```

Admin state is down
Hardware: Ethernet, MAC Address: 1c:98:ec:e3:6a:00
MTU 1500
Full-duplex
rate-limit multicast 4000 pps (4000 actual)

Speed 0 Mb/s
Auto-Negotiation is turned on
Input flow-control is off, output flow-control is off
RX
    0 input packets          0 bytes
    0 input error            0 dropped
    0 CRC/FCS
L3:
    ucast: 0 packets, 0 bytes
    mcast: 0 packets, 0 bytes
TX
    0 output packets         0 bytes
    0 input error            0 dropped
    0 collision
L3:
    ucast: 0 packets, 0 bytes
    mcast: 0 packets, 0 bytes

```

Limiting **broadcast** traffic to **50 pps** on LAG 100:

```

switch# config
switch(config)# interface 1/1/3
switch(config-if)# interface lag 100
switch(config-if)# rate-limit broadcast 50 pps

```

Configuring an unknown-unicast rate-limit in packets per second:

```

switch(config)# interface 1/1/3
switch(config-if)# rate-limit unknown-unicast 100 pps

```

Configuring a multicast rate-limit as a percentage of link bandwidth:

```

switch(config)# interface 1/1/4
switch(config-if)# rate-limit multicast 1 percent

```

Command History

Release	Modification
10.13	Added the percent parameter.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320	config-if	Administrators or local user group members with execution

Platforms	Command context	Authority
8325 9300 10000		rights for this command.

show dcbx interface

```
show dcbx interface <IFNAME> [peer | vsx-peer]
```

Description

Shows the current DCBx status and the configuration of PFC, ETS, and application priority applied on the interface and the status of the TLVs received from the peer.

Parameter	Description
interface <IFNAME>	Specifies the interface name.
peer	Shows peer DCBx information.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Examples

Showing DCBx on interface 1/1/1 with default DCBx version:

```
switch# show dcbx interface 1/1/1
DCBx admin state      : enabled
DCBx operational state : active
DCBx version          : local = IEEE, remote = IEEE
PFC operational state  : active

Mismatch  Advertisement      Local      Peer
-----  -
Priority Flow Control (PFC)
-> Willing:                  No          Yes
   MACsec Bypass Capability: Yes          Yes
   Max PFC Traffic Classes:  1             1
   Priority 0:                False         False
   Priority 1:                False         False
   Priority 2:                False         False
   Priority 3:                False         False
   Priority 4:                True          True
   Priority 5:                False         False
   Priority 6:                False         False
   Priority 7:                False         False

Enhanced Transmission Selection (ETS)
-> Willing:                  No          Yes
   Credit-Based Shaper:      No          No
   Max Traffic Classes:      8           8
   Priority 0:                Class 1      Class 1
   Priority 1:                Class 0      Class 0
   Priority 2:                Class 2      Class 2
```


Priority 3:	Class 3	Class 3
Priority 4:	Class 4	Class 4
Priority 5:	Class 5	Class 5
Priority 6:	Class 6	Class 6
Priority 7:	Class 7	Class 7
Class 0:	ETS 5	ETS 5
Class 1:	ETS 30	ETS 30
Class 2:	ETS 10	ETS 10
-> Class 3:	ETS 10	ETS 25
-> Class 4:	ETS 25	ETS 10
Class 5:	ETS 10	ETS 10
Class 6:	ETS 10	ETS 10
Class 7:	Strict	Strict

Application Priority Map:

Mismatch	Protocol	Protocol ID	Local Priority	Peer Priority
->	iscsi			5
	tcp-sctp	1001 (0x03E9)	1	1
->	tcp-sctp	1002 (0x03EA)	2	7
	EtherType	2000 (0x07D0)	6	6

Showing DCBx on interface 1/1/1 with CEE version:

```
switch# show dcbx interface 1/1/1
```

```
DCBx admin state      : enabled
DCBx operational state : active
DCBx version          : local = CEE, remote = CEE
PFC operational state  : active
```

Mismatch	Advertisement	Local	Peer

Control			
	Operating Version:	0	0
	Max Version:	0	0
	Sequence Number:	1	1
->	Acknowledgement Number:	1	0
Priority Flow Control (PFC)			
	Operating Version:	0	0
	Max Version:	0	0
	Enabled:	Yes	Yes
->	Willing:	No	Yes
	Error:	No	No
	Max PFC Traffic Classes:	8	8
	Priority 0:	False	False
	Priority 1:	False	False
	Priority 2:	False	False
	Priority 3:	False	False
	Priority 4:	True	True
	Priority 5:	False	False
	Priority 6:	False	False
	Priority 7:	False	False
Priority Group			
	Operating Version:	0	0
	Max Version:	0	0
->	Enabled:	Yes	No
->	Willing:	No	Yes
	Error:	No	No

```

Max Traffic Classes:           8           8
Priority 0:                     PGID 1       PGID 1
Priority 1:                     PGID 0       PGID 0
Priority 2:                     PGID 2       PGID 2
Priority 3:                     PGID 3       PGID 3
Priority 4:                     PGID 4       PGID 4
Priority 5:                     PGID 5       PGID 5
Priority 6:                     PGID 6       PGID 6
Priority 7:                     PGID 7       PGID 7
-> PG0 Percentage:             5           10
-> PG1 Percentage:             10          25
PG2 Percentage:                30          30
PG3 Percentage:                30          30
PG4 Percentage:                30          30
PG5 Percentage:                30          30
PG6 Percentage:                30          30
PG7 Percentage:                30          30

Application Protocol
Operating Version:             0           0
Max Version:                   0           0
Enabled:                       Yes         Yes
-> Willing:                    No          Yes
Error:                         No          No

Application Priority Map:
Mismatch Protocol Protocol ID Local Peer
                        Priority Priority
-----
-> iscsi
   tcp/udp      1001 (0x03E9)  1      1
-> tcp/udp      1002 (0x03EA)  2      7
   EtherType    2000 (0x07D0)  6      6

```

Showing DCBx on interface 1/1/1 with mismatched version:

```

switch# show dcbx interface 1/1/1
DCBx admin state      : enabled
DCBx operational state : version_mismatch
DCBx version          : local = IEEE, remote = CEE
PFC operational state : active

```

Showing DCBx peer connected to interface 1/1/1 using CEE version:

```

switch# show dcbx interface 1/1/1 peer
DCBx version: CEE
Control
  Operating Version      : 0
  Max Version            : 0
  Sequence Number        : 1
  Acknowledgement Number : 1

Priority Flow Control (PFC)
  Operating Version      : 0
  Max Version            : 0
  Enabled                : Yes
  Willing                : No
  Error                  : No
  Max PFC Traffic Classes: 8

```

```

Priority 0      : False
Priority 1      : False
Priority 2      : False
Priority 3      : False
Priority 4      : True
Priority 5      : False
Priority 6      : False
Priority 7      : False

Priority Group
  Operating Version : 0
  Max Version       : 0
  Enabled           : No
  Willing           : Yes
  Error             : No
  Max Traffic Classes : 8
  Priority 0        : PGID 1
  Priority 1        : PGID 0
  Priority 2        : PGID 2
  Priority 3        : PGID 3
  Priority 4        : PGID 4
  Priority 5        : PGID 5
  Priority 6        : PGID 6
  Priority 7        : PGID 7
  PG0 Percentage    : 25
  PG1 Percentage    : 25
  PG2 Percentage    : 10
  PG3 Percentage    : 10
  PG4 Percentage    : 5
  PG5 Percentage    : 5
  PG6 Percentage    : 10
  PG7 Percentage    : 10

Application Protocol
  Operating Version : 0
  Max Version       : 255
  Enabled           : Yes
  Willing           : No
  Error             : No

Application Priority Map:
  Protocol      Protocol ID      Priority
  -----
  iscsi         1001 (0x03E9)      1
  tcp/udp       1002 (0x03EA)      7
  tcp/udp       2000 (0x07D0)      6
  EtherType

```

Command History

Release	Modification
10.08	Updated the output to show the DCBx version enhancement.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show interface queues

show interface <INTERFACE-NAME> queues

Description

Displays interface-level queue statistics.

Parameter	Description
<INTERFACE-NAME>	Specifies the name of an Ethernet port or LAG on the switch. Format: member/slot/port or lag number .

Usage

Statistics include:

- **Tx Bytes:** Total bytes transmitted. The byte count may include packet headers and internal metadata that are removed before the packet is transmitted. Packet headers added when the packet is transmitted may not be included.
- **Tx Packets:** Total packets transmitted.
- **Tx Drops:** Total packets dropped by an egress queue due to insufficient capacity.

Examples

Showing queue statistics for interface **1/1/5**:

```
switch# show interface 1/1/5 queues
Interface 1/1/5 is down
Admin state is up
      Tx Bytes      Tx Packets      Tx Drops
Q0              0             0             3
Q1          15356             73             0
Q2              0             0             0
Q3              0             0             0
Q4              0             0             0
Q5              0             0             0
Q6              0             0             0
Q7              0             0             0
```

Showing queue statistics for interface **lag 1**:

```
switch# show interface lag 1 queues
Aggregate-name lag1
Aggregated-interfaces :
1/1/6 1/1/7
Speed 20000 Mb/s
      Tx Bytes      Tx Packets      Tx Drops
```

Q0	0	0	0
Q1	0	0	0
Q2	0	0	0
Q3	0	0	0
Q4	0	0	0
Q5	0	0	0
Q6	0	0	0
Q7	3450	25	0

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show interface qos

show interface <INTERFACE-NAME> qos

Description

Shows various QoS settings for a specific interface.

Parameter	Description
<INTERFACE-NAME>	Specifies the name of an interface on the switch. Format: member/slot/port or lag number .

Examples

Showing QoS settings for interface **1/1/5**:

```
switch# show interface 1/1/5 qos
Interface 1/1/5 is down
Admin state is up
qos trust none (global)
qos queue-profile factory-default (global)
qos schedule-profile EQSEExample (override)
qos dscp 47
rate-limit unknown-unicast 64 pps (64 actual)
rate-limit broadcast 500 pps (500 actual)
qos shape 200000 kbps (200064 kbps actual) burst 70 (70 actual)
```

Queue	Maximum Bandwidth	Bandwidth Units	Burst (KB)
-----	-----	-----	-----
Q1	10082461	kbps	120

Q4	20164923	kbps	32
Q7	30247384	kbps	120

Showing QoS settings for interface **1/1/6**:



The following example applies only to the Aruba 8320, 8325, and 1000 Switch Series.

```
switch# show interface 1/1/6 qos
Interface 1/1/6 is down
Admin state is up
qos trust none (global)
qos queue-profile factory-default (global)
qos schedule-profile EQSExample (override)
rate-limit multicast 1 percent (100000 actual)
qos shape 200000 kbps (200064 kbps actual) burst 70 (70 actual)
```

Queue	Maximum Bandwidth	Bandwidth Units	Burst (KB)
Q1	10082461	kbps	120
Q4	20164923	kbps	32
Q7	30247384	kbps	120

Showing QoS settings for interface **1/1/6**:



The following example applies only to the Aruba 9300 Switch Series.

```
switch# show interface 1/1/6 qos
Interface 1/1/6 is down
Admin state is up
qos trust none (global)
qos queue-profile factory-default (global)
qos schedule-profile EQSExample (override)
rate-limit multicast 1 percent (400000 actual)
qos shape 200000 kbps (200064 kbps actual) burst 70 (70 actual)
```

Queue	Maximum Bandwidth	Bandwidth Units	Burst (KB)
Q1	10082461	kbps	120
Q4	20164923	kbps	32
Q7	30247384	kbps	120

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show interface qos buffers

show interface <INTERFACE-NAME> qos buffers

Description

Shows interface-level buffer usage statistics.

Parameter	Description
<INTERFACE-NAME>	Specifies the name of an Ethernet port on the switch.

Usage

Usage statistics include:

- **In-Use:** Current buffer memory consumption in kilobytes for each priority.
- **Recent-Peak:** Maximum amount of buffer memory consumed by each priority since the last time the statistics were queried.
- **Max-Use:** Maximum amount of buffer memory consumed by each priority since the statistics were last cleared.
- **Limit:** Specifies the hardware-programmed usage limit for the row. Priorities configured for lossless behavior display two rows:
 - First row displays lossless pool usage and configured XOFF limit.
 - Second row displays headroom pool usage and the configured limit.

The hardware-programmed headroom pool limit is based on configuration of either the **flow-control buffer headroom** or **flow-control buffer cable-length** commands on the interface. If both configurations are specified on the same interface, the headroom buffer limit configuration takes precedence over the cable length. If auto is configured, then the hardware interface detected cable-length is used. If no headroom limit or cable length configuration is specified (or cable-length auto is specified but hardware is not reporting a cable length) then a default cable length of 350 meters is used to compute the headroom limit.

Examples

Showing buffer statistics for interface **1/1/1**:

```
switch# show interface 1/1/1 qos buffers
Buffer usage for 1/1/1 in kbytes:

Cable Length: 350 meters (default)
Recent Peak Interval: 50 seconds
Dynamic thresholds limit = 2^<Alpha>*FreeSpace

Ingress:
  Pri  Pool type      Pool Id  In Use  Recent Peak  Max Use  Limit
```

---	-----	-----	-----	-----	-----	-----
0	Lossy	0	1486	2357	4852	
1	Lossless	1	542	1468	5123	Dyn:1
1	Headroom	1	0	0	6	65
2	Lossless	2	4825	7658	12240	Dyn:1
2	Headroom	2	0	0	0	65
3	Lossy	0	0	0	0	
4	Lossy	0	0	0	0	
5	Lossy	0	0	0	0	
6	Lossy	0	0	0	0	
7	Lossy	0	0	0	0	

Showing buffer statistics for interface **1/1/2**:

```
switch# show interface 1/1/2 qos buffers
Buffer usage for 1/1/2 in kbytes:

Cable Length: 400 meters (user configured)
Recent Peak Interval: 50 seconds
Dynamic thresholds limit = 2^<Alpha>*FreeSpace
```

```
Ingress:
  Pri  Pool type      Pool Id  In Use  Recent Peak  Max Use  Limit
  ---  -
  0    Lossy          0        1486    2357      4852
  1    Lossless       1         542    1468      5123    Dyn:1
  1    Headroom       1          0         0         6        66
  2    Lossless       2        4825    7658     12240    Dyn:1
  2    Headroom       2          0         0         0        66
  3    Lossy          0          0         0         0
  4    Lossy          0          0         0         0
  5    Lossy          0          0         0         0
  6    Lossy          0          0         0         0
  7    Lossy          0          0         0         0
```

Showing buffer statistics for interface **1/1/3**:

```
switch# show interface 1/1/3 qos buffers
Buffer usage for 1/1/3 in kbytes:

Cable Length: 50 meters (hardware detected)
Recent Peak Interval: 50 seconds
Dynamic thresholds limit = 2^<Alpha>*FreeSpace
```

```
Ingress:
  Pri  Pool type      Pool Id  In Use  Recent Peak  Max Use  Limit
  ---  -
  0    Lossy          0        1486    2357      4852
  1    Lossless       1         542    1468      5123    Dyn:1
  1    Headroom       1          0         0         6        55
  2    Lossless       2        4825    7658     12240    Dyn:1
  2    Headroom       2          0         0         0        55
  3    Lossy          0          0         0         0
  4    Lossy          0          0         0         0
  5    Lossy          0          0         0         0
  6    Lossy          0          0         0         0
  7    Lossy          0          0         0         0
```

Command History

Release	Modification
10.11	Command introduced.

Command Information

Platforms	Command context	Authority
8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos cos-map

show qos cos-map [default] [vsx-peer]

Description

Shows the global QoS CoS code point settings, or the factory default settings.

Parameter	Description
default	Shows the factory default CoS code point settings.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Examples

Showing the current CoS map:

```
switch# show qos cos-map
code_point local_priority color name
-----
0          2             green Best_Effort
1          0             green Background
2          1             green Spare
3          3             green Excellent_Effort
4          4             green Controlled_Load
5          5             green Video
6          6             green Voice
7          7             green Network_Control
```

Showing the default CoS map:

```
switch# show qos cos-map default
code_point local_priority color name
-----
0          1             green Best_Effort
1          0             green Background
2          2             green Excellent_Effort
3          3             green Critical_Applications
4          4             green Video
```

5	5	green	Voice
6	6	green	Internetwork_Control
7	7	green	Network_Control

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8320 8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos dscp-map

show qos dscp-map [default] [vsx-peer]

Description

Displays the current or default global QoS dscp-map.

Parameter	Description
default	Shows the factory default DSCP code point settings.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Examples

Showing the current QoS DSCP map:

```
switch# show qos dscp-map
code_point  local_priority  color  name
-----
0           1              green  CS0
1           1              green
2           1              green
3           1              green
4           1              green
5           1              green
6           1              green
7           1              green
8           0              green  CS1
...
```

45	5	green	
46	7	green	EF
47	5	green	
48	6	green	CS6
...			
61	7	green	
62	7	green	
63	7	green	

Showing the default QoS DSCP map:

```
switch# show qos dscp-map default
code_point local_priority color name
-----
0          1              green CS0
1          1              green
2          1              green
3          1              green
4          1              green
5          1              green
...
45         5              green
46         5              green EF
47         5              green
48         6              green CS6
...
61         7              green
62         7              green
63         7              green
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos pool

```
show qos pool [config [vsx-peer] | statistics]
```

Description

Shows the QoS pool buffer configuration and analysis information.

Parameter	Description
config	Shows the current pool configuration.
statistics	Shows the pool statistics.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Examples

Showing the QoS global pool configuration and analysis:

```
switch# show qos pool

Global Packet-Buffer Overview (in Kbytes)

Packet Buffer Pools      Size      Port Limits      Max Usage
-----
Lossy Pool               4105      --              3242

Lossless Pool 1          13080 (50.00%)  --              9235
  Headroom Pool 1        2058      --              0
    Sum of port headroom limits  --              198      --

Lossless Pool 2          8973 (37.10%)  --              8467
  Headroom Pool 2        2058      --              974
    Sum of port headroom limits  --              265      --
-----
TOTAL:                   32768

Analysis:
  Lossy Pool      was 79% utilized
  Lossless pool 1 was 71% utilized
  Headroom pool 1 is 90% over-provisioned
  Lossless pool 2 was 94% utilized
  Headroom pool 2 is 87% over-provisioned

...
```

Showing the QoS current pool configuration:

```
switch# show qos pool config

Global Packet-Buffer Configuration (in Kbytes)

Packet Buffer Pools      Size      Port Limits
-----
Lossy Pool               4105      --

Lossless Pool 1          13080 (50.00%)  --
  Headroom Pool 1        2058      --
    Sum of port headroom limits  --              198

Lossless Pool 2          8973 (37.10%)  --
```

```

      Headroom Pool 2                2058                --
      Sum of port headroom limits                265
-----
TOTAL:                32768

Lossless Pool 1                XOFF  Headroom
-----
1/1/1 priority 3                Dynamic 2                49
1/1/4 priority 3                Dynamic 2                49
1/1/50:1-1/1/50:2 priority 3  Dynamic 2                50
-----
TOTAL:                --                198

Lossless Pool 2                XOFF  Headroom
-----
1/1/4 priority 5                Dynamic 2                49
1/1/16 priority 5              Dynamic 2                24
1/1/32-1/1/33 priority 5      Dynamic 2                96
-----
TOTAL:                --                265
...

```

Showing the QoS pool statistics:

```

switch# show qos pool statistics

Packet-Buffer Pool Statistics

Recent-Peak Interval: 37 seconds

Packet Buffer Pools    Total Size    Peak Use    Recent Peak    Current Use
-----
Lossy Pool             6144         4852         2357           1486
Lossless Pool 1        10240        5123         1468           542
  Headroom 1           2048          6            0              0
Lossless Pool 2        12280       12240        7658          4825
  Headroom 2           2048          0            0              0

All buffer sizes are in Kbytes.
...

```

Command History

Release	Modification
10.10.1000	Support for 9300 series switches added.
10.09	Support for 10000 series switches added.
10.08	Command introduced.

Command Information

Platforms	Command context	Authority
8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos queue-profile

```
show qos queue-profile [<NAME> | factory-default] [vsx-peer]
```

Description

Shows the status of all queue profiles, or a specific queue profile.

Parameter	Description
<NAME>	Specifies the name of a queue profile. Range 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).
[factory-default]	Specifies the factory default queue profile.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Usage

The status of a queue profile can be:

- Applied - The profile is actively being used by the switch.
- Complete - The profile meets the criteria to be applied.
- Incomplete - The profile does not meet the criteria to be applied.

For a queue profile to be complete and ready to be applied:

- All eight local priorities must be mapped to some queue.
- There can be 1 to 8 queues.

Examples

Showing the settings of the factory default queue profile:

```
switch# show qos queue-profile factory-default
queue_num local_priorities name
-----
0          0
1          1
2          2
3          3
4          4
5          5
6          6
7          7
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos schedule-profile

show qos schedule-profile [<NAME> | factory-default | strict] [vsx-peer]

Description

Shows the status of all schedule profiles, or a specific schedule profile.

Parameter	Description
<NAME>	Specifies the name of a queue or schedule profile. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).
[factory-default]	Specifies the factory default queue profile.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Usage

The status of a schedule profile can be:

- Applied - The profile is actively being used by one or more ports.
- Complete - The profile meets the criteria to be applied.
- Incomplete - The profile does not meet the criteria to be applied.

For a schedule profile to be complete and ready to be applied it must have:

- An algorithm for each queue defined by the applied queue profile.
- All queues must use the same algorithm except for the highest numbered queue, which may be strict.

Example

Showing the status of all schedule profiles:

```
switch# show qos schedule-profile
profile_status profile_name
-----
```

applied	MySchedule
complete	factory-default
complete	Test

Showing the configuration of factory default schedule profile:

```
switch# show qos schedule-profile factory-default
Queue
Number  Algorithm      Weight
-----  -
0       dwrr           1
1       dwrr           1
2       dwrr           1
3       dwrr           1
4       dwrr           1
5       dwrr           1
6       dwrr           1
7       dwrr           1
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos threshold-profile

```
show qos threshold-profile [<NAME>] [vsx-peer]
```

Description

Shows the status of all threshold profiles, or a specific threshold profile.

Parameter	Description
<NAME>	Specifies the name of a threshold profile. Range: 1 to 64 alphanumeric characters, including period (.), underscore (_), and hyphen (-).
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Usage

Status fields are:

- applied: The threshold profile is applied to all configured ports.
- partially applied: The threshold profile is applied to some configured ports but failed on other configured ports.
- not applied: The threshold profile could not be applied to any configured ports.
- blank field: The threshold profile is not applied in the configuration (globally or as a port override).

Examples

Showing the status of threshold profile **mythreshold**:

On the 8325, 9300, and 10000 switches:

```
switch# show qos threshold-profile mythreshold
```

Queue	Action	Color	Minimum	Maximum	Max Probability
2	wred-resp	green	1000	1200	30
2	wred-non-resp	red	1100	1150	40
3	ecn	all	1200	1300	

Port	Status
1/1/1	applied
1/1/8	applied
1/2/2	applied

On the 8100 and 8360 switches:

```
switch# show qos threshold-profile mythreshold
```

queue_num	action	Threshold
5	ecn	40
7	ecn	50

Ports	Status
1/1/1	applied
1/1/8	applied
1/2/2	applied

Command History

Release	Modification
10.08	Updated the example.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show qos trust

show qos trust [default] [vsx-peer]

Description

Shows the global QoS trust settings, or the factory default settings.

Parameter	Description
default	Shows the factory default QoS trust settings.
vsx-peer	Shows the output from the VSX peer switch. If the switches do not have the VSX configuration or the ISL is down, the output from the VSX peer switch is not displayed. This parameter is available on switches that support VSX.

Examples

Showing the current QoS trust settings:

```
switch# show qos trust
qos trust cos
```

Showing the default QoS trust settings:

```
switch# show qos trust default
qos trust none
```

Command History

Release	Modification
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

show queue-monitor

```
show [<IFNAME>] queue-monitor
  list [min-depth kbytes <value>]
  table [min-depth kbytes <value>]
  histogram
  status
```

Description

Queue statistics history can be displayed in the CLI as either a list or a histogram reflecting the last eight hours of queue statistics history, or as a table, showing the previous five minutes of queue statistics history, allowing network administrators to view any of the following:

- A summary of congestion events across the switch
- Per-interface histograms of congestion data
- Per-interface detailed view of congestion data

Parameter	Description
list	Displays queue history information in a list format
table	Displays queue history information in a table.
histogram	Displays queue history information in a histogram.
min-depth kbytes <value>	Displays queue information only for those queues larger than the specified minimum depth in Kilobytes.
status	Show whether the feature is enabled or disabled, and information about memory consumption and monitored interfaces.

Examples

Displaying the status of the queue statistics feature, when the feature is disabled (the default status).

```
switch# show queue-monitor status
Feature Status: Stopped
Memory Consumption: 0 MB
Monitored Interfaces/Max: 0/52
```

Displaying the status of the queue statistics feature, when the feature is enabled.

```
switch# show queue-monitor status
Feature Status: Running
Memory Consumption: 72 MB
Statistics Collected: Queue Tx Recent Peak Depth
Monitored Interfaces/Max: 52/52
Monitored Interfaces: 1/1/1-1/1/52
```

Displaying the queue statistics history data for all monitored interfaces in list format:

```
switch# show queue-monitor list
Interface Queue Statistics History
Time Range: 2023-08-07 09:11:30 to 2023-08-08 17:11:30 (UTC)
Sort: Time-ascending, Interface, Queue
Min-Depth Filter: 1 KB
```

Time	Interface	Queue	Depth (B)
2023-08-08 17:11:10	1/1/8	4	2048
2023-08-08 17:11:10	1/1/8	7	2048
2023-08-08 17:11:10	1/1/9	4	2048

2023-08-08 17:11:10	1/1/9	7	2048
2023-08-08 17:11:15	1/1/8	4	2048
2023-08-08 17:11:15	1/1/8	7	2048
2023-08-08 17:11:15	1/1/9	4	2048
2023-08-08 17:11:15	1/1/9	7	2048
2023-08-08 17:11:20	1/1/8	4	2048
2023-08-08 17:11:20	1/1/9	4	2048
2023-08-08 17:11:25	1/1/8	4	2048
2023-08-08 17:11:25	1/1/9	4	2048

The following example displays a portion of output of the **show interface 1/1/8 queue-monitor** command, showing queue statistics history data in tabular format for interface 1/1/8 with a min-depth of 2 kbytes.

```
switch# show interface 1/1/8 queue-monitor table min-depth kbytes 2
1/1/8 Statistics History
Collection Interval: 10 seconds
Time Range: 2023-08-08 17:07:00 to 2023-08-08 17:12:00 (UTC)
Min-Depth Filter: 2 KB
```

Time (s)	Queue Depth - bytes							
	Q0	Q1	Q2	Q3	Q4	Q5	Q6	Q7
-300	-	-	-	-	-	-	-	-
-295	-	-	-	-	-	-	-	-
-290	-	-	-	-	-	-	-	-
-285	-	-	-	-	-	-	-	-
-280	-	-	-	-	-	-	-	-
-275	-	-	-	-	-	-	-	-
-270	-	-	-	-	-	-	-	-
-265	-	-	-	-	-	-	-	-
-260	-	-	-	-	-	-	-	-
-255	-	-	-	-	-	-	-	-
-250	-	2078	-	-	-	-	-	-
-245	-	9914	-	-	-	-	-	-
-240	-	171102	-	-	-	-	-	-
-235	-	996262	-	-	-	-	-	-
-230	-	5444	-	-	-	-	-	-
-225	-	-	-	-	-	-	-	-
-220	-	-	-	-	-	-	-	-
-215	-	-	-	-	-	-	-	-
-210	-	-	-	-	-	-	-	-
-205	-	-	-	-	-	-	-	-
-200	-	-	-	-	-	-	-	-
...								

Displaying the queue statistics history data for interface 1/1/8 in histogram format.

```
switch(config)# show interface 1/1/8 queue-monitor histogram
1/1/8 Statistics History
Time Range: 2023-08-08 09:12:15 to 2023-08-08 17:12:15 (UTC)
Samples per Queue: 13
```

Depth (KB)	Count at Depth							
	Q0	Q1	Q2	Q3	Q4	Q5	Q6	Q7
0	13	13	13	13	7	13	13	10
1-64	0	0	0	0	6	0	0	3
65-128	0	0	0	0	0	0	0	0

129-256	0	0	0	0	0	0	0	0
257-512	0	0	0	0	0	0	0	0
513-1024	0	0	0	0	0	0	0	0
1025-2048	0	0	0	0	0	0	0	0
2049-4096	0	0	0	0	0	0	0	0
4097-8192	0	0	0	0	0	0	0	0
8193-16384	0	0	0	0	0	0	0	0
16385-32768	0	0	0	0	0	0	0	0
32769-65536	0	0	0	0	0	0	0	0

Related Commands

Command	Description
queue-monitor	This command enables the queue monitoring feature, allowing the switch to collect queue statistics at 10-second time intervals. These aggregate statistics can be used to monitor network health, troubleshoot network issues, and identify normal network behavior patterns and performance anomalies. This feature is disabled by default, but can be enabled on a maximum of 52 interfaces per switch.

Command History

Release	Modification
10.14.1000	Command introduced on 9300 Switch series
10.14	Command introduced on 8325 and 10000 Switch series

Command Information

Platforms	Command context	Authority
8325 9300 10000	Operator (>) or Manager (#)	Operators or Administrators or local user group members with execution rights for this command. Operators can execute this command from the operator context (>) only.

strict queue

```
strict queue <0-7> [max-bandwidth <RATE> [kbps|percent]][burst <SIZE>]]
no strict queue <0-7> [max-bandwidth <RATE> [kbps|percent]][burst <SIZE>]]
```

Description

Assigns the strict priority algorithm to a queue. Strict priority services all packets waiting in a queue, before servicing the packets in lower priority queues.

Egress queue shaping can be configured using the **max-bandwidth** and **burst** options to limit the amount of traffic transmitted per output queue. The buffer associated with each egress queue stores the excess traffic to absorb bursts and smooth the output rate. Sustained rates of traffic above the maximum bandwidth will eventually fill the output queue, causing tail drops. Use **show interface <IF-NAME> queues** to determine if any tail-drop errors have occurred.

The **no** form of this command removes the the queue configuration from the schedule profile. To remove only egress queue shaping, re-enter the strict queue command without the **max-bandwidth** and **burst** parameters.

Parameter	Description
<code><QUEUE-NUMBER></code>	Specifies the number of the queue. Range: 0 to 7.
<code>max-bandwidth <BANDWIDTH></code>	Specifies the maximum bandwidth allowed on the queue in Kbps. Range: 8 to 100000000. Alternatively, the maximum bandwidth rate can also be configured on the queue as a percentage of the port shape or link bandwidth if a port shape is not configured. The allowed range is 1-100.
<code>burst <SIZE></code>	Specifies the maximum burst allowed on the queue. Aruba 8320, 8325 and 10000 Switch Series: Range: 1 to 1073. Default: 32. Aruba 9300 Switch Series: Range: 1 to 1024. Default: 32. Aruba 8400 Switch Series: Range: 1 to 32. Default: 32.

Usage

Either all the queues of the schedule profile can be *strict* or just the highest numbered queue. When applied to a LAG, each member Ethernet port independently schedules its egress transmissions using the strict settings. Only limited changes can be made to a *strict* queue that is part of an applied schedule profile:

- The max-bandwidth settings.
- The highest numbered queue can be swapped between *strict* and *dwrr*

Any other changes or removing a queue (**no strict queue**) will result in an unusable schedule profile. If that schedule profile is applied in the interface context, the switch will revert to the schedule profile applied in the global context until the profile is corrected. If that schedule profile is applied in the global context, the switch will revert to using the factory-default profile until the profile is corrected.

It is possible for the following errors to occur:

- **The max-bandwidth cannot be greater than 100 percent.**

This error occurs when a max-bandwidth value greater than 100 percent is configured on a queue.

- **The max-bandwidth cannot be less than <NUM> kbps.**

This error occurs when a kbps max-bandwidth value less than the supported minimum kbps shape value is configured on a queue. The supported minimum kbps shape value can be retrieved using the **show capacities** command.

Examples

Assigning strict priority to queue **7** in the schedule profile **MySchedule**:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)# strict queue 7
```

Deleting strict priority from queue **7** in the schedule profile **MySchedule**:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)# no strict queue 7
```

Assigning strict priority to queue **7** in the schedule profile **MySchedule** with a maximum bandwidth of 10000 Kbps and a burst of 62 Kbytes:

```
switch(config)# qos schedule-profile myschedule
switch(config-schedule)# strict queue 7 max-bandwidth 10000 kbps
```

Assigning strict priority to queue 7 in the schedule profile MySchedule with a maximum bandwidth of 40 percent and a burst of 62 Kbytes:

```
switch(config)# qos schedule-profile MySchedule
switch(config-schedule)# strict queue 7 max-bandwidth 40 percent
```

Command History

Release	Modification
10.13	Added the kbps and percent parameters.
10.07 or earlier	--

Command Information

Platforms	Command context	Authority
All platforms	config-schedule- <i><NAME></i>	Administrators or local user group members with execution rights for this command.

Accessing HPE Aruba Networking Support

HPE Aruba Networking Support Services	https://www.arubanetworks.com/support-services/
AOS-CX Switch Software Documentation Portal	https://www.arubanetworks.com/techdocs/AOS-CX/help_portal/Content/home.htm
HPE Aruba Networking Support Portal	https://networkingsupport.hpe.com/home
North America telephone	1-800-943-4526 (US & Canada Toll-Free Number) +1-408-754-1200 (Primary - Toll Number) +1-650-385-6582 (Backup - Toll Number - Use only when all other numbers are not working)
International telephone	https://www.arubanetworks.com/support-services/contact-support/

Be sure to collect the following information before contacting Support:

- Technical support registration number (if applicable)
- Product name, model or version, and serial number
- Operating system name and version
- Firmware version
- Error messages
- Product-specific reports and logs
- Add-on products or components
- Third-party products or components

Other useful sites

Other websites that can be used to find information:

HPE Aruba Networking Developer Hub	https://developer.arubanetworks.com/hpe-aruba-networking-aoscx/docs/about
Airheads social forums and Knowledge Base	https://community.arubanetworks.com/
HPE Aruba Networking Hardware	https://www.arubanetworks.com/techdocs/hardware/DocumentationPortal/Content/home.htm

Documentation
and Translations
Portal

HPE Aruba
Networking
software <https://networkingsupport.hpe.com/downloads>

Software
licensing and
Feature Packs <https://lms.arubanetworks.com/>

End-of-Life
information <https://www.arubanetworks.com/support-services/end-of-life/>

Accessing Updates

You can access updates from the HPE Aruba Networking Support Portal at <https://networkingsupport.hpe.com>.

Some software products provide a mechanism for accessing software updates through the product interface. Review your product documentation to identify the recommended software update method.

To subscribe to eNewsletters and alerts:

<https://networkingsupport.hpe.com/notifications/subscriptions> (requires an active HPE Aruba Networking Support Portal account to manage subscriptions). Security notices are viewable without an HPE Aruba Networking Support Portal account.

Warranty Information

To view warranty information for your product, go to <https://www.arubanetworks.com/support-services/product-warranties/>.

Regulatory Information

To view the regulatory information for your product, view the *Safety and Compliance Information for Server, Storage, Power, Networking, and Rack Products*, available at <https://www.hpe.com/support/Safety-Compliance-EnterpriseProducts>

Additional regulatory information

HPE Aruba Networking is committed to providing our customers with information about the chemical substances in our products as needed to comply with legal requirements, environmental data (company programs, product recycling, energy efficiency), and safety information and compliance data, (RoHS and WEEE). For more information, see <https://www.arubanetworks.com/company/about-us/environmental-citizenship/>.

Documentation Feedback

HPE Aruba Networking is committed to providing documentation that meets your needs. To help us improve the documentation, send any errors, suggestions, or comments to Documentation Feedback (docsfeedback-switching@hpe.com). When submitting your feedback, include the document title, part number, edition, and publication date located on the front cover of the document. For online help

content, include the product name, product version, help edition, and publication date located on the legal notices page.